

TornadoVM: Transparent Hardware Acceleration for Java...and Beyond!

Juan Fumero, PhD

Research Fellow at The University of Manchester, UK

@snatverk

JavaZone#Oslo 2021

9th Dec 2021

MANCHESTER
1824

The University of Manchester



TORNADOVM

Outline

1. Background
 1. Why TornadoVM, why now?
2. Overview TornadoVM
 1. TornadoVM as multi-backend
 2. Discussion of each backend
3. Multi-Backend: OpenCL, PTX & SPIR-V
4. Performance Evaluation
5. Interoperability with Python, R, ...
6. Reusability and Modularity (Opening TornadoVM Components)
7. Use cases and how TornadoVM is being piloted in Industry
8. Conclusions



Motivation

Current Computer Systems



GPUs



Integrated GPUs

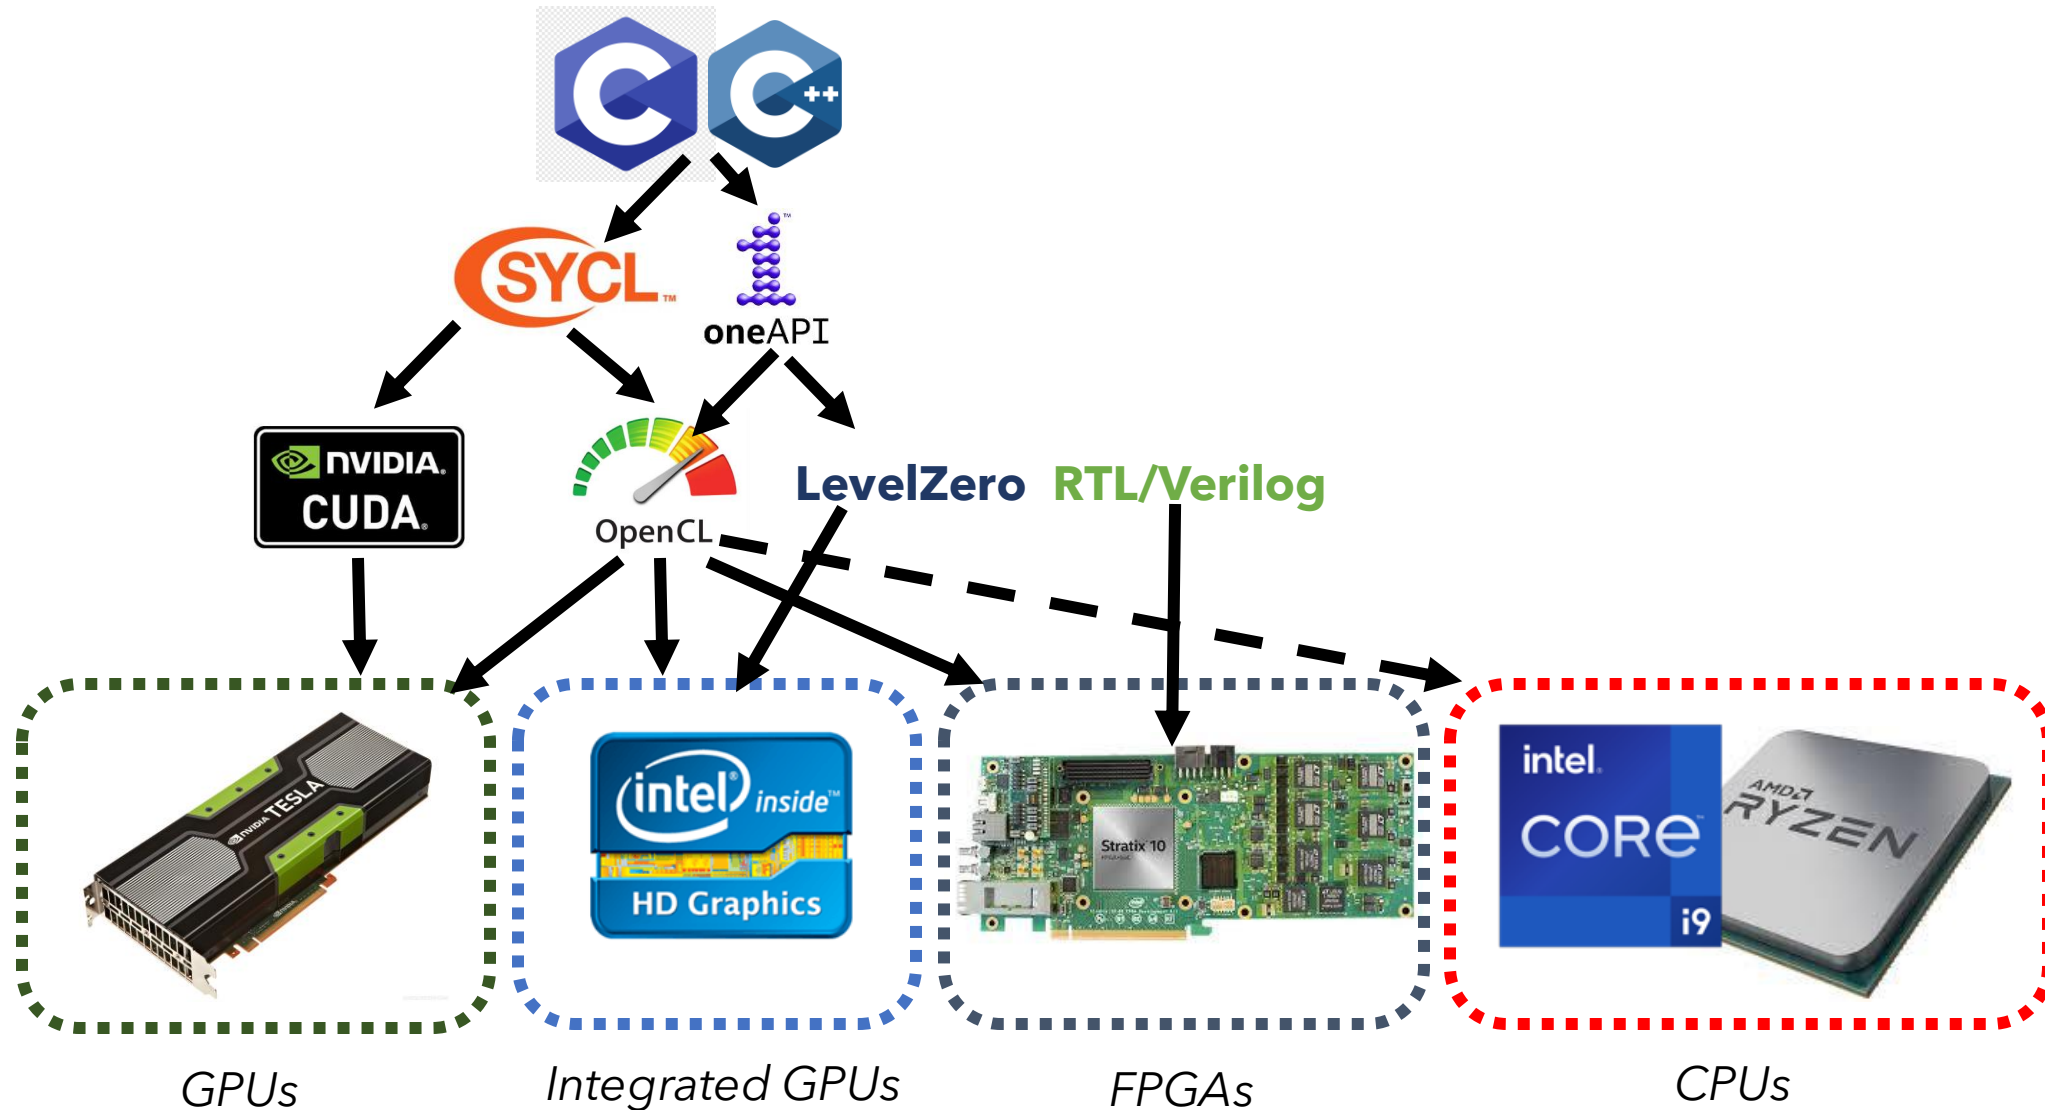


FPGAs

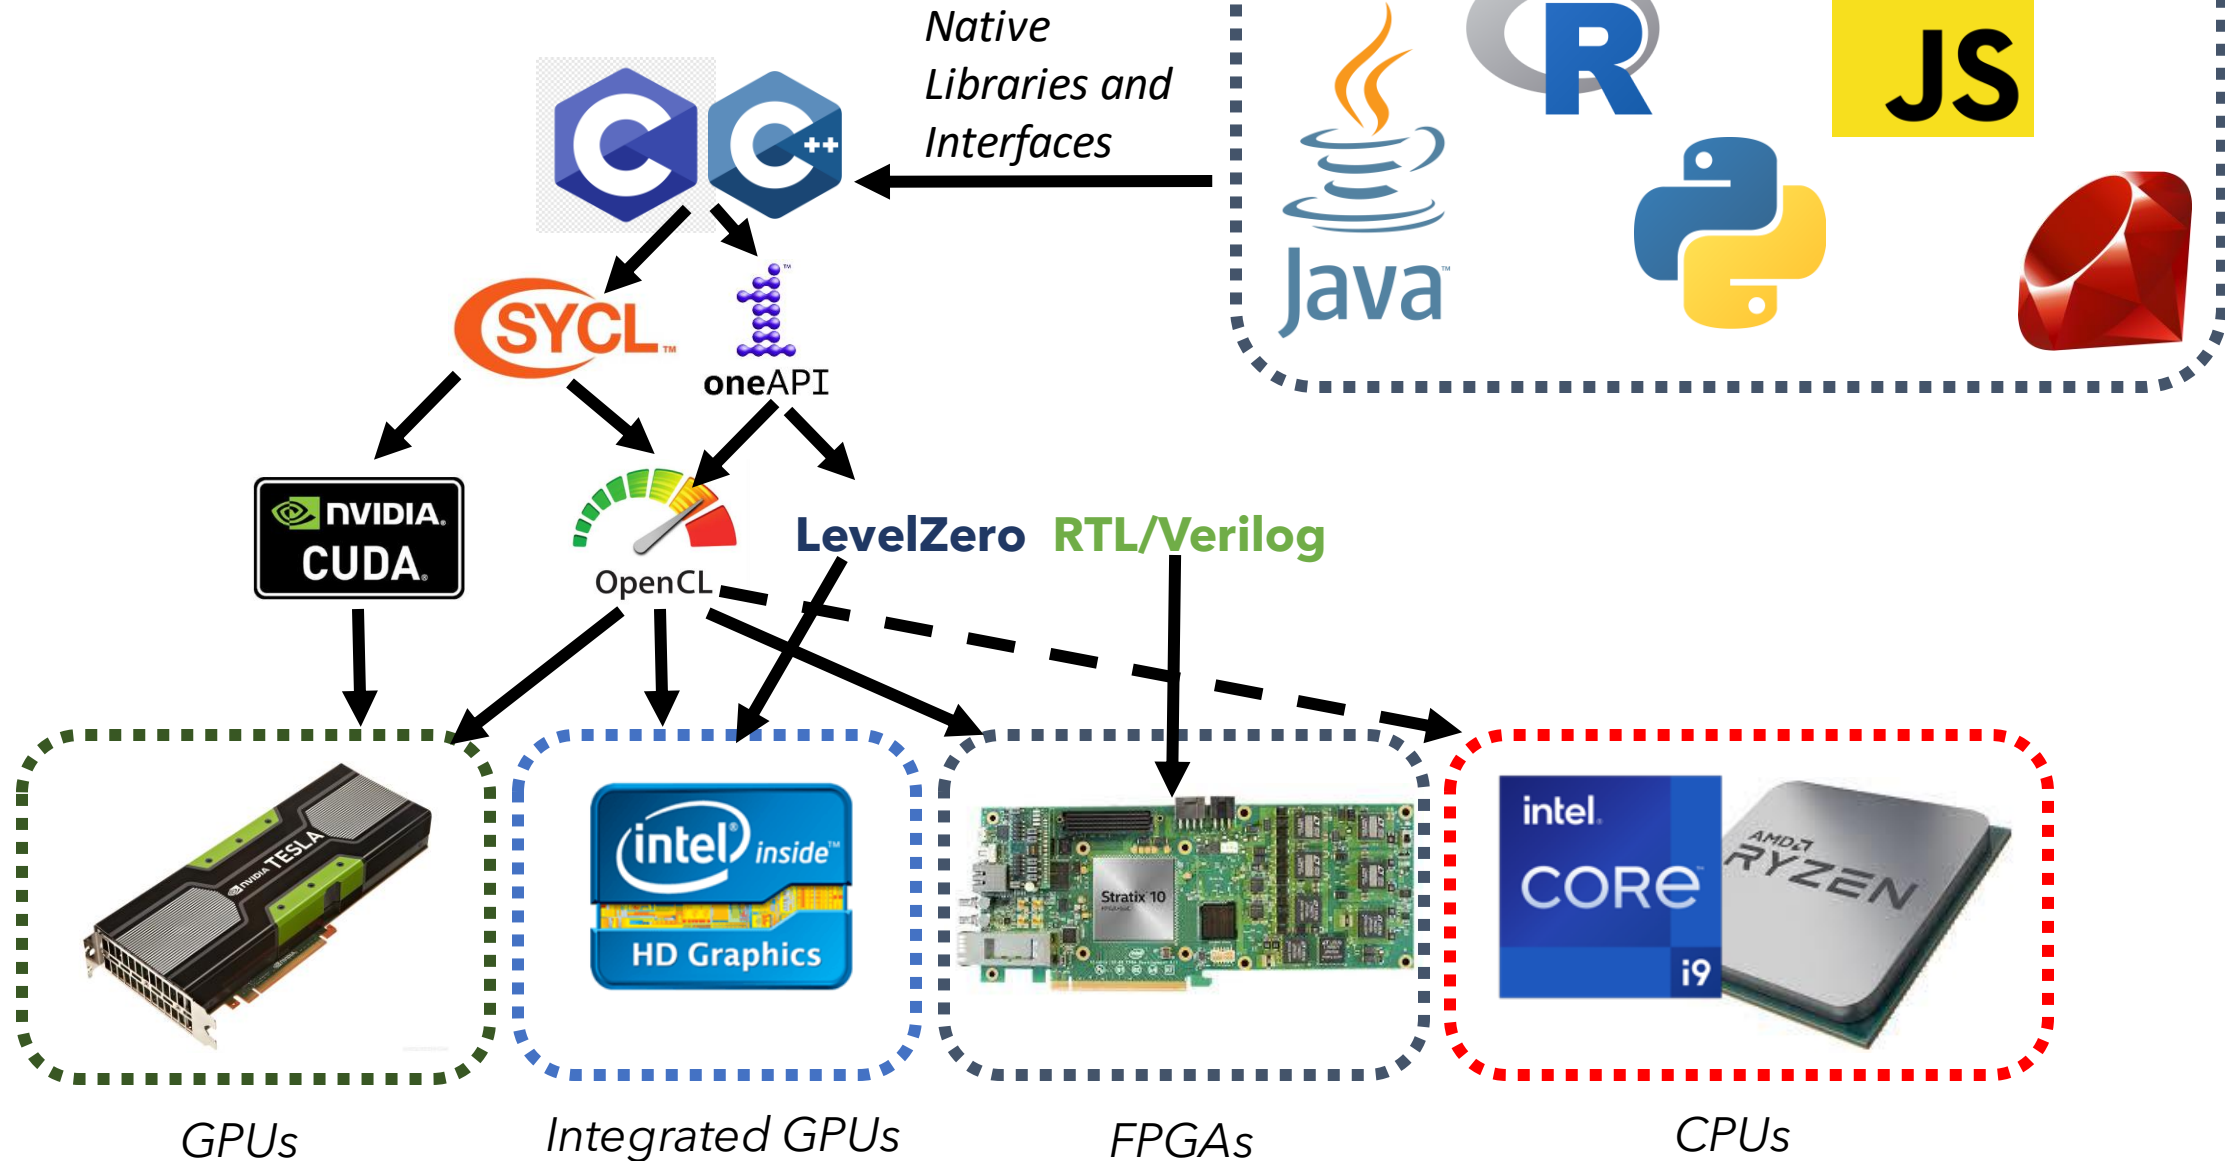


CPUs

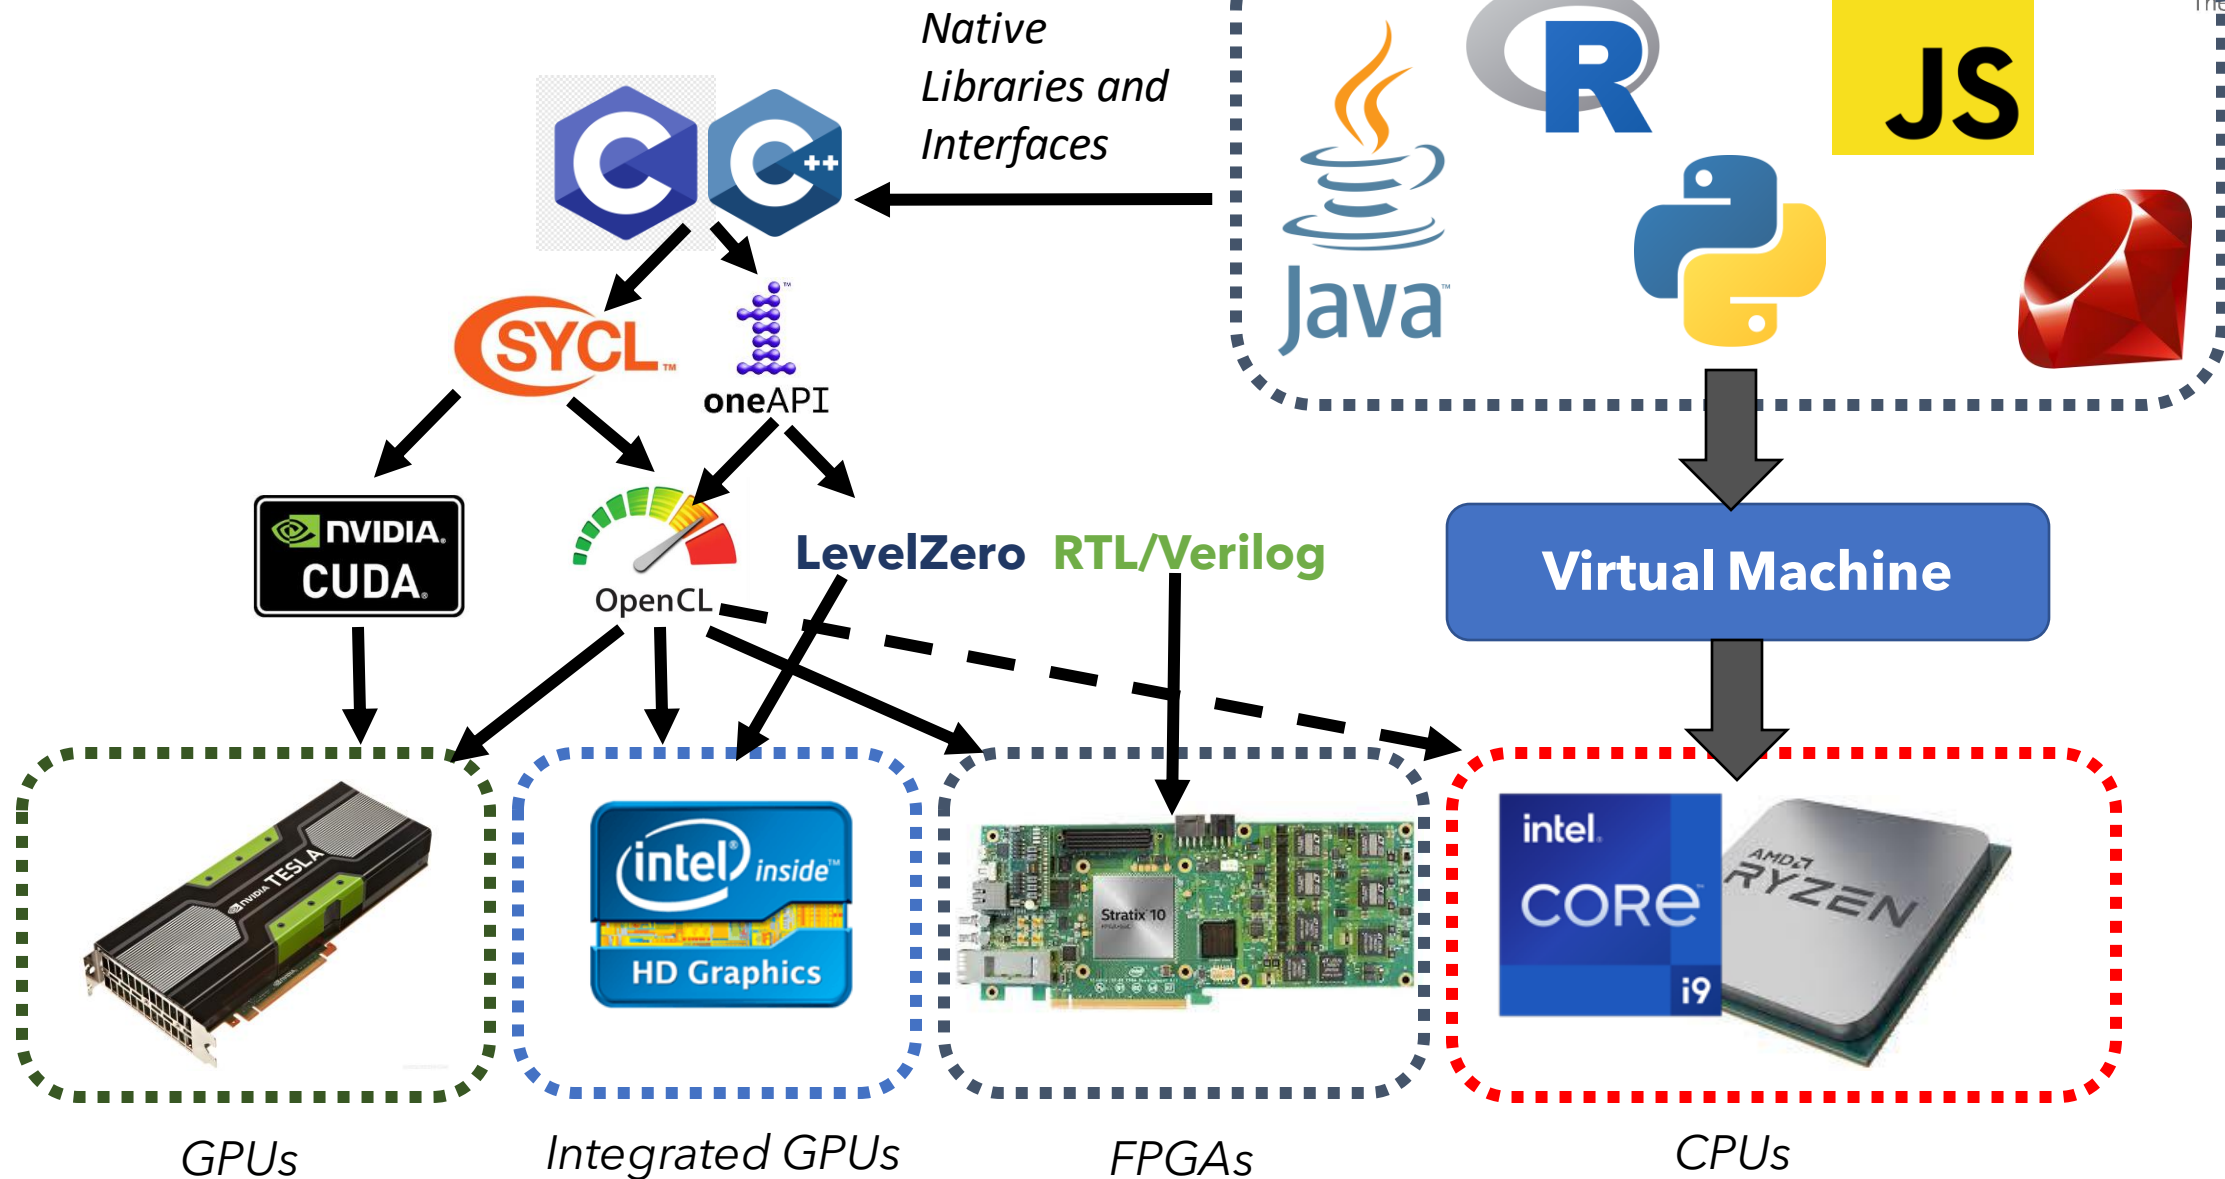
Current Computer Systems



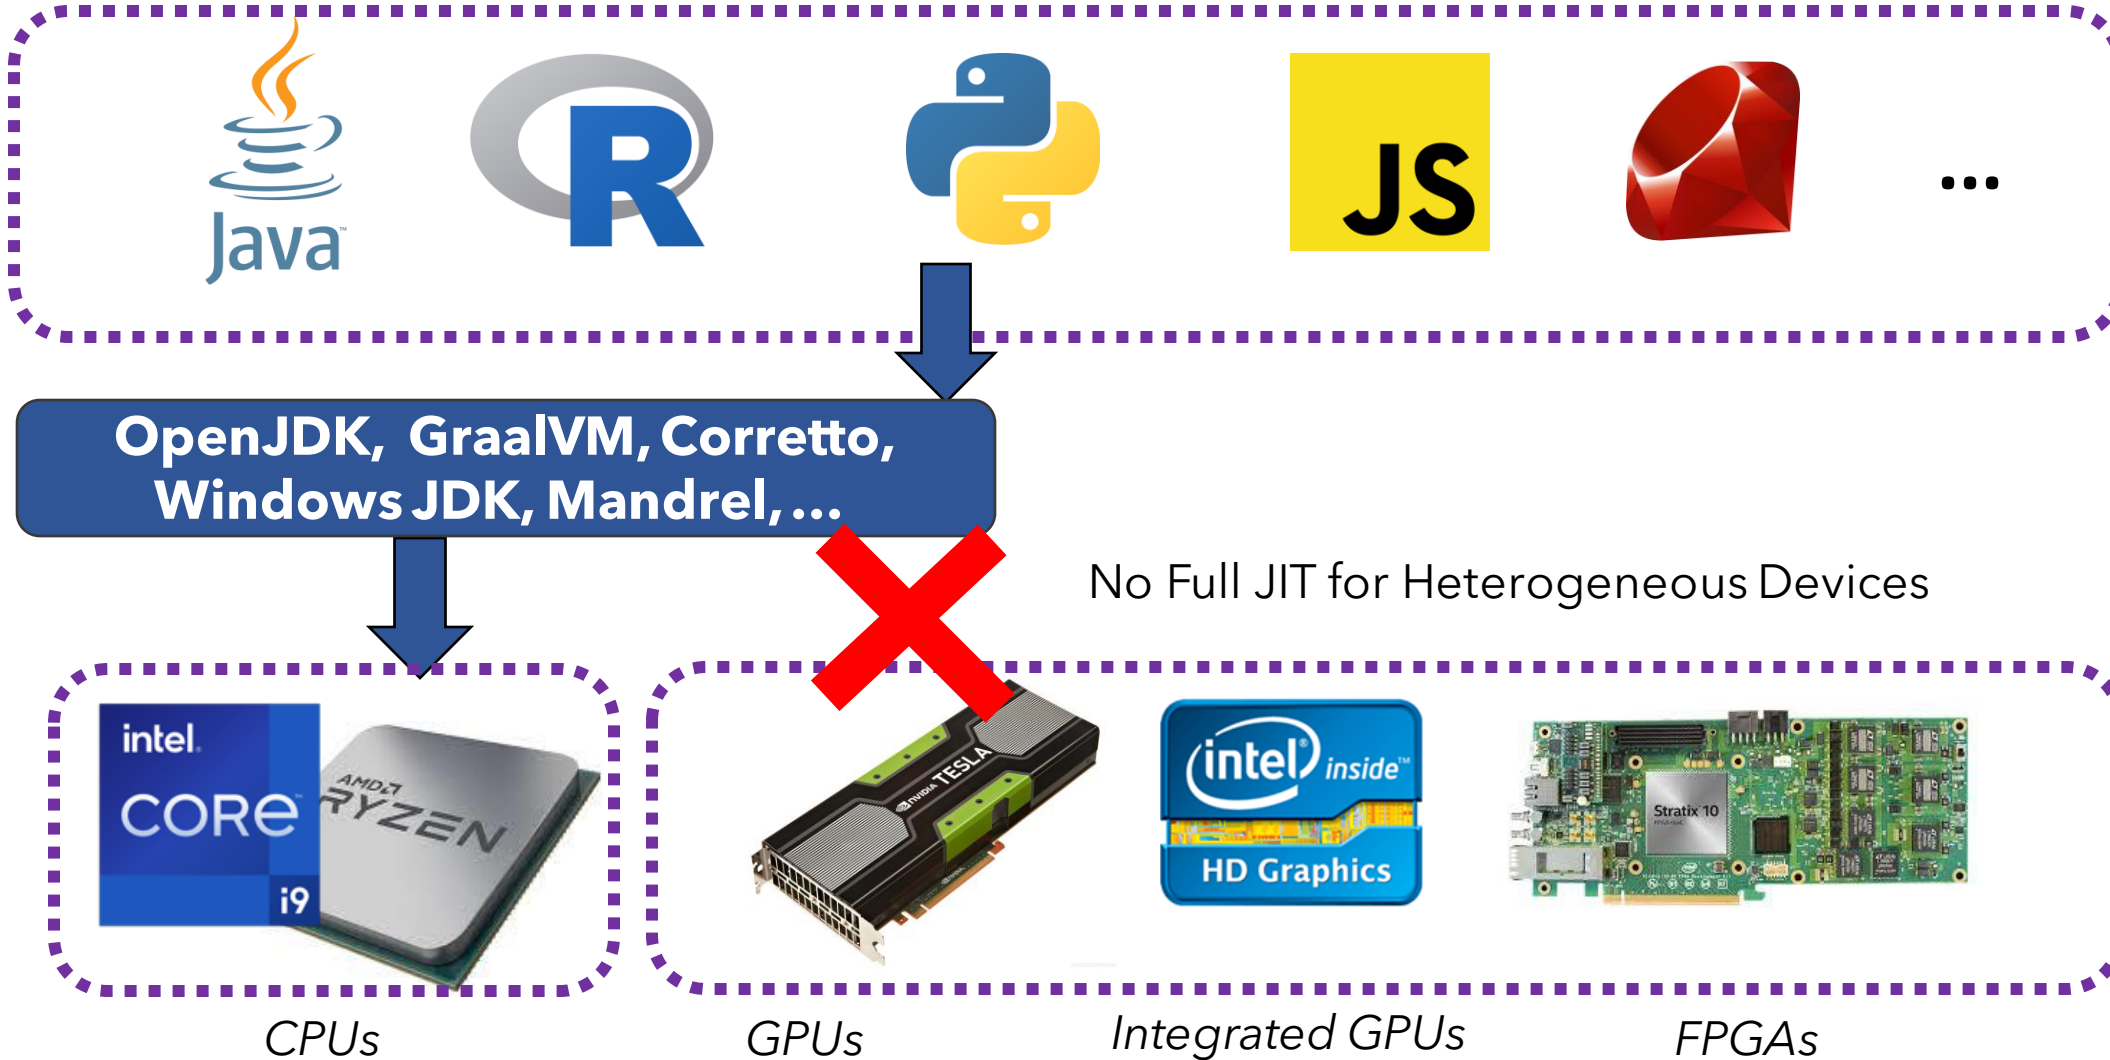
Current Computer Systems



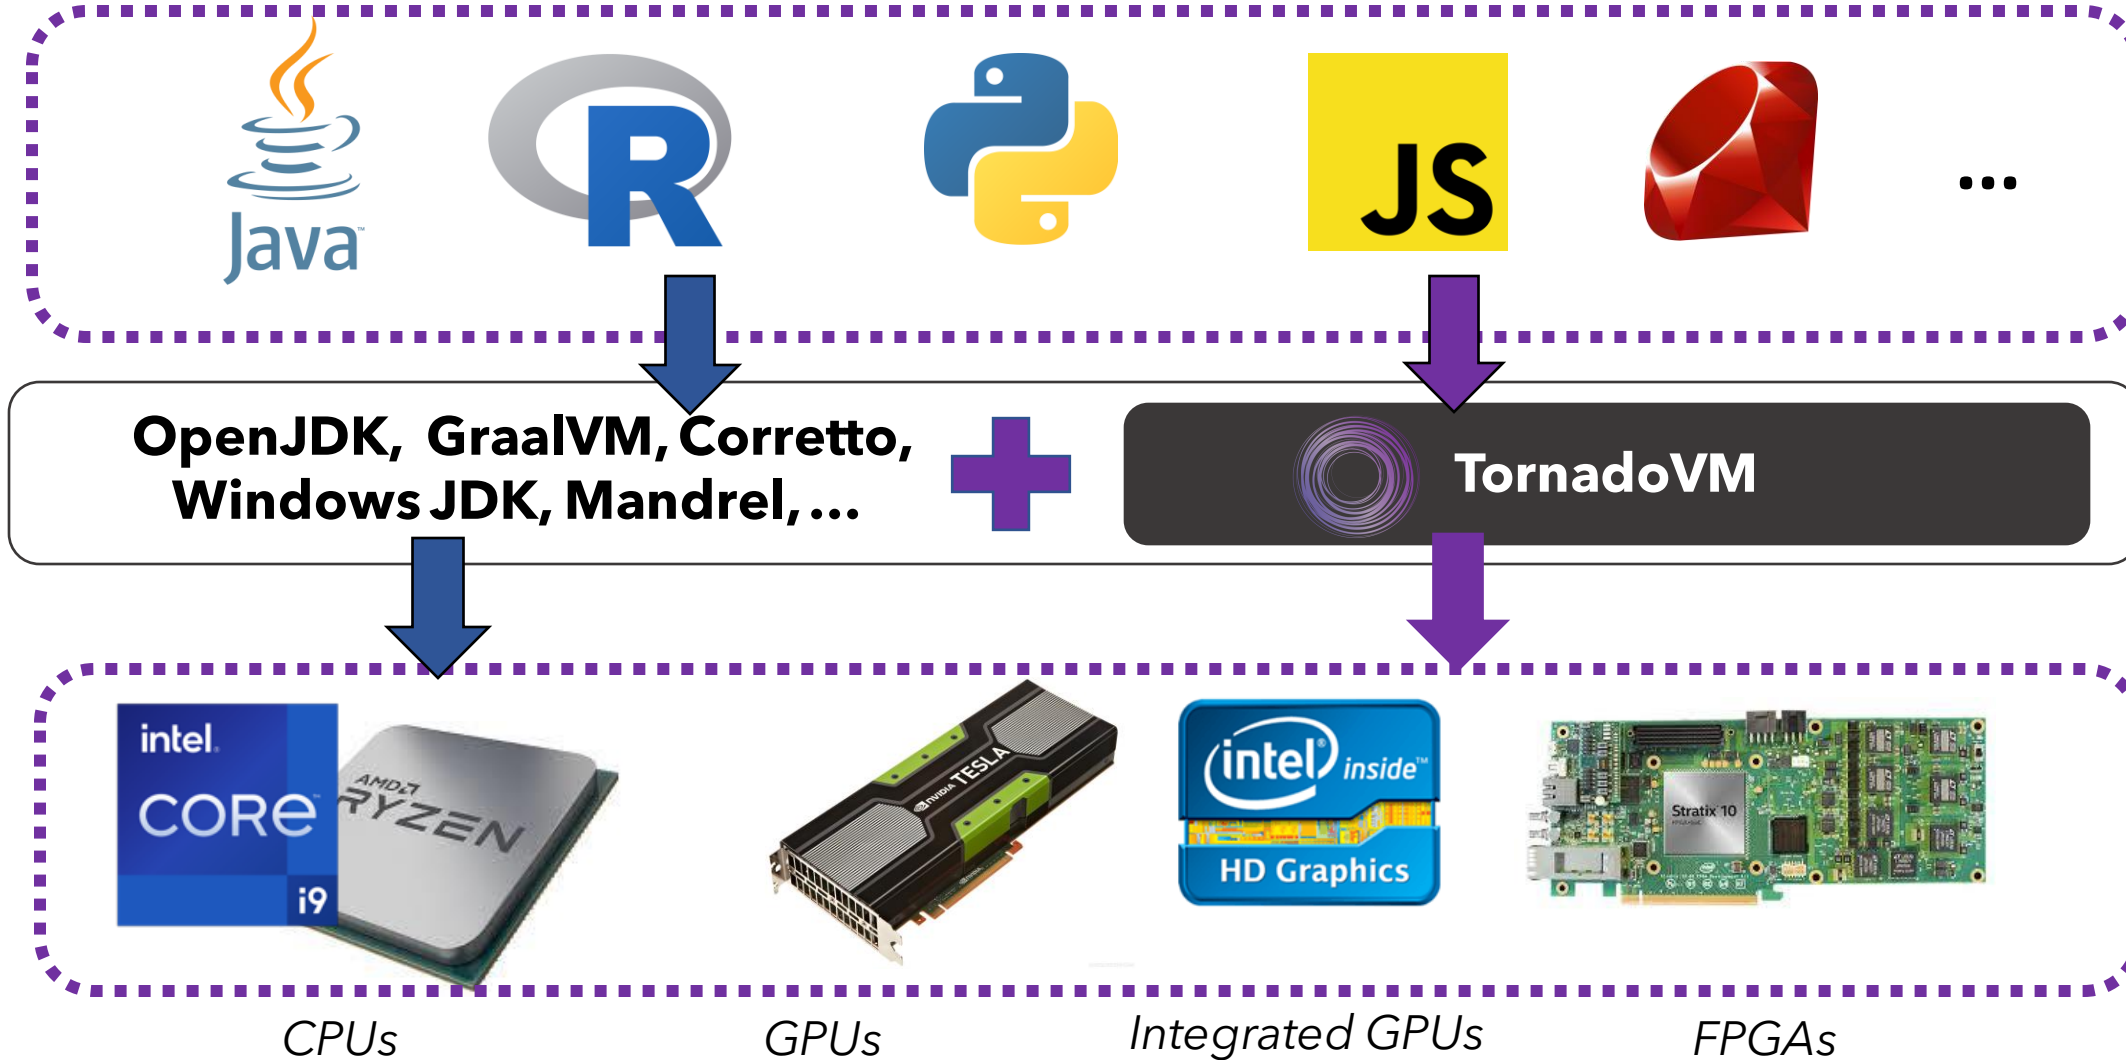
Current Computer Systems



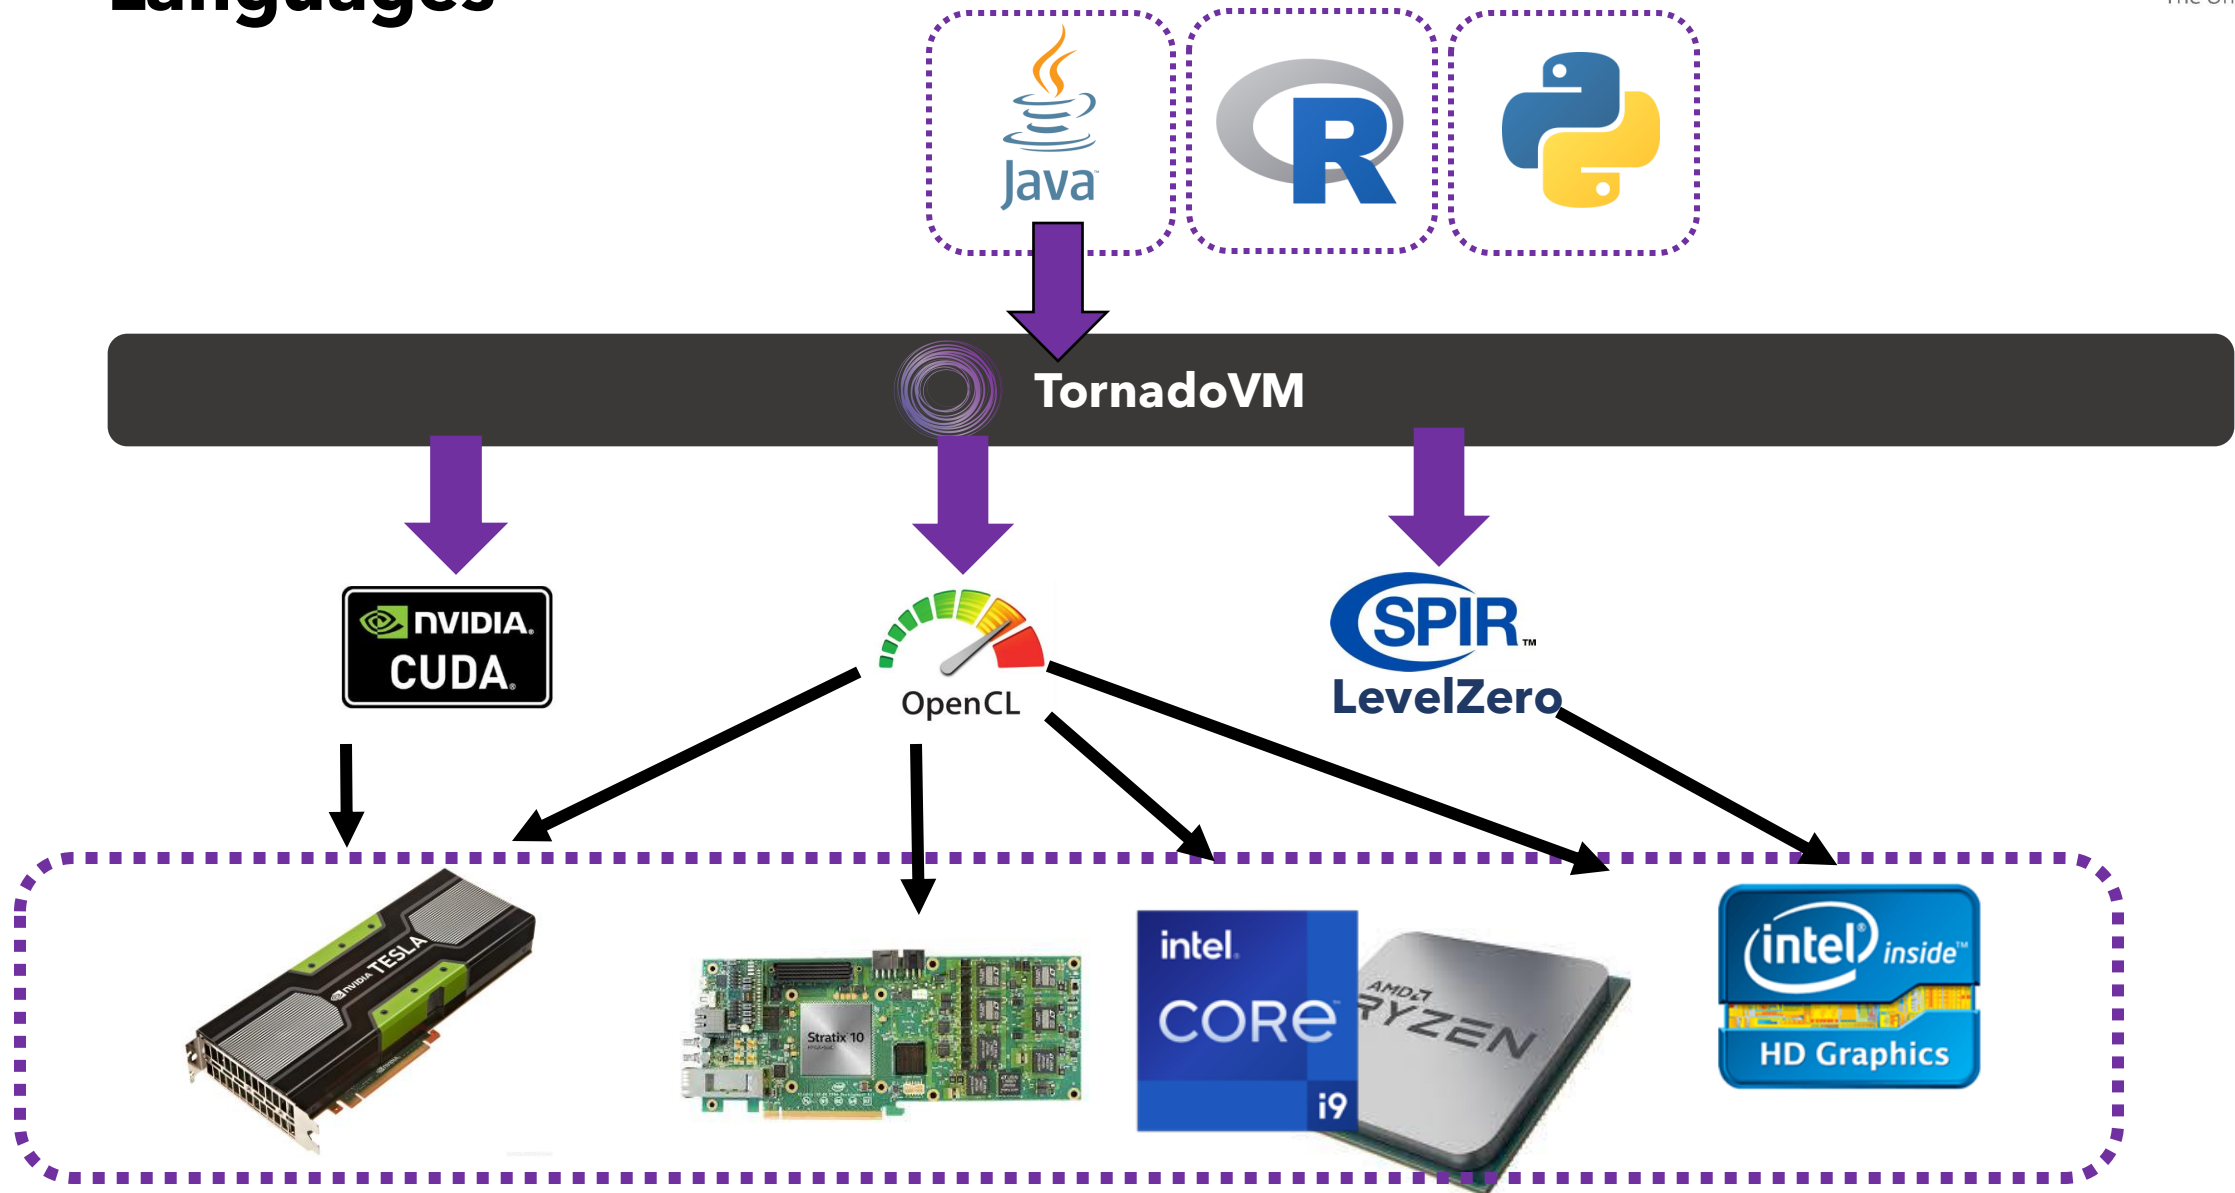
Fast Path to GPUs and FPGAs



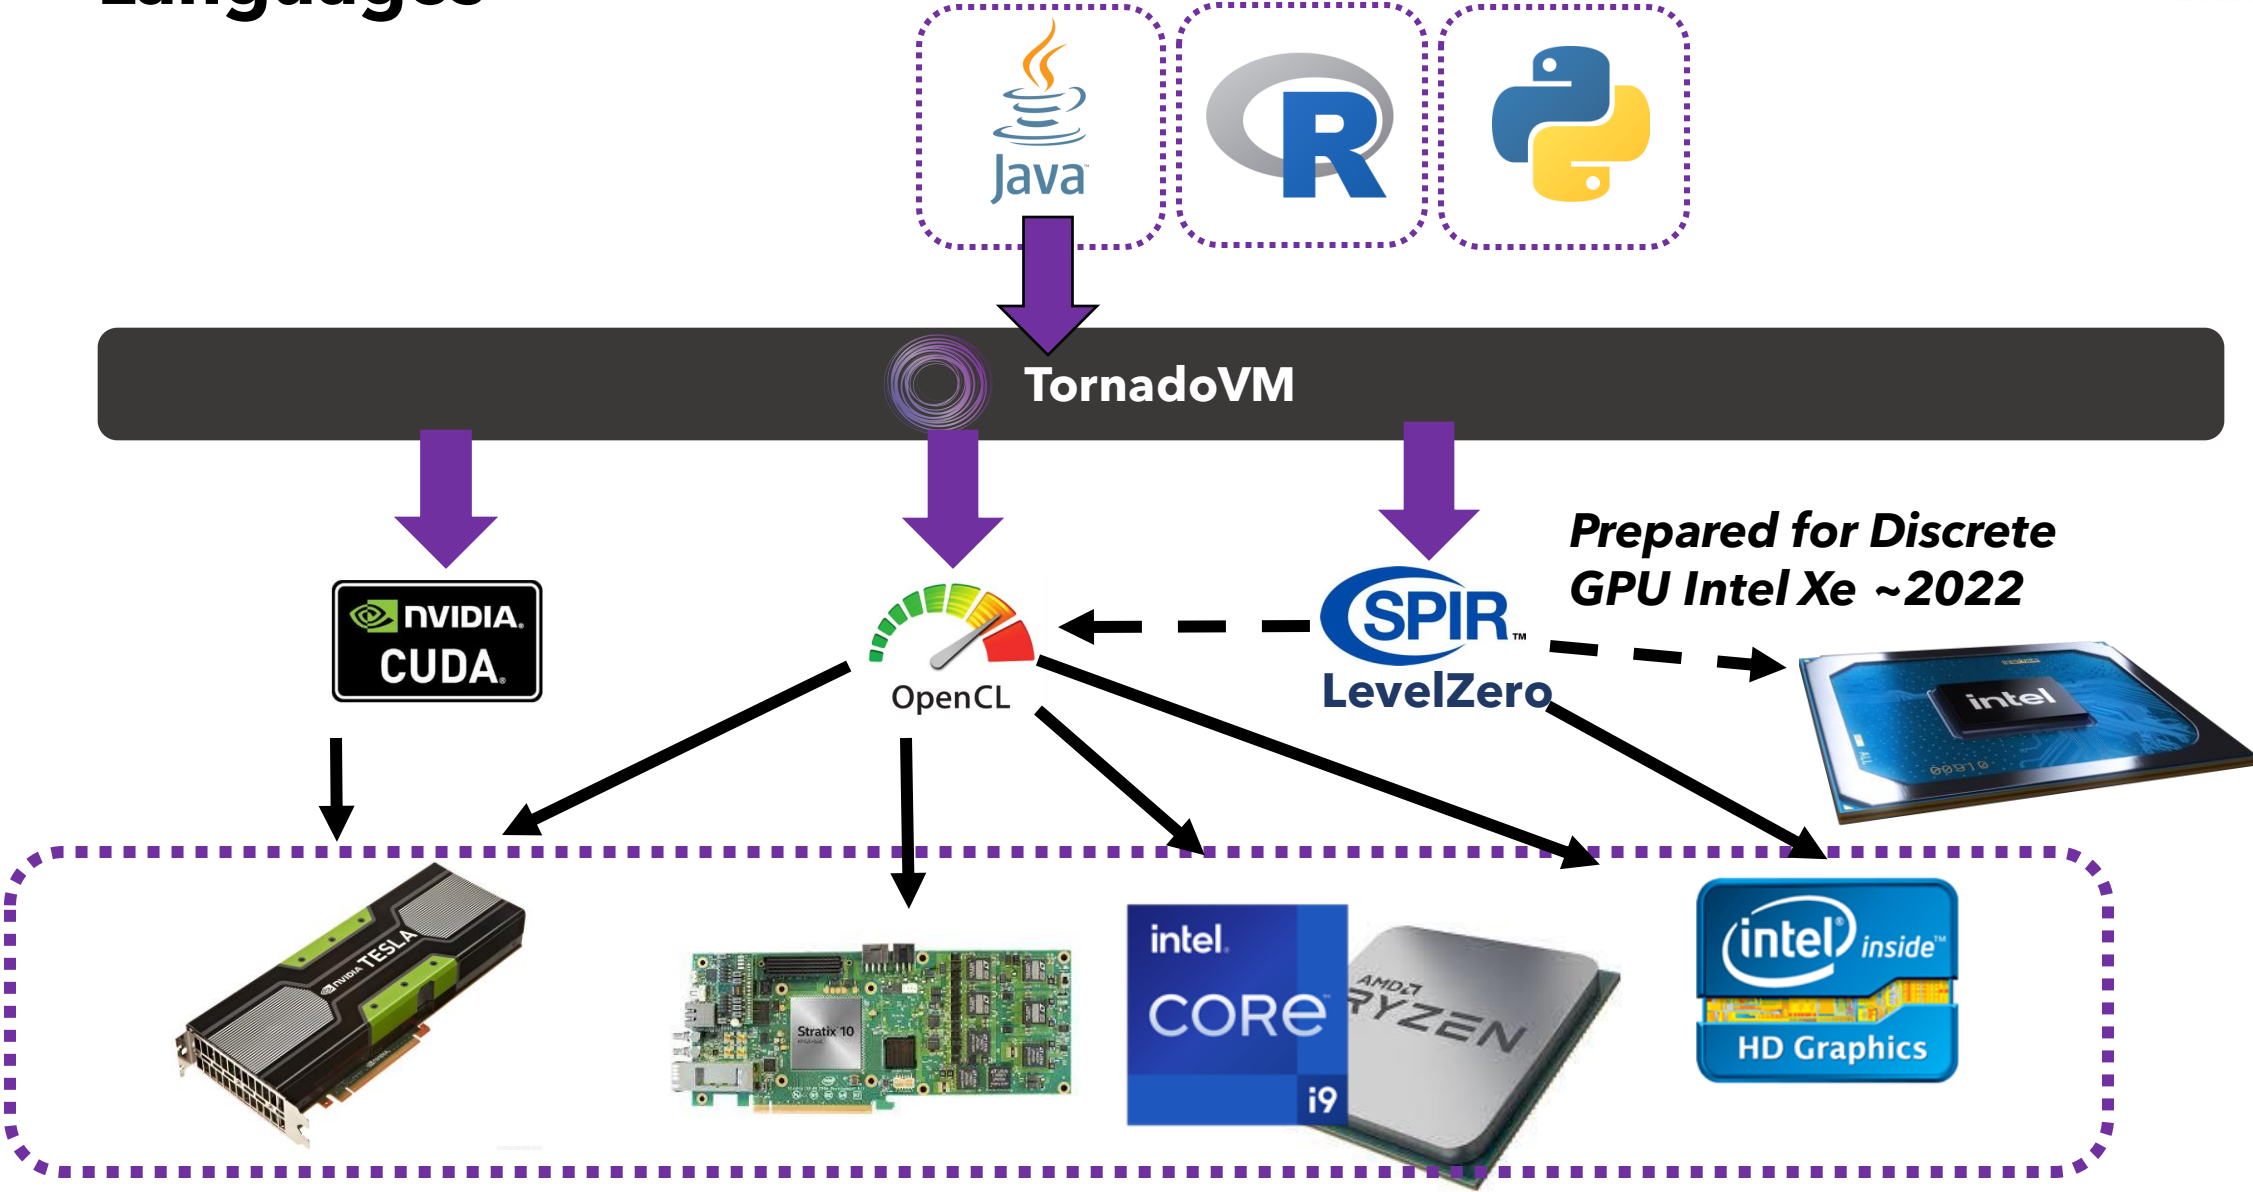
Fast Path to GPUs and FPGAs

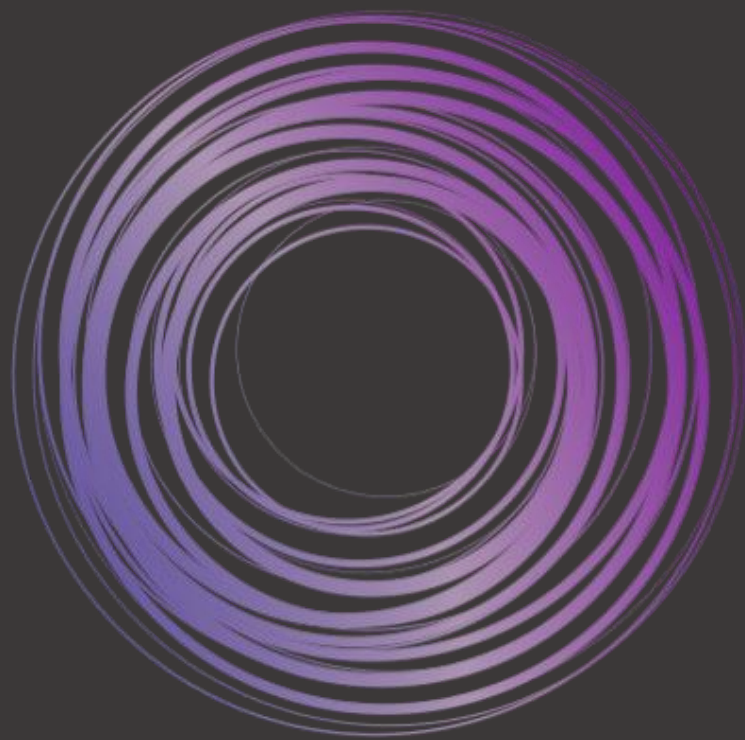


Enabling Acceleration for Managed Runtime Languages



Enabling Acceleration for Managed Runtime Languages





TORNADOVM

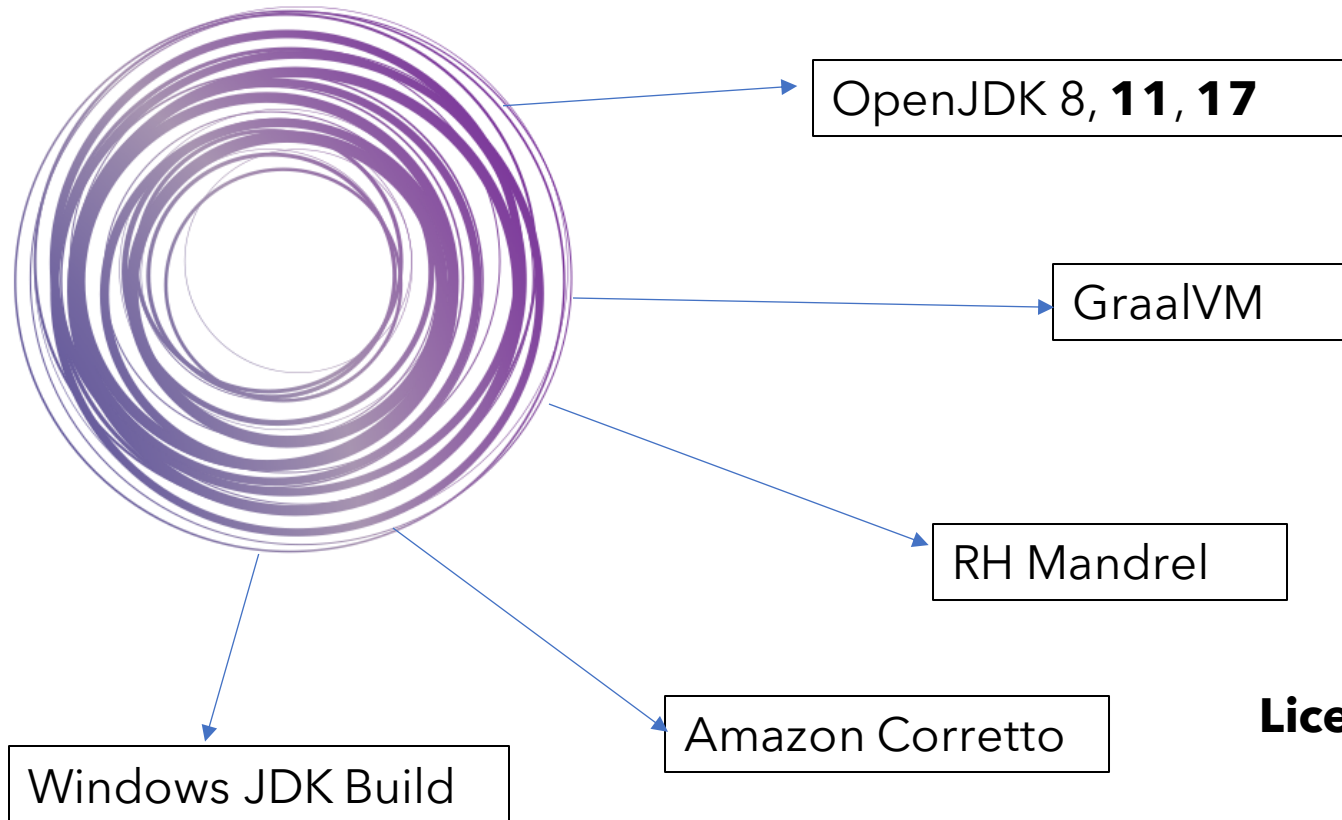
www.tornadovm.org

TornadoVM Overview

www.tornadovm.org



<https://github.com/beehive-lab/TornadoVM>

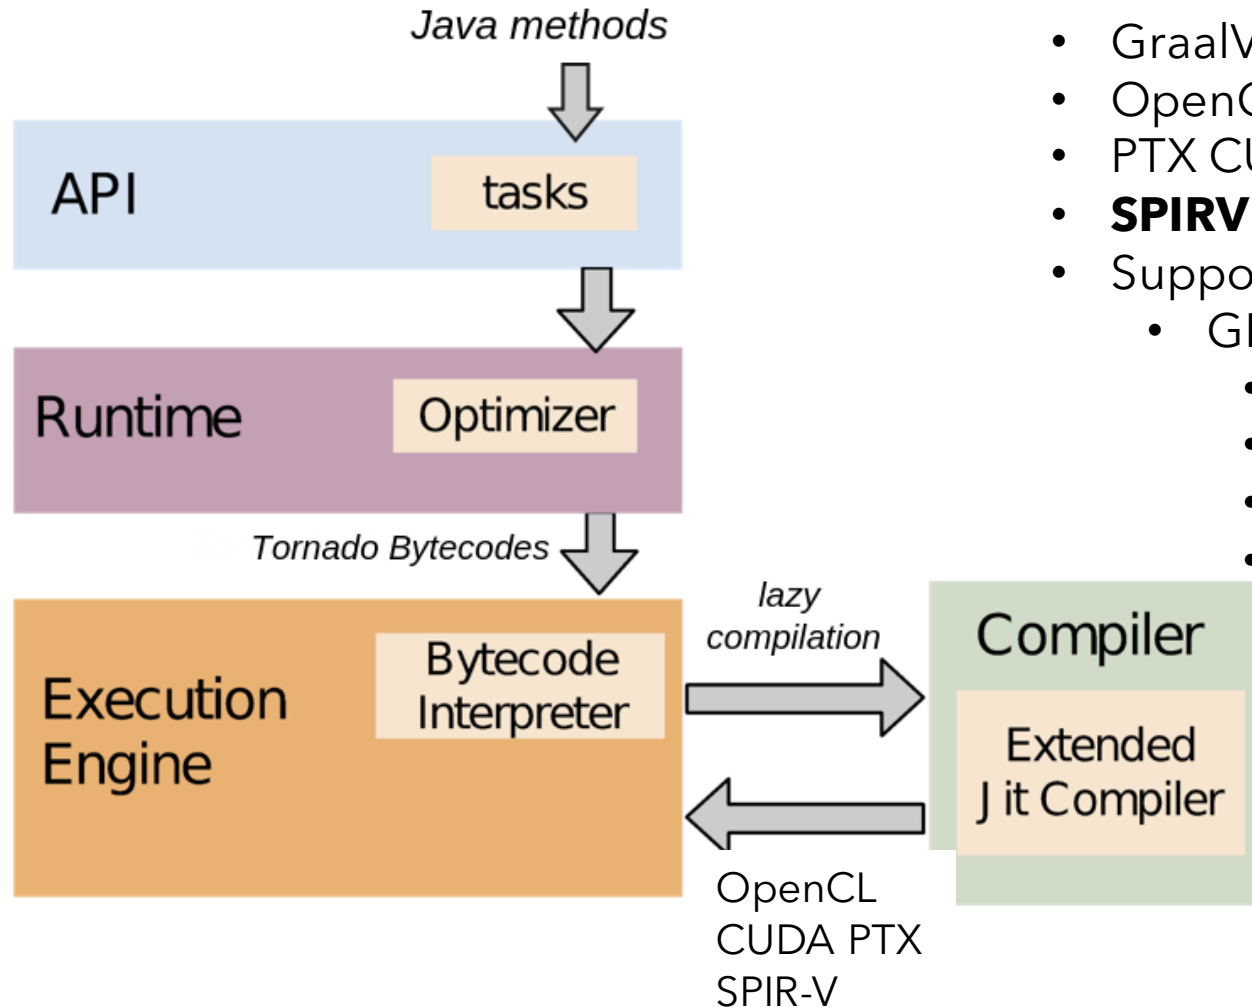


> Open-source Plug-in to multiple JVMs that allows developers to run JVM based programs on heterogeneous hardware

- Perform Automatic Task Migration
- Optimising JIT Compiler for GPUs/FPGAs
- Vendor agnostic, GPUs, CPUs, FPGAs within the same source

License: GPLv2 + CE

TornadoVM Overview



- GraalVM 21.2.0
- OpenCL ≥ 1.2
- PTX CUDA ≥ 10.0
- **SPIRV 1.2 (Prototype)**
- Support for:
 - GPUs:
 - NVIDIA
 - AMD
 - Intel
 - ARM Mali
 - FPGAs:
 - Xilinx
 - Intel
 - CPUs:
 - Intel/AMD

Different Backends



OpenCL

Open Computing **L**anguage

Open Standard - Khronos
Group (non-profit tech
consortium)

Writing programs portable*
across platforms
(source code portability)

**Run on CPUs, GPUs,
DSPs, FPGAs**

Different Backends



OpenCL

Open Computing **L**anguage

Open Standard – Khronos Group (non-profit tech consortium)

Writing programs portable* across platforms (source code portability)

Run on CPUs, GPUs, DSPs, FPGAs



PTX: **P**arallel **T**hread **eX**ecution

ISA used in NVIDIA CUDA's programming model

Developed by NVIDIA

Only for NVIDIA GPUs

Different Backends



OpenCL

Open Computing **L**anguage

Open Standard - Khronos Group (non-profit tech consortium)

Writing programs portable* across platforms (source code portability)

Run on CPUs, GPUs, DSPs, FPGAs



PTX: **P**arallel **T**hread **eX**ecution

ISA used in NVIDIA CUDA's programming model

Developed by NVIDIA

Only for NVIDIA GPUs



Any OpenCL ≥ 2.1 device

Standard Portable
Intermediate **R**epresentation

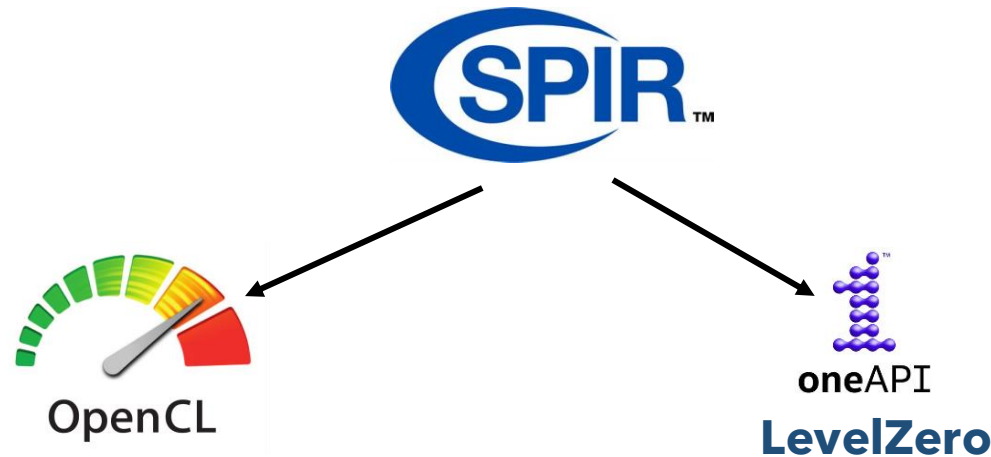
Standard IR binary originally created for OpenCL (≥ 2.1)

Enables distribution of compute binaries for OpenCL

Shared IR with Vulkan for Graphics

And Intel Level Zero?

- It a brand new baremetal API for low-level programming of heterogeneous architectures.
- It is part of the Intel oneAPI ecosystem and can be used as a standalone library.
- Level Zero consumes SPIRV binaries for compute



But ... why Level Zero?

- Clearly influenced by OpenCL
- It can evolve independently
- It supports:
 - Low latency command queues
 - **Virtual functions**
 - **Memory visibility control, caching control**
 - Unified memory
 - Device partitioning
 - Instrumentation and debugging
 - **Control of power management**
 - **Control of frequency**
 - **Hardware diagnostics**
 - ...
- **This level of control is very appealing for system programming, runtime systems and compilers**



It is part of the oneAPI stack and can be accessed as a standalone library:
<https://github.com/oneapi-src/level-zero>

More Info:

- Level Zero Spec: <https://spec.oneapi.io/level-zero/latest/index.html>
- <https://jffumero.github.io/posts/2021/09/introduction-to-level-zero/>

Comparisons

Advantages

Disadvantages



OpenCL

- Easier to write than other alternatives
- Source code portable
- Wide variety of devices

- Performance is not portable (hard to know what the compiler driver will do)



- Highly Tuned for NVIDIA GPUs
- High Performance
- Low-level features

- Only works for NVIDIA GPUs.
- No control over the final compilation (PTX -> bin)



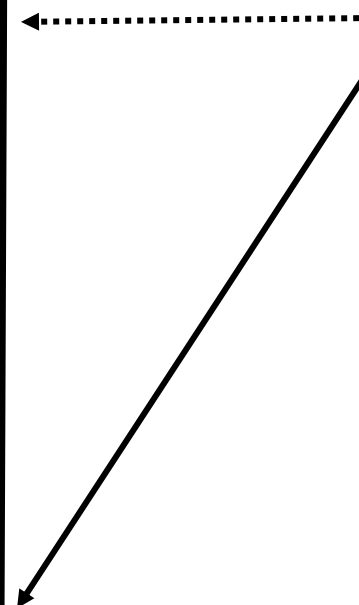
- Very low-level control of the hardware resources
- It dispatches SPIR-V kernels
- Higher control of execution
- Prepared for a wide set of devices

- Exposed to users but designed for coupling with runtimes/compilers (by design)
- New technology

... and now experimenting with

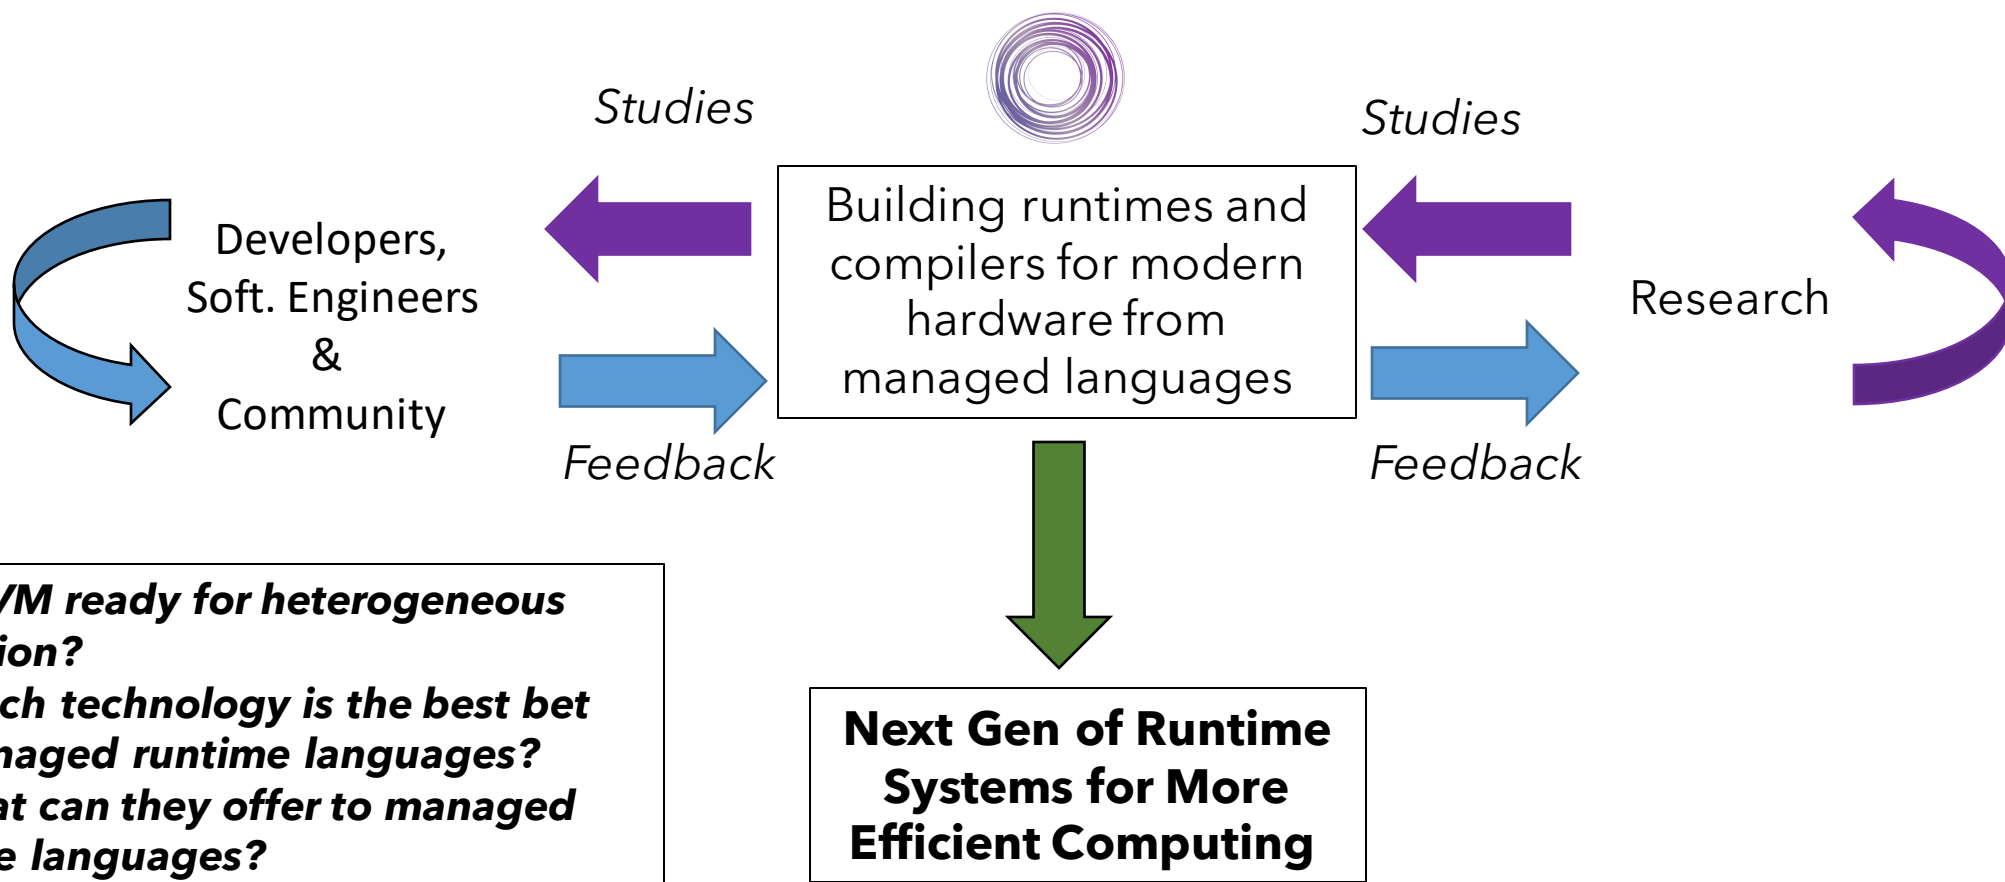


SPIR-V Kernels can be consumed by OpenCL runtime and Intel Level Zero API





So why all of these backends?



Q) Is JVM ready for heterogeneous execution?

Q) Which technology is the best bet for managed runtime languages?

Q) What can they offer to managed runtime languages?



But, how TornadoVM
compiles parallel code
from Java?

Multi-backend JIT Compiler Workflow

Programmer's view



```
public static void saxpy(int[] a, int[] b, int[] c, int alpha) {  
    for (@Parallel int i = 0; i < a.length; i++) {  
        a[i] = alpha * b[i] + c[i];  
    }  
}
```

Multi-backend JIT Compiler Workflow

Programmer's view



```
public static void saxpy(int[] a, int[] b, int[] c, int alpha) {  
    for (@Parallel int i = 0; i < a.length; i++) {  
        a[i] = alpha * b[i] + c[i];  
    }  
}
```

javac



Java Bytecodes

TornadoVM JIT Compiler

Multi-backend JIT Compiler Workflow

Programmer's view



```
public static void saxpy(int[] a, int[] b, int[] c, int alpha) {  
    for (@Parallel int i = 0; i < a.length; i++) {  
        a[i] = alpha * b[i] + c[i];  
    }  
}
```

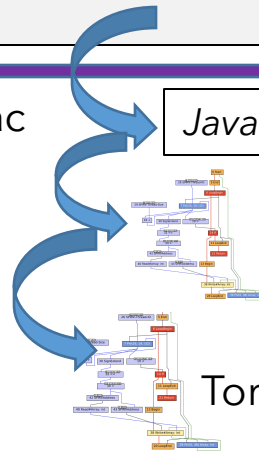
javac

Java Bytecodes

TornadoVM JIT Compiler

Graal IR

TornadoVM Common IR



Multi-backend JIT Compiler Workflow

Programmer's view



```
public static void saxpy(int[] a, int[] b, int[] c, int alpha) {  
    for (@Parallel int i = 0; i < a.length; i++) {  
        a[i] = alpha * b[i] + c[i];  
    }  
}
```

javac

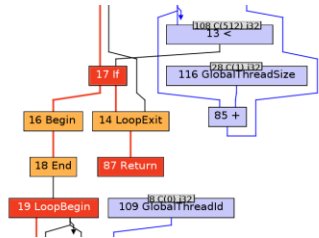
Java Bytecodes

TornadoVM JIT Compiler

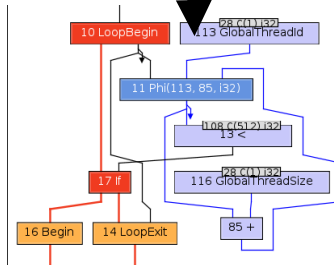
Graal IR

TornadoVM Common IR

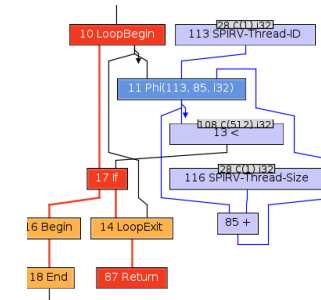
TornadoVM IR for PTX



TornadoVM IR for OpenCL



TornadoVM IR for SPIR-V



Multi-backend JIT Compiler Workflow

Programmer's view



```
public static void saxpy(int[] a, int[] b, int[] c, int alpha) {  
    for (@Parallel int i = 0; i < a.length; i++) {  
        a[i] = alpha * b[i] + c[i];  
    }  
}
```

javac

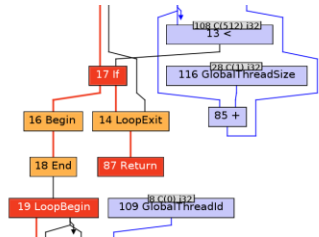
Java Bytecodes

TornadoVM JIT Compiler

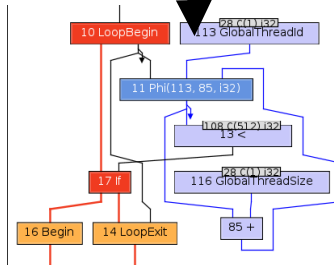
Graal IR

TornadoVM Common IR

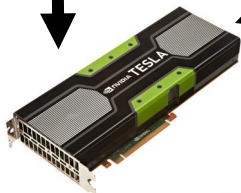
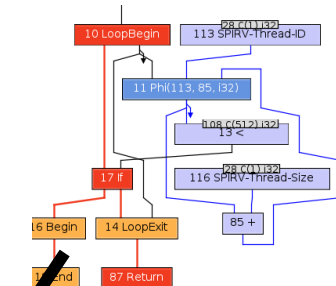
TornadoVM IR for PTX



TornadoVM IR for OpenCL



TornadoVM IR for SPIR-V



Multi-backend JIT Compiler Workflow

Programmer's view



```
public static void saxpy(int[] a, int[] b, int[] c, int alpha) {  
    for (@Parallel int i = 0; i < a.length; i++) {  
        a[i] = alpha * b[i] + c[i];  
    }  
}
```

javac

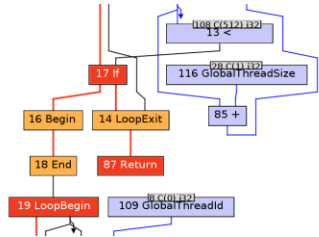
Java Bytecodes

TornadoVM JIT Compiler

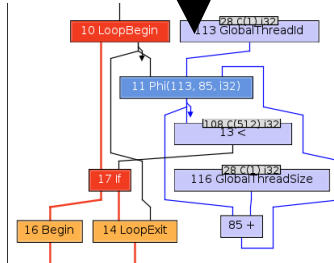
Graal IR

TornadoVM Common IR

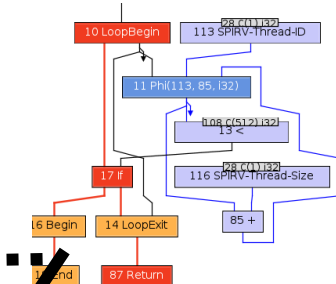
TornadoVM IR for PTX



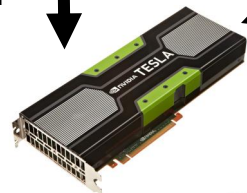
TornadoVM IR for OpenCL



TornadoVM IR for SPIR-V



OpenCL

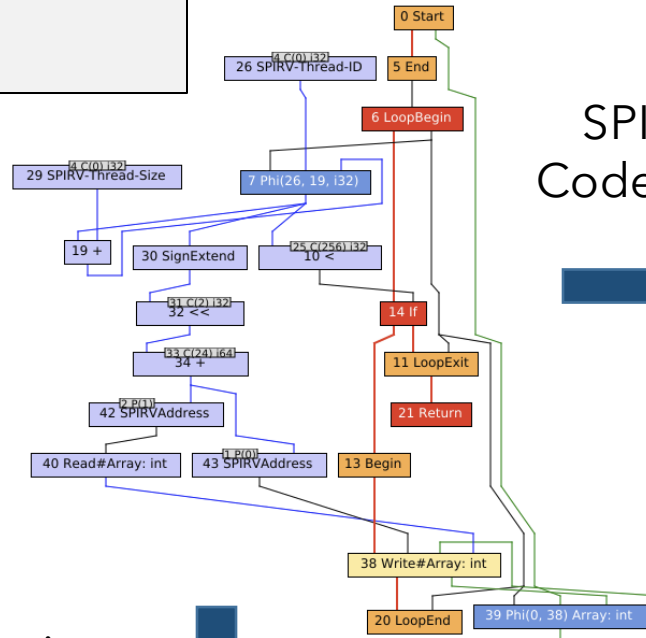
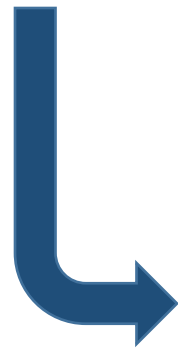


SPIR-V JIT compiler (and runtime) for TornadoVM



```
public static void saxpy(int[] a,
                        int[] b,
                        int[] c,
                        int alpha) {
    for (@Parallel int i = 0; i < a.length; i++) {
        a[i] = alpha * b[i] + c[i];
    }
}
```

Build
Graal/Tornado IR



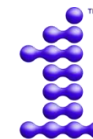
JIT Optimizer
For SPIR-V



SPIR-V
Code-Gen



```
...
%B2 = OpLabel
%77 = OpLoad %uint %spirv_i_4 Aligned 4
%78 = OpSConvert %ulong %77
OpStore %spirv_l_6 %78 Aligned 8
%79 = OpLoad %ulong %spirv_l_6 Aligned 8
%80 = OpShiftLeftLogical %ulong %79 %uint_2
OpStore %spirv_l_7 %80 Aligned 8
%81 = OpLoad %ulong %spirv_l_7 Aligned 8
%82 = OpIAdd %ulong %81 %ulong_24
OpStore %spirv_l_8 %82 Aligned 8
%83 = OpLoad %ulong %spirv_l_1 Aligned 8
%84 = OpLoad %ulong %spirv_l_8 Aligned 8
%85 = OpIAdd %ulong %83 %84
OpStore %spirv_l_9 %85 Aligned 8
%86 = OpLoad %ulong %spirv_l_9 Aligned 8
%87 = OpConvertUToPtr %_ptr_CrossWorkgroup_uint %86
%88 = OpLoad %uint %87 Aligned 4
OpStore %spirv_i_10 %88 Aligned 4
%89 = OpLoad %ulong %spirv_l_2 Aligned 8
%90 = OpLoad %ulong %spirv_l_8 Aligned 8
...
```

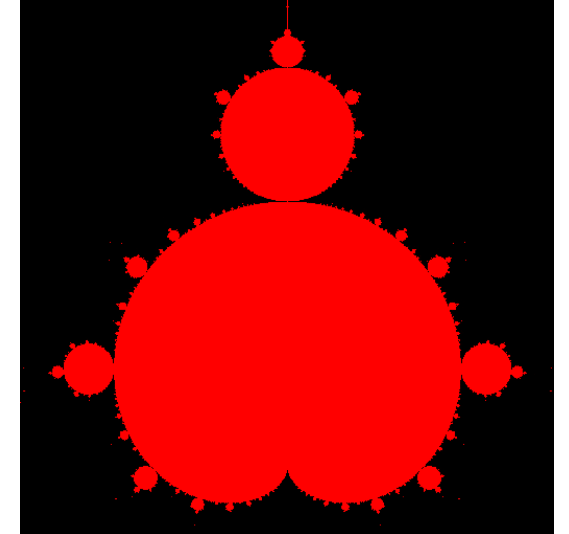


oneAPI

LevelZero-JNI dispatch

Example – Mandelbrot Computation

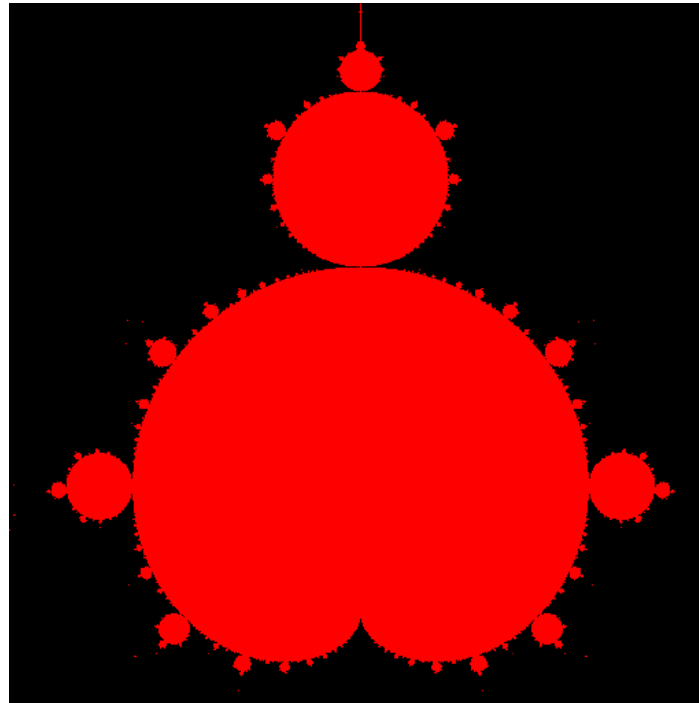
```
public class Mandelbrot {  
  
    static void mandelbrotFractal(final int size, short[] output) {  
        for (@Parallel int i = 0; i < size; i++) {  
            for (@Parallel int j = 0; j < size; j++) {  
                // Mandelbrot computation  
                // Compute the value of each pixel (x, y)  
                // Check example on Github for the specifics  
            }  
        }  
    }  
  
    void createTaskAndRun(int size) {  
        mandelbrotImage = new short[size * size];  
  
        TaskSchedule ts = new TaskSchedule("s0")  
            .task("t0", Mandelbrot::mandelbrotFractal, size, mandelbrotImage)  
            .streamOut(mandelbrotImage);  
  
        ts.execute();  
    }  
}
```



<https://github.com/jjfumero/tornadovm-examples>

Live Demo with the SPIR-V Backend

Mandelbrot
computation

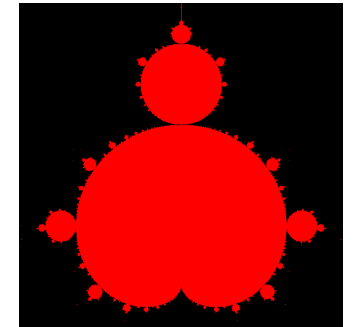
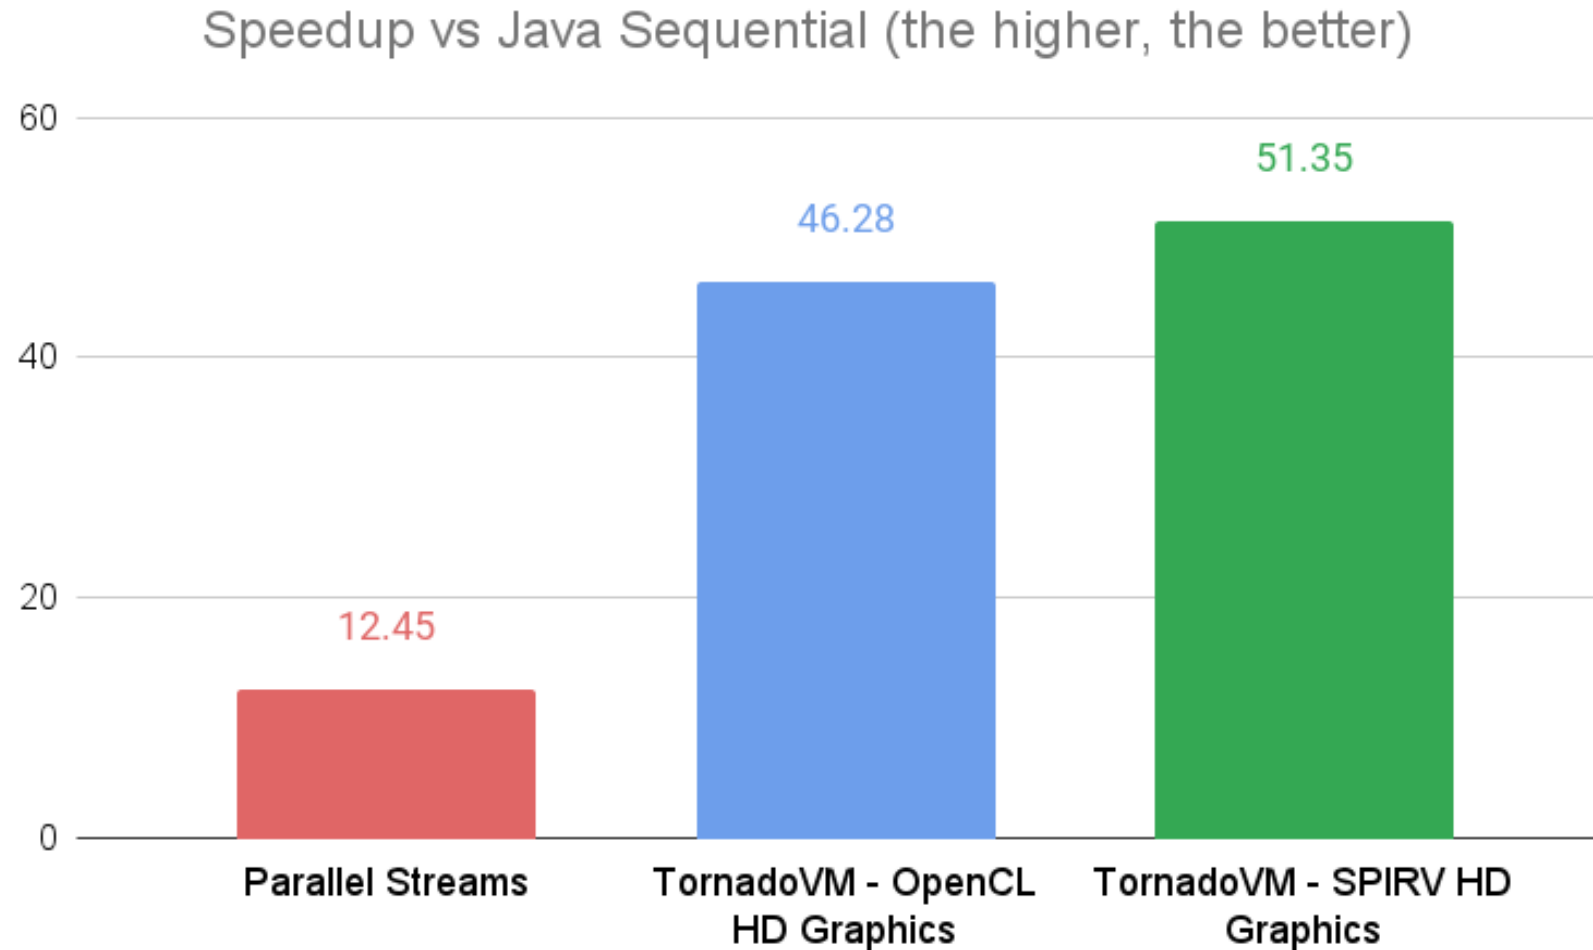


<https://github.com/jjfumero/tornadovm-examples>

Performance



<https://github.com/jjfumero/tornadovm-examples>

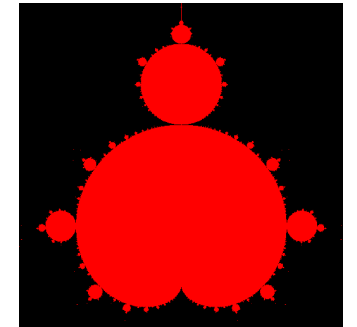
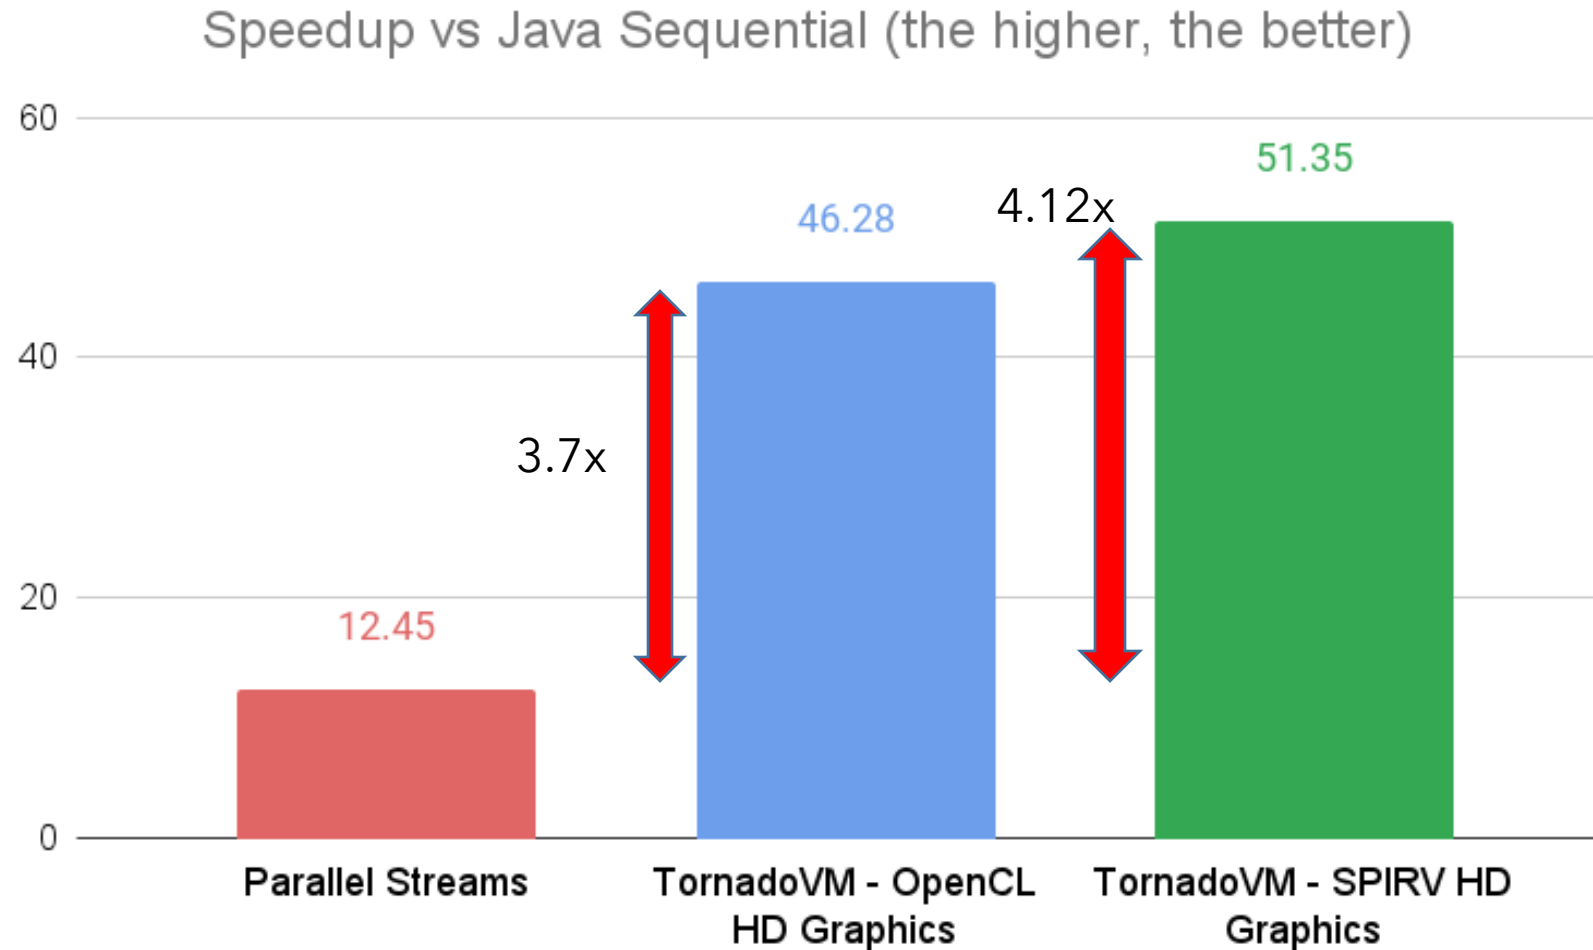


* CPU: Intel(R) Core(TM) i9-10885H
* GPU: Intel HD Graphics
* Java: 1.8.0_302
* LevelZero: 21.38.21026
* TornadoVM: 0.12

Performance



<https://github.com/jjfumero/tornadovm-examples>



* CPU: Intel(R) Core(TM) i9-10885H
* GPU: Intel HD Graphics
* Java: 1.8.0_302
* LevelZero: 21.38.21026
* TornadoVM: 0.12



Profiling

Understanding Performance with the Profiler

```
$ tornado --enableProfiler console Program
```

Understanding Performance with the Profiler

\$ tornado --enableProfiler console Program

*All times are in
nanoseconds*

```
{
  "s0": {
    "TOTAL_KERNEL_TIME": "58591028",
    "COPY_OUT_TIME": "55693",
    "TOTAL_GRAAL_COMPILE_TIME": "179950755",
    "TOTAL_DISPATCH_DATA_TRANSFERS_TIME": "0",
    "TOTAL_TASK_SCHEDULE_TIME": "388705840",
    "COPY_IN_TIME": "50547",
    "TOTAL_BYTE_CODE_GENERATION": "6230794",
    "TOTAL_DRIVER_COMPILE_TIME": "58653972",
    "TOTAL_COPY_IN_SIZE_BYTES": "1048624",
    "TOTAL_COPY_OUT_SIZE_BYTES": "524312",
    "s0.t0": {
      "METHOD": "Mandelbrot.mandelbrotFractal",
      "DEVICE_ID": "0:0",
      "DEVICE": "Intel(R) UHD Graphics [0x9bc4]",
      "TASK_KERNEL_TIME": "58591028",
      "TASK_COMPILE_GRAAL_TIME": "179950755",
      "TASK_COMPILE_DRIVER_TIME": "58653972"
    }
  }
}
```

Task Scheduler's Name

Task-Name

Java Method Compiled

```
TaskSchedule ts = new TaskSchedule("s0")
    .task("t0", Mandelbrot::mandelbrotFractal, size, mandelbrotImage)
    .streamOut(mandelbrotImage);
```

Understanding Performance

```
$ tornado --enableProfiler console Program
```

```
{
  "s0": {
    "TOTAL_KERNEL_TIME": "58591028",
    "COPY_OUT_TIME": "55693",
    "TOTAL_GRAAL_COMPILE_TIME": "179950755",
    "TOTAL_DISPATCH_DATA_TRANSFERS_TIME": "0",
    "TOTAL_TASK_SCHEDULE_TIME": "388705840",
    "COPY_IN_TIME": "50547",
    "TOTAL_BYTE_CODE_GENERATION": "6230794",
    "TOTAL_DRIVER_COMPILE_TIME": "58653972",
    "TOTAL_COPY_IN_SIZE_BYTES": "1048624",
    "TOTAL_COPY_OUT_SIZE_BYTES": "524312",
    "s0.t0": {
      "METHOD": "Mandelbrot.mandelbrotFractal",
      "DEVICE_ID": "0:0",
      "DEVICE": "Intel(R) UHD Graphics [0x9bc4]",
      "TASK_KERNEL_TIME": "58591028",
      "TASK_COMPILE_GRAAL_TIME": "179950755",
      "TASK_COMPILE_DRIVER_TIME": "58653972"
    }
  }
}
```

Total Time including
data transfers,
execution and
TornadoVM runtime to
dispatch the kernels.

Understanding Performance

```
$ tornado --enableProfiler console Program
```

```
{  
  "s0": {  
    "TOTAL_KERNEL_TIME": "58591028",  
    "COPY_OUT_TIME": "55693",  
    "TOTAL_GRAAL_COMPILE_TIME": "179950755",  
    "TOTAL_DISPATCH_DATA_TRANSFERS_TIME": "0",  
    "TOTAL_TASK_SCHEDULE_TIME": "388705840",  
    "COPY_IN_TIME": "50547",  
    "TOTAL_BYTE_CODE_GENERATION": "6230794",  
    "TOTAL_DRIVER_COMPILE_TIME": "58653972",  
    "TOTAL_COPY_IN_SIZE_BYTES": "1048624",  
    "TOTAL_COPY_OUT_SIZE_BYTES": "524312",  
    "s0.t0": {  
      "METHOD": "Mandelbrot.mandelbrotFractal",  
      "DEVICE_ID": "0:0",  
      "DEVICE": "Intel(R) UHD Graphics [0x9bc4]",  
      "TASK_KERNEL_TIME": "58591028",  
      "TASK_COMPILE_GRAAL_TIME": "179950755",  
      "TASK_COMPILE_DRIVER_TIME": "58653972"  
    }  
  }  
}
```

Compilation with Graal
+ code generation

(Java byte code ->
Graal IR -> Tornado IR ->
optimizations + code
generation)

Internal Byte-Code Generation

**Driver JIT compiler
(e.g., SPIR-V -> final
GPU binary)**

Understanding Performance

\$ tornado --enableProfiler console Program

```
{
  "s0": {
    "TOTAL_KERNEL_TIME": "58591028",
    "COPY_OUT_TIME": "55693",
    "TOTAL_GRAAL_COMPILE_TIME": "179950755",
    "TOTAL_DISPATCH_DATA_TRANSFERS_TIME": "0",
    "TOTAL_TASK_SCHEDULE_TIME": "388705840",
    "COPY_IN_TIME": "50547",
    "TOTAL_BYTE_CODE_GENERATION": "6230794",
    "TOTAL_DRIVER_COMPILE_TIME": "58653972",
    "TOTAL_COPY_IN_SIZE_BYTES": "1048624",
    "TOTAL_COPY_OUT_SIZE_BYTES": "524312",
    "s0.t0": {
      "METHOD": "Mandelbrot.mandelbrotFractal",
      "DEVICE_ID": "0:0",
      "DEVICE": "Intel(R) UHD Graphics [0x9bc4]",
      "TASK_KERNEL_TIME": "58591028",
      "TASK_COMPILE_GRAAL_TIME": "179950755",
      "TASK_COMPILE_DRIVER_TIME": "58653972"
    }
  }
}
```

Total Kernel Time

Total Copy Out (Device -> Java Heap)

Total Copy In (Java Heap -> Device)

Kernel Time For each task

Understanding Performance

\$ tornado --enableProfiler console Program

```
{
  "s0": {
    "TOTAL_KERNEL_TIME": "58591028",
    "COPY_OUT_TIME": "55693",
    "TOTAL_GRAAL_COMPILE_TIME": "179950755",
    "TOTAL_DISPATCH_DATA_TRANSFERS_TIME": "0",
    "TOTAL_TASK_SCHEDULE_TIME": "388705840",
    "COPY_IN_TIME": "50547",
    "TOTAL_BYTE_CODE_GENERATION": "6230794",
    "TOTAL_DRIVER_COMPILE_TIME": "58653972",
    "TOTAL_COPY_IN_SIZE_BYTES": "1048624",
    "TOTAL_COPY_OUT_SIZE_BYTES": "524312",
    "s0.t0": {
      "METHOD": "Mandelbrot.mandelbrotFractal",
      "DEVICE_ID": "0:0",
      "DEVICE": "Intel(R) UHD Graphics [0x9bc4]",
      "TASK_KERNEL_TIME": "58591028",
      "TASK_COMPILE_GRAAL_TIME": "179950755",
      "TASK_COMPILE_DRIVER_TIME": "58653972"
    }
  }
}
```

Ideally, most of the time should be spent in Kernel Execution

- * Take advantage of the device's computing power
- * Keep transfers to minimum

If the application has a lot of data transfers, it is worth trying with shared memory devices (e.g., Integrated GPU) --> In TornadoVM this is not currently handled (WIP)



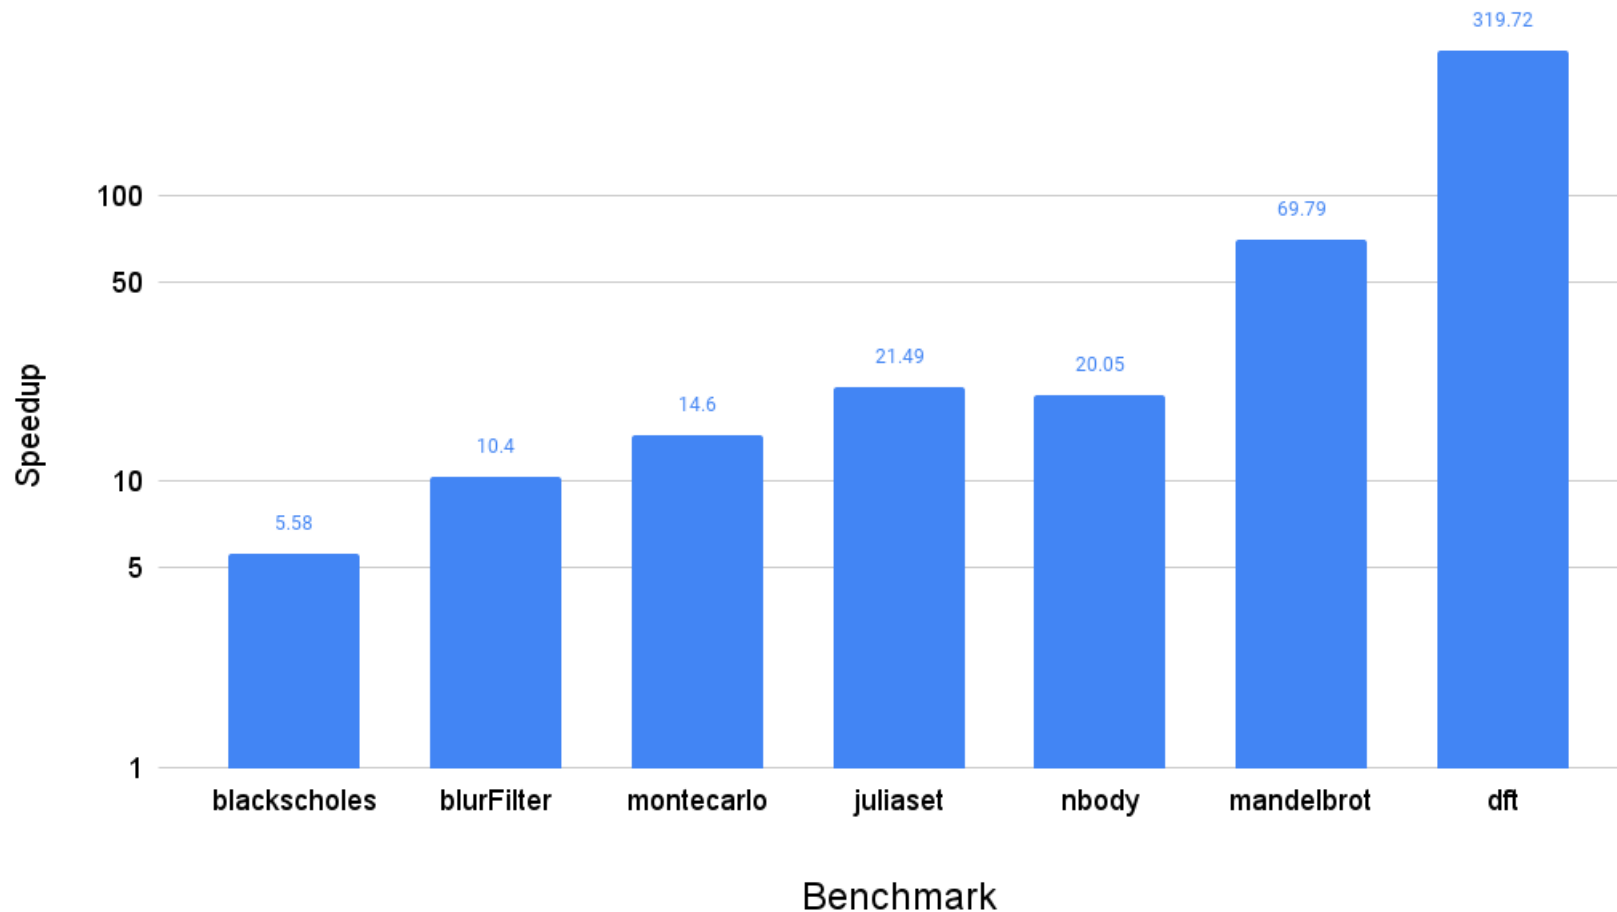
<https://github.com/jjfumero/tornadovm-examples>



Performance

Performance - JMH Benchmarking

TornadoVM SPIR-V Speedup vs OpenJDK 8 C2 compiler (the higher, the better)



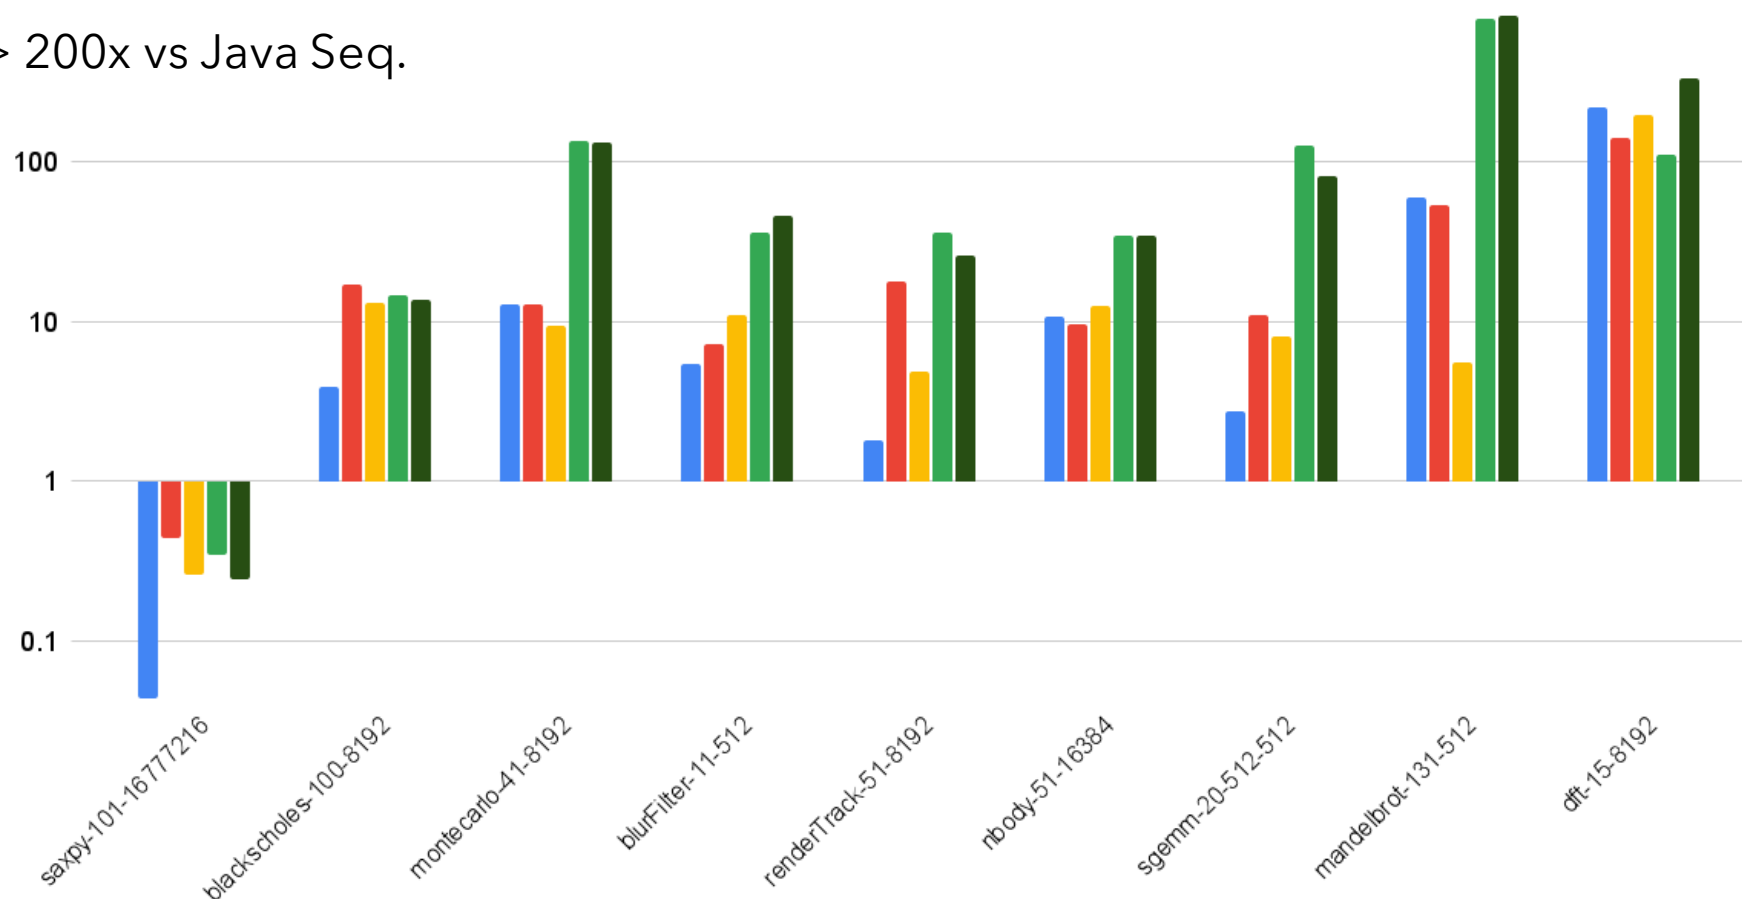
- Intel HD Graphics 630 (Intel i7-7700HQ)
- Running for ~4h - Report from JMH
- Up to 320x performance
- Level-Zero: 21.38.21026
- SPIRV-1.2
- TornadoVM v0.12

Performance vs OpenCL Backend

Peak Speedup of each the SPIR-V and OpenCL Backends vs Java Sequential

■ SPIR-V HD Graphics ■ OCL HD Graphics ■ OCL Intel CPU i9 ■ OCL GTX 2060 ■ PTX GTX 2060

> 200x vs Java Seq.



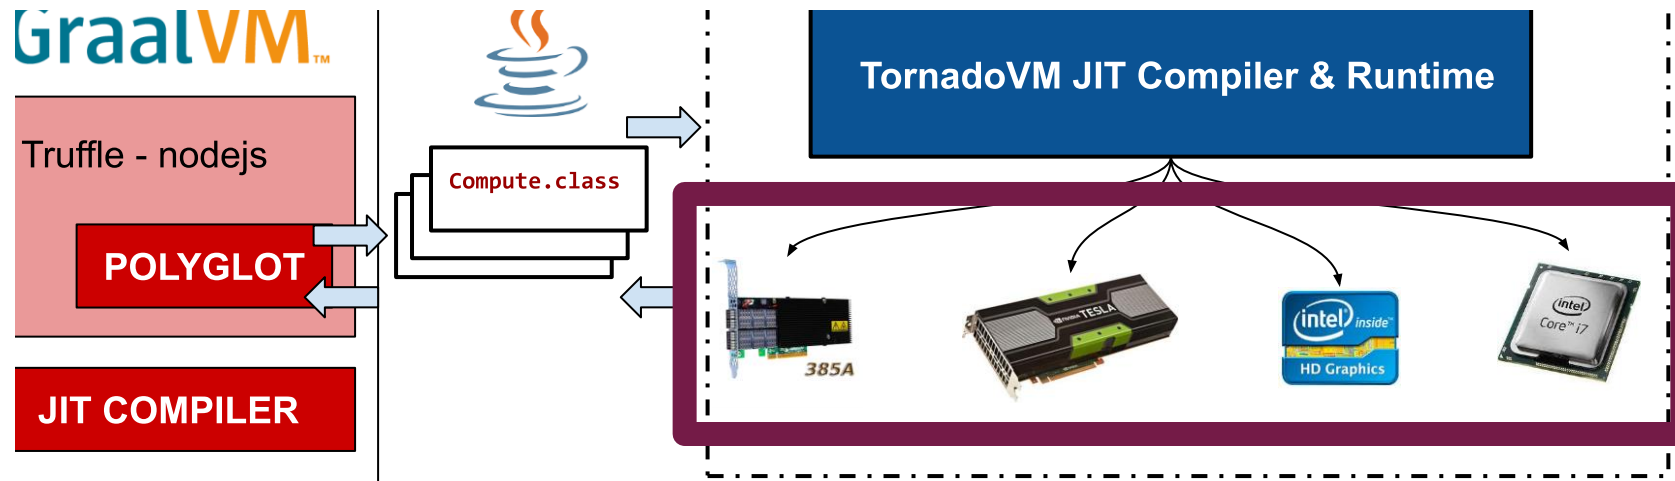
SPIR-V Backend and Level Zero is competitive with the PTX and OpenCL backends

- Intel HD Graphics 630 (Intel i9-10885H)
- GTX 2060
- Level-Zero: 21.38.21026



Running other Programming Languages?

Support for other dynamic languages



Support for other dynamic languages

```
$ ./graalvm-ce-java8-21.2.0/bin/graalpython [params] mxmWithTornadoVM.py
```

Running with tornadoVM

Task info: s0.t0

Backend : SPIRV

Device : SPIRV LevelZero - Intel(R) UHD Graphics [0x9bc4] GPU

Dims : 2

Global work offset: [0, 0]

Global work size : [256, 256]

Local work size : [256, 1, 1]

Number of workgroups : [1, 256]

```
#!/usr/bin/python
print("Running with tornadoVM")
import java
myclass = java.type('MyCompute')
output = myclass.compute()
```



<https://www.tornadovm.org/resources>



Standalone library for low- level GPU programming

LevelZero JNI Library for TornadoVM

- Level Zero Bridge for TornadoVM
 - Since LevelZero is not stable yet, we tried to do a 1-1 mapping between the Java API and C-LevelZero.
 - Easy for us to adapt to new changes
 - In near future, we will leverage this API

```
// Create the Level Zero Driver
LevelZeroDriver driver = new LevelZeroDriver();
int result =
driver.zeInit(ZeInitFlag.ZE_INIT_FLAG_GPU_ONLY);
LevelZeroUtils.errorLog("zeInit", result);

// Get the number of drivers
int[] numDrivers = new int[1];
result = driver.zeDriverGet(numDrivers, null);
LevelZeroUtils.errorLog("zeDriverGet", result);
```

The Intel Level Zero Spec: <https://spec.oneapi.io/level-zero/latest/index.html>

LevelZero JNI Library for TornadoVM

- Level Zero Bridge for TornadoVM
 - Since LevelZero is not stable, we tried to do a 1-1 mapping between the Java API and C-LevelZero.
 - Easy for us to adapt to new changes
 - In near future, we will leverage this API

```
// Create the Level Zero Driver
LevelZeroDriver driver = new LevelZeroDriver();
int result =
driver.zeInit(ZeInitFlag.ZE_INIT_FLAG_GPU_ONLY);
LevelZeroUtils.errorLog("zeInit", result);

// Get the number of drivers
int[] numDrivers = new int[1];
result = driver.zeDriverGet(numDrivers, null);
LevelZeroUtils.errorLog("zeDriverGet", result);
```

```
// Create buffer
LevelZeroBufferInteger bufferA = new LevelZeroBufferInteger();
// Declare buffer as a shared memory
result = context.zeMemAllocShared(context.getContextHandle(),
                                deviceMemAllocDesc,
                                hostMemAllocDesc,
                                bufferSize,
                                1,
                                device.getDeviceHandlerPtr(),
                                bufferA);
LevelZeroUtils.errorLog("zeMemAllocShared", result);

// Level Zero Context
// Device descriptor
// Host Descriptor
// Buffer size in Bytes
// Alignment
// Device pointer
// Buffer to use
```

LevelZero JNI Library for TornadoVM

- This library dispatches SPIR-V kernels
- It does not support full LevelZero, just what we need for TornadoVM, although it could be easily extensible
- It is open source under:
 - **MIT License**



<https://github.com/beehive-lab/levelzero-jni/>

The screenshot shows the GitHub repository page for `beehive-lab / levelzero-jni`. The repository is public and has 1 star and 0 forks. The main branch is `master`. The repository contains a table of files and folders, including `levelZeroLib`, `scripts`, `src`, `.gitignore`, `LICENSE`, `README.md`, `copy_data.cl`, and `pom.xml`. The `README.md` file is selected, showing the title `LevelZero JNI` and the description: "Baremetal GPU and FPGA programming for Java using the LevelZero API." The repository also has a "Languages" section showing the distribution of code by language: Java (66.2%), C++ (24.7%), C (8.6%), and Other (0.5%).

File/Folder	Description	Time
levelZeroLib	[LevelZero] Driver UUID stored	last month
scripts	Run script for timing data transfers added	3 months ago
src	VPU device type added	8 days ago
.gitignore	[JNI][LO] Initial prototype imported from Inter...	9 months ago
LICENSE	[LICENSE] MIT license added	3 months ago
README.md	LevelZero version supported documented in...	20 days ago
copy_data.cl	[JNI][LO] Initial prototype imported from Inter...	9 months ago
pom.xml	[JNI][LO] Initial prototype imported from Inter...	9 months ago

LevelZero JNI

Baremetal GPU and FPGA programming for Java using the [LevelZero API](#).

This project is a Java Native Interface (JNI) binding for Intel's Level Zero. This library is as designed to be as closed as possible to the LevelZero API for C++.

Subset of LevelZero 1.2.2 supported (LevelZero Feb 2021 version)

Compilation & Configuration of the JNI Level-Zero API

Languages

Language	Percentage
Java	66.2%
C++	24.7%
C	8.6%
Other	0.5%



Open-up the code the *SPIRV* generator

SPIR-V Beehive Toolkit for code-gen within TornadoVM

- Java Library for SPIR-V code generation
- Works totally independent from TornadoVM
- It implements **full SPIR-V 1.2**
- Open-source it as a stand-alone library

SPIR-V Beehive Toolkit for code-gen within TornadoVM

- Java Library for SPIR-V code generation
- Works totally independent from TornadoVM
- It implements **full SPIR-V 1.2**
- Open-source it as a stand-alone library

```
// SPIR-V Header
asm.module = new SPIRVModule(
    new SPIRVHeader(
        1, // Major Version
        2, // Minor Version
        33, // ID-Generator (new one)
        0, // Bounds
        0)); // Schema
```

SPIR-V Beehive Toolkit for code-gen within TornadoVM

- Java Library for SPIR-V code generation
- Works totally independent from TornadoVM
- It implements **full SPIR-V 1.2**
- Open-source it as a stand-alone library

```
// SPIR-V Header
asm.module = new SPIRVModule(
    new SPIRVHeader(
        1, // Major Version
        2, // Minor Version
        33, // ID-Generator (new one)
        0, // Bounds
        0)); // Schema
```



```
; SPIR-V
; Version: 1.2
; Generator: Khronos; 33
; Bound: 77
; Schema: 0
```

SPIR-V Beehive Toolkit for code-gen within TornadoVM

ADD: a + b

```
SPIRVId add = module.getNextId();  
blockScope.add(new SPIRVOpIAdd(  
    uint,    // type ID  
    add,     // result  
    id74,    // a  
    id75));  // b
```



```
%add = OpIAdd %uint %74 %75
```

SPIR-V Beehive Toolkit for code-gen within TornadoVM

ADD: a + b

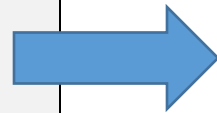
```
SPIRVId add = module.getNextId();  
blockScope.add(new SPIRVOpIAdd(  
    uint,    // type ID  
    add,     // result  
    id74,    // a  
    id75));  // b
```



```
%add = OpIAdd %uint %74 %75
```

Load a[i]

```
SPIRVId idLoad = module.getNextId();  
blockScope.add(new SPIRVOpLoad(  
    ptrCrossGroupUint,  
    idLoad,  
    a_addr, // Load A[i]  
    new SPIRVOptionalOperand<>(  
        SPIRVMemoryAccess.Aligned(  
            new SPIRVLiteralInteger(8)))  
));
```



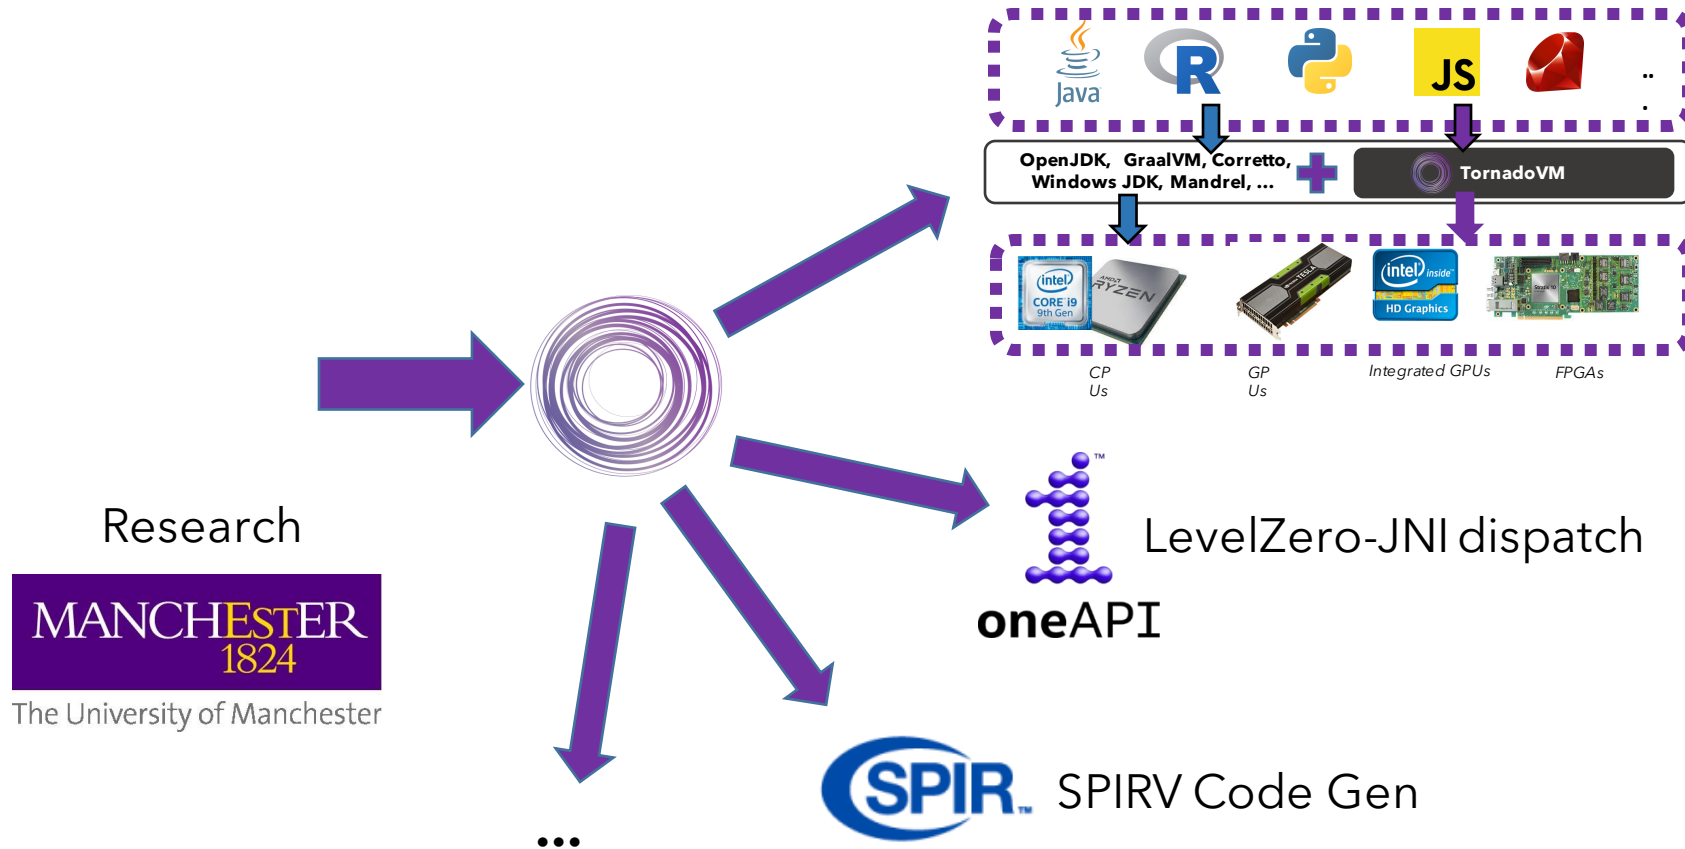
```
%idLoad = OpLoad %_ptr_CrossWorkgroup_uint %addr Aligned 8
```




Creating Value for the Community

MANCHESTER
1824

The University of Manchester





Final remarks



Areas of Interest





<https://www.tornadovm.org/use-cases>

Levenshtein Distance	9.8x
K-Means	9.7x
Hierarchical Clustering	28x

Natural Language Processing

DeepNets	6x and 88x (kernel)
Logistic Regression	14x

Machine Learning and Deep Learning

 449x	 1264 endpoints
Overall Performance Gains	Hardware IoT Deployment

IoT and smart buildings

SLAM-Bench	150x
BlurFilter	> 300x
ViolaJones	22x
RenderTrack	80x

Computer Vision

DFT	4500x
-----	-------

Digital Signal Processing

NBody	> 2000x
-------	---------

Physics Simulation

FinTech

BlackScholes	> 100x
MonteCarlo	> 10x

<https://e2data.eu/blog>

<https://e2data.eu/> (Deliverable 6.3)

MPLR 2020: Transparent acceleration of Java-based deep learning engines

VEE 2019: Dynamic application reconfiguration on heterogeneous hardware

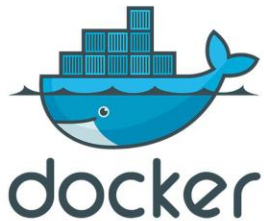
TornadoVM is Open Source and available on GitHub



GPLv2 + CE

<https://github.com/beehive-lab/TornadoVM>

<https://github.com/beehive-lab/tornadovm-installer>



<https://github.com/beehive-lab/docker-tornado>

```
$ docker pull beehivelab/tornado-gpu
```

#And run!

```
$ ./run_nvidia.sh javac.py YourApp
```

```
$ ./run_nvidia.sh tornado YourApp
```





Team

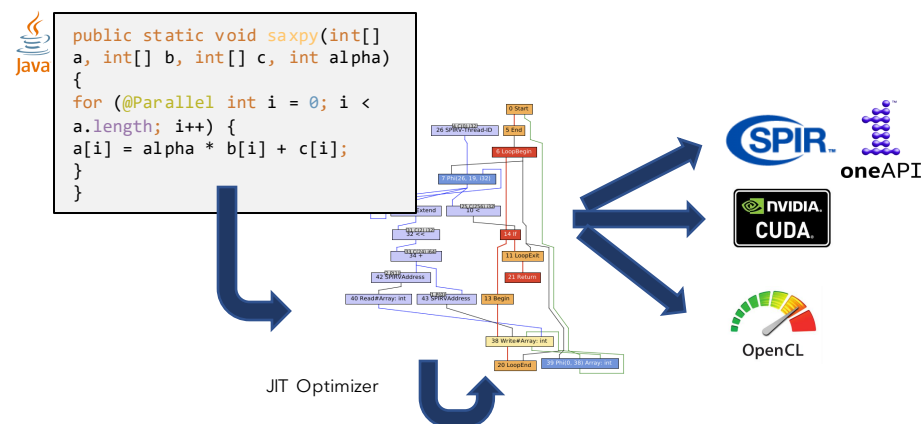
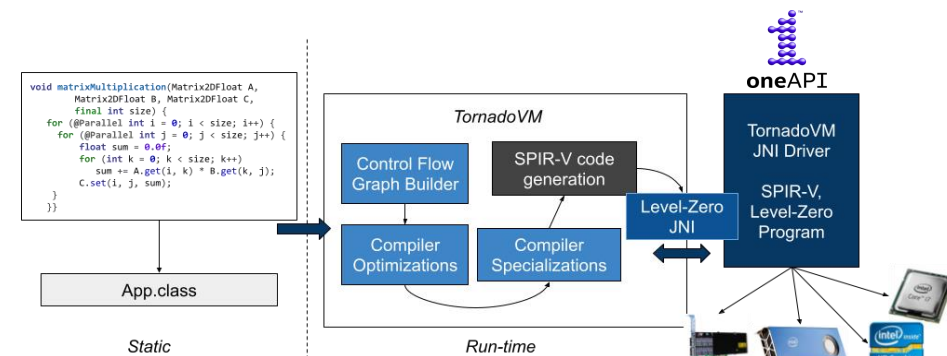
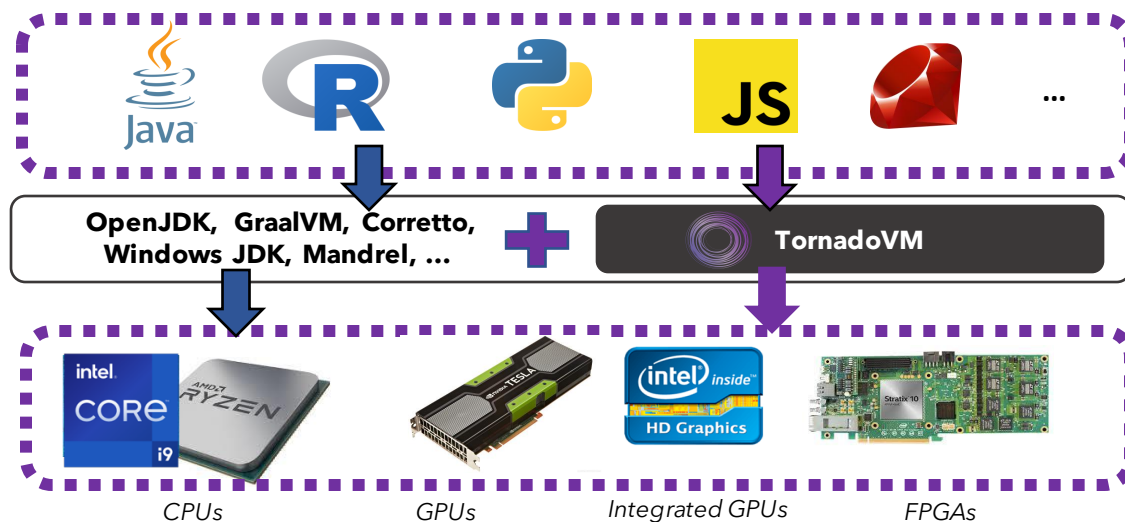
- Academic staff:
Christos Kotselidis
- Research staff:
Juan Fumero
Thanos Stratikopoulos
- PhD Students:
Maria Xekalaki
- Undergraduate Students:
Vinh Pham Van
- Master Students:
Florin Blanaru
- Alumni:
Michail Papadimitriou
James Clarkson
Benjamin Bell
Amad Aslam
Foivos Zakkak
Gyorgy Rethy
Mihai-Christian Olteanu
Ian Vaughan
Ales Kubicek

**We are looking for collaborations
(industrial & academics) -> Let's talk!**





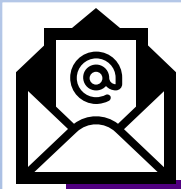
Takeaways



> 100x vs standard JVMs



tornadovm.org

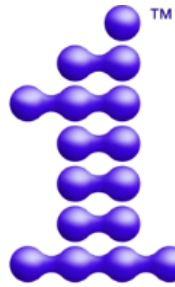


MANCHESTER
1824

The University of Manchester



ELEGANT



oneAPI



European
Commission

Thank you so much for your attention

- Partially supported by the EU Horizon 2020:
 - E2Data 780245
 - ELEGANT 957286
- Partially supported by Intel Grant

Juan Fumero: juan.fumero@manchester.ac.uk



@snatverk