

# Accelerating Java Programs on RISC-V with Vector Instructions via TornadoVM and OCK

Juan Fumero

Research Fellow & Intel Innovator  
TornadoVM's Software Architect



The University of Manchester



@jfumero.bsky.social

J Extension Task Group  
27th February 2025



**TORNADOVM**

# Outline

1. Motivation
2. Overview of TornadoVM
3. Java Acceleration on RISC-V
  - TornadoVM and oneAPI Construction Kit (OCK)
4. OCK for RISC-V
5. Performance Numbers on RISC-V RVV 1.0 – Feb 2025
6. Conclusions & Discussions



# Who am I?

*Dr. Juan Fumero*



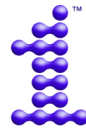
*Research Fellow @ University of Manchester*

Architect and Developer of TornadoVM

oneAPI **Intel Innovator**

- oneAPI Lang SIG

- oneAPI Hardware SIG



**oneAPI**

[juan.fumero@manchester.ac.uk](mailto:juan.fumero@manchester.ac.uk)

[@jfumero.bsky.social](https://bsky.app/profile/jfumero.bsky.social)



*Background:*

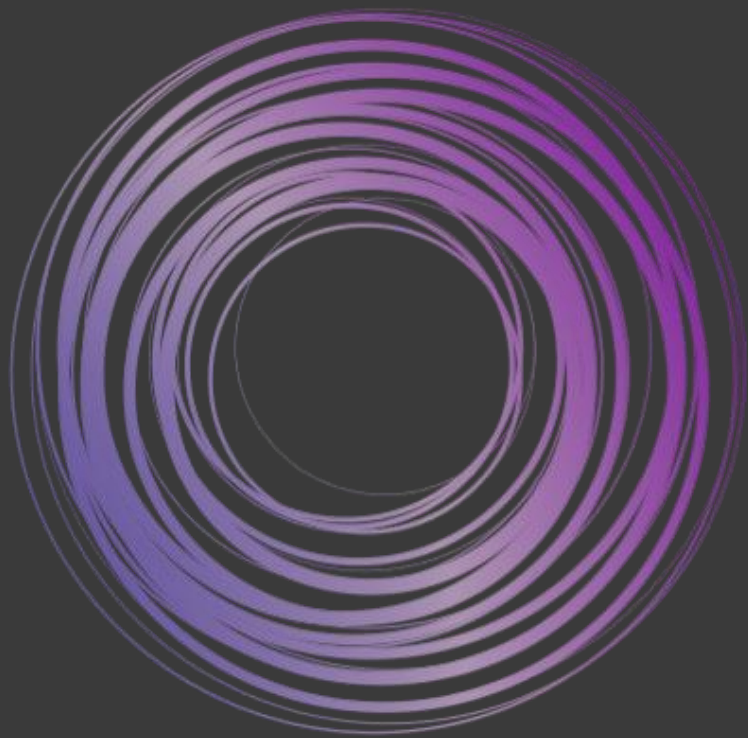


*PhD: Java JIT  
Compilers for GPUs*

**Oracle Labs** *GraalVM/Truffle*



*Intel CilkPlus  
Vectorization for  
Root and GeantV*

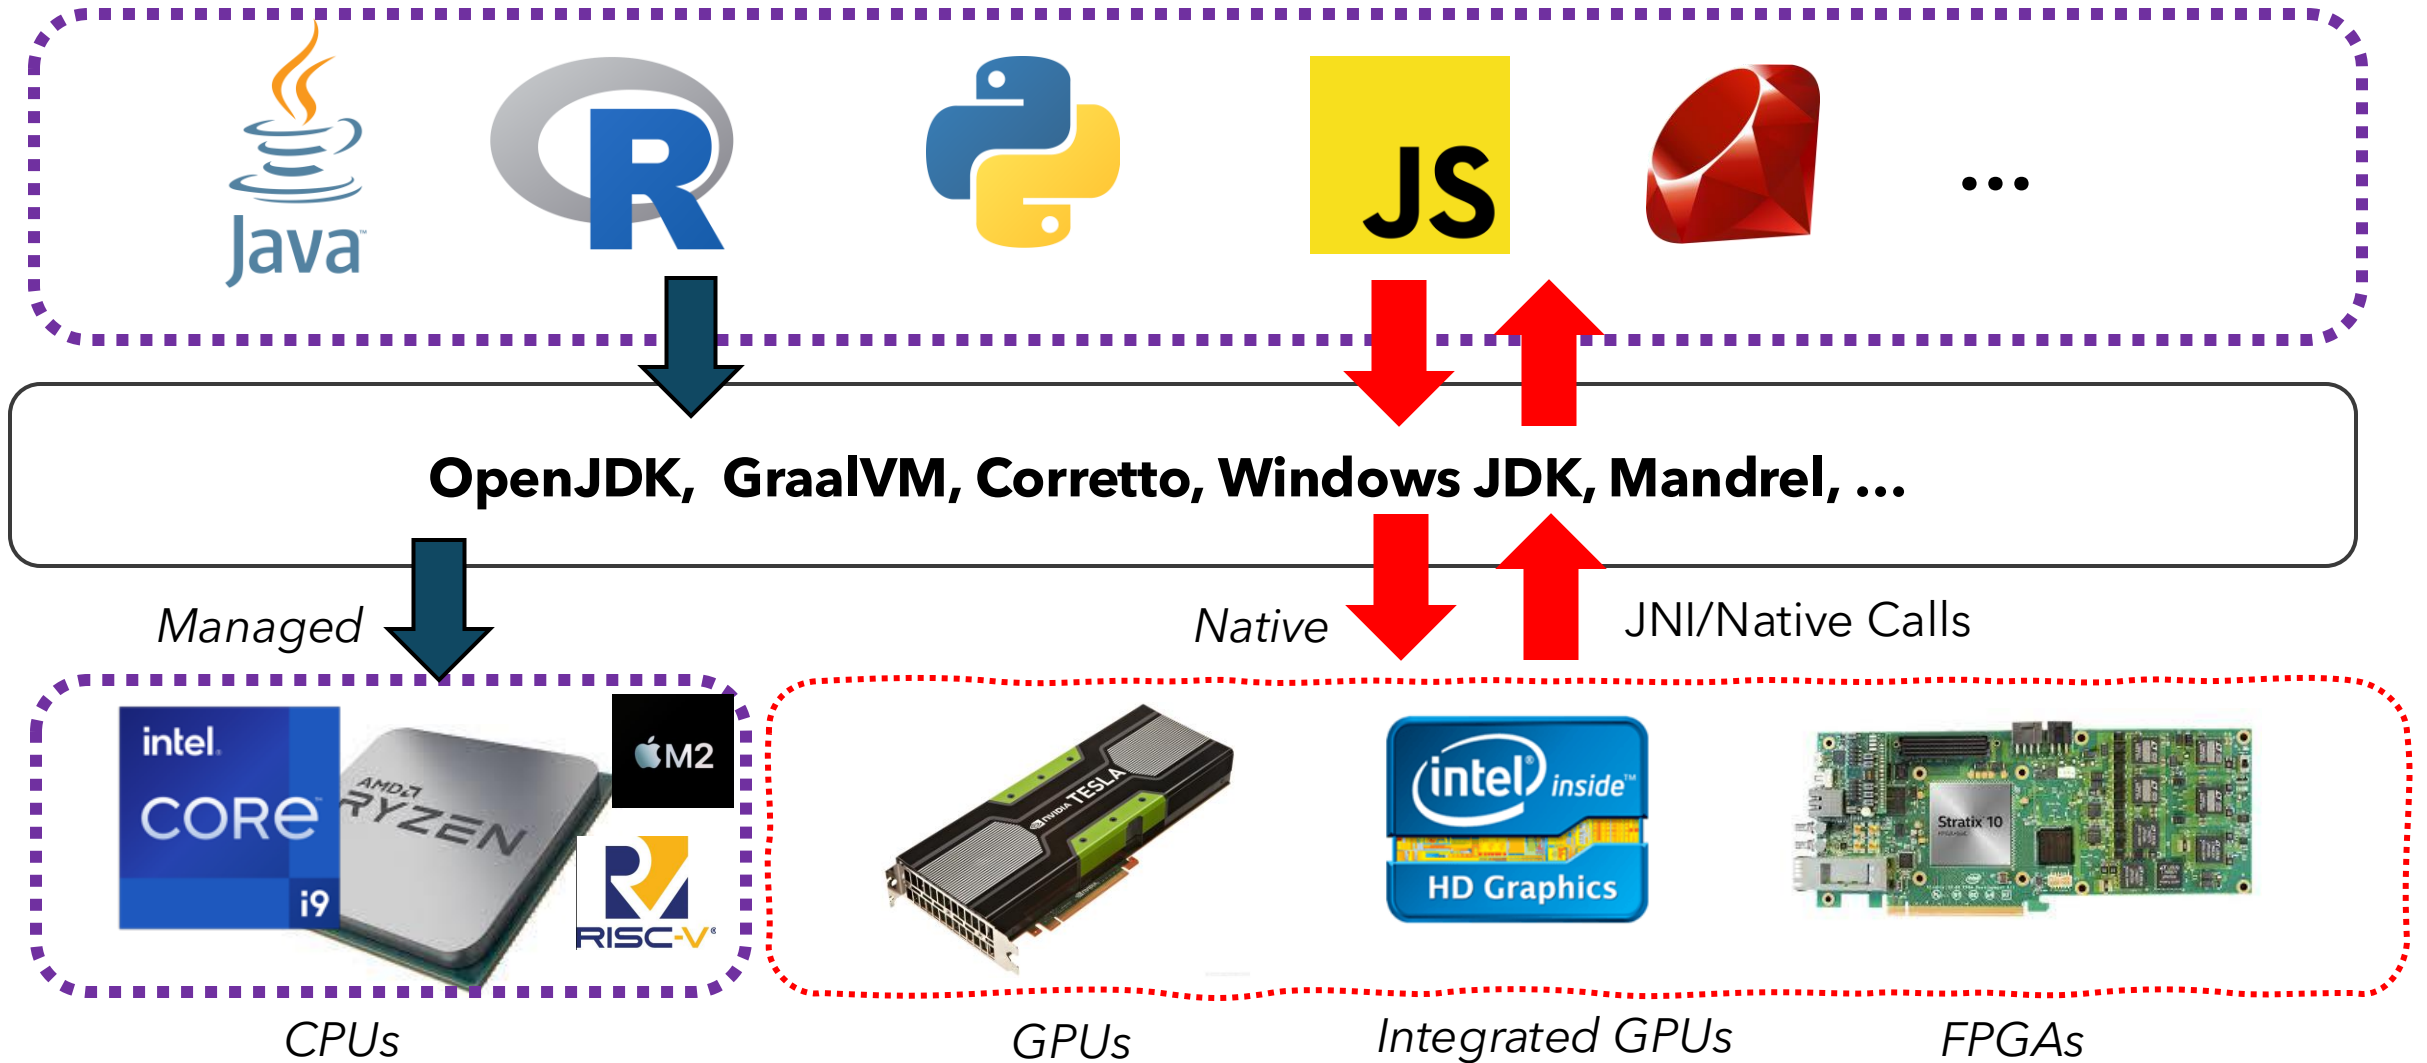


**TORNADOVM**

[www.tornadovm.org](http://www.tornadovm.org)

Motivation

# Enabling Acceleration for Managed Runtime Languages



But what if?



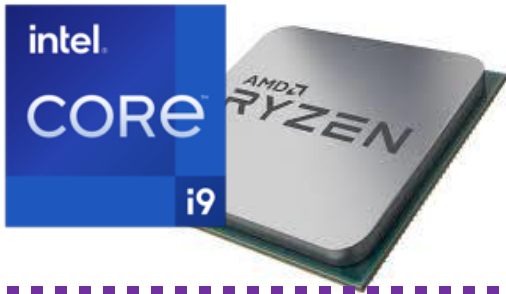
...



**OpenJDK ++ (with Modern Hardware Support)**



*Managed*



*CPUs*



*GPUs*



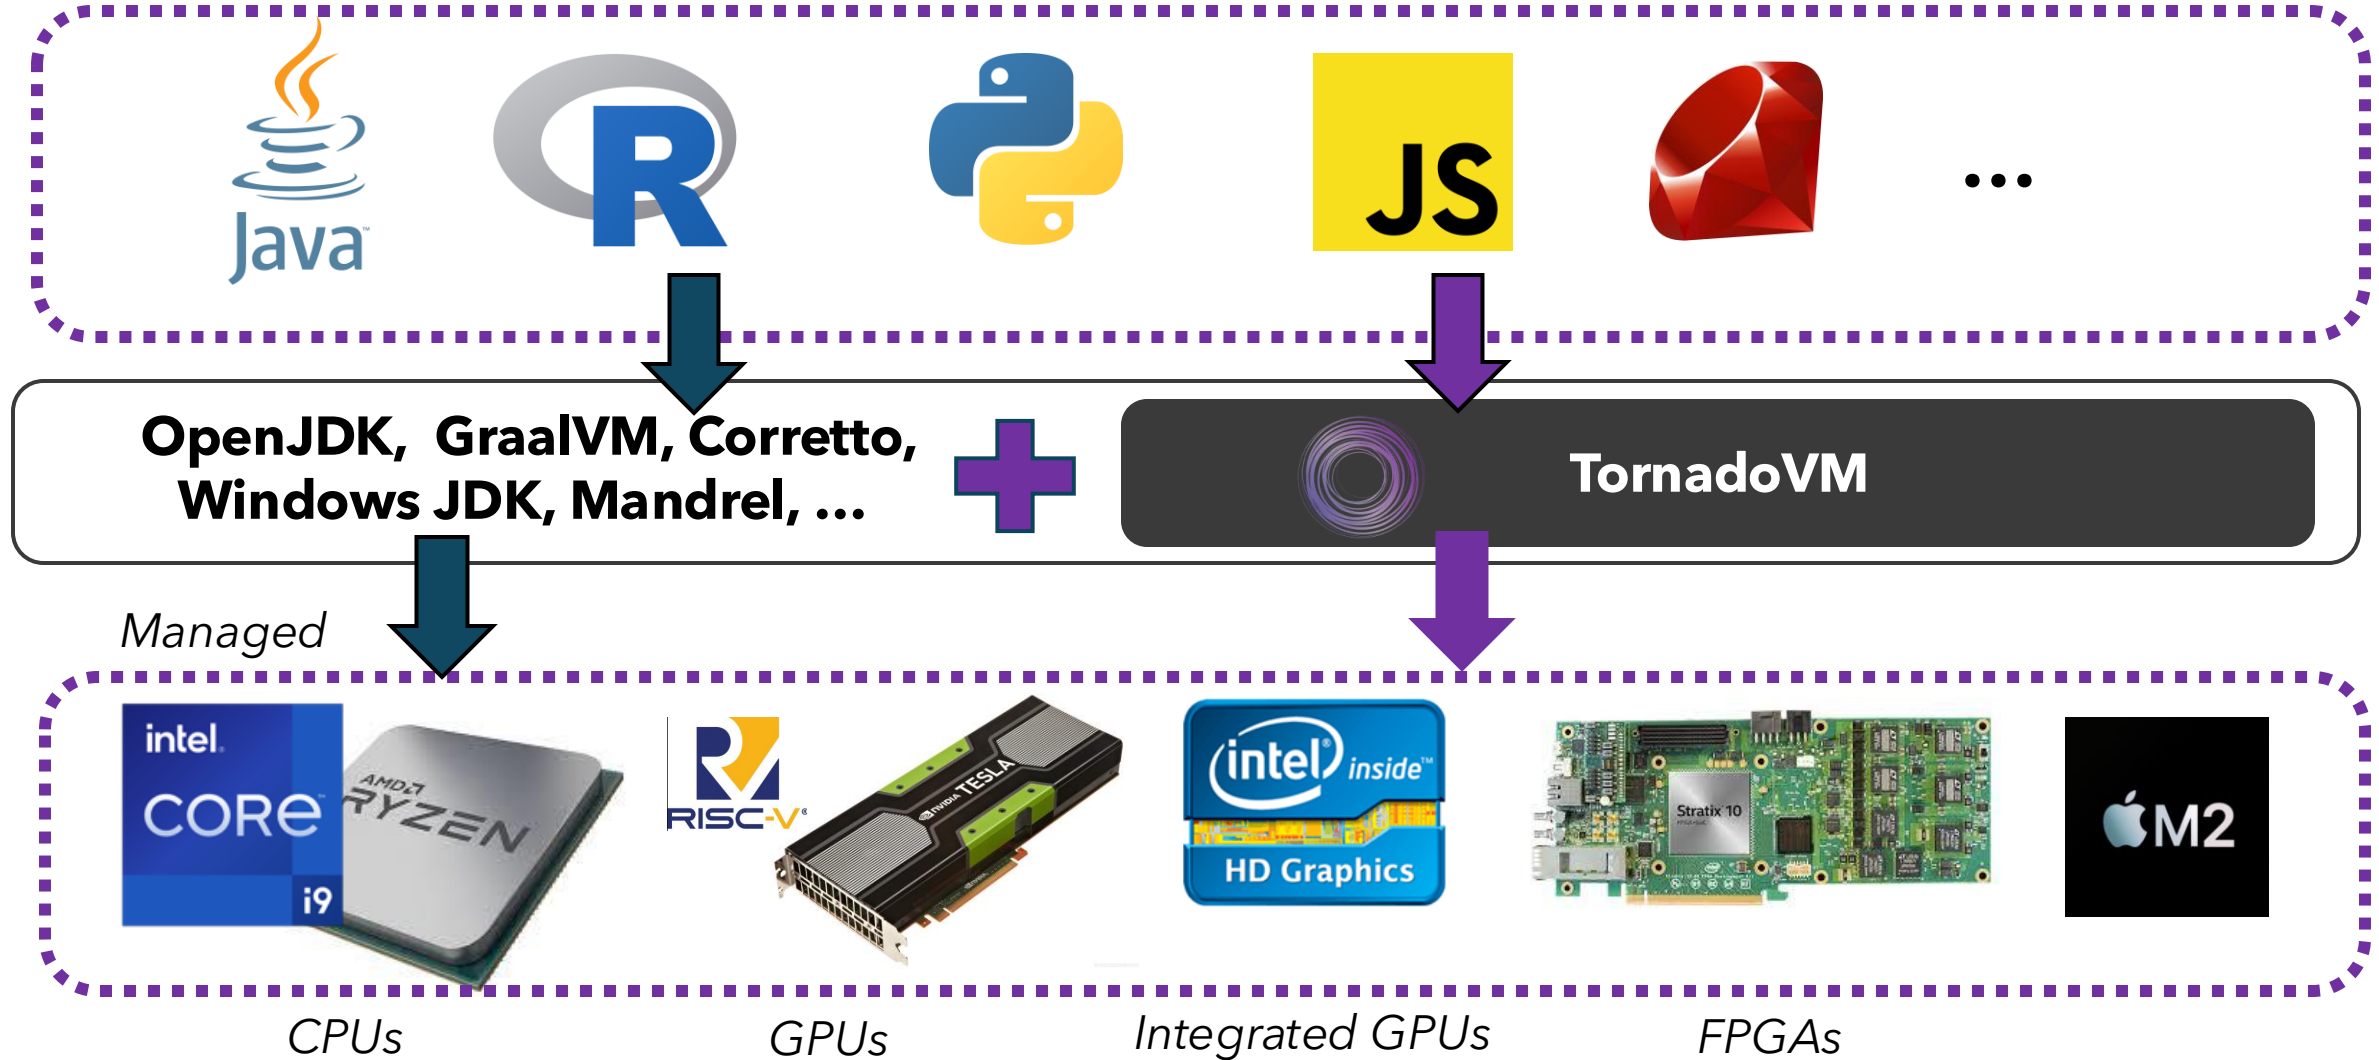
*Integrated GPUs*



*FPGAs*

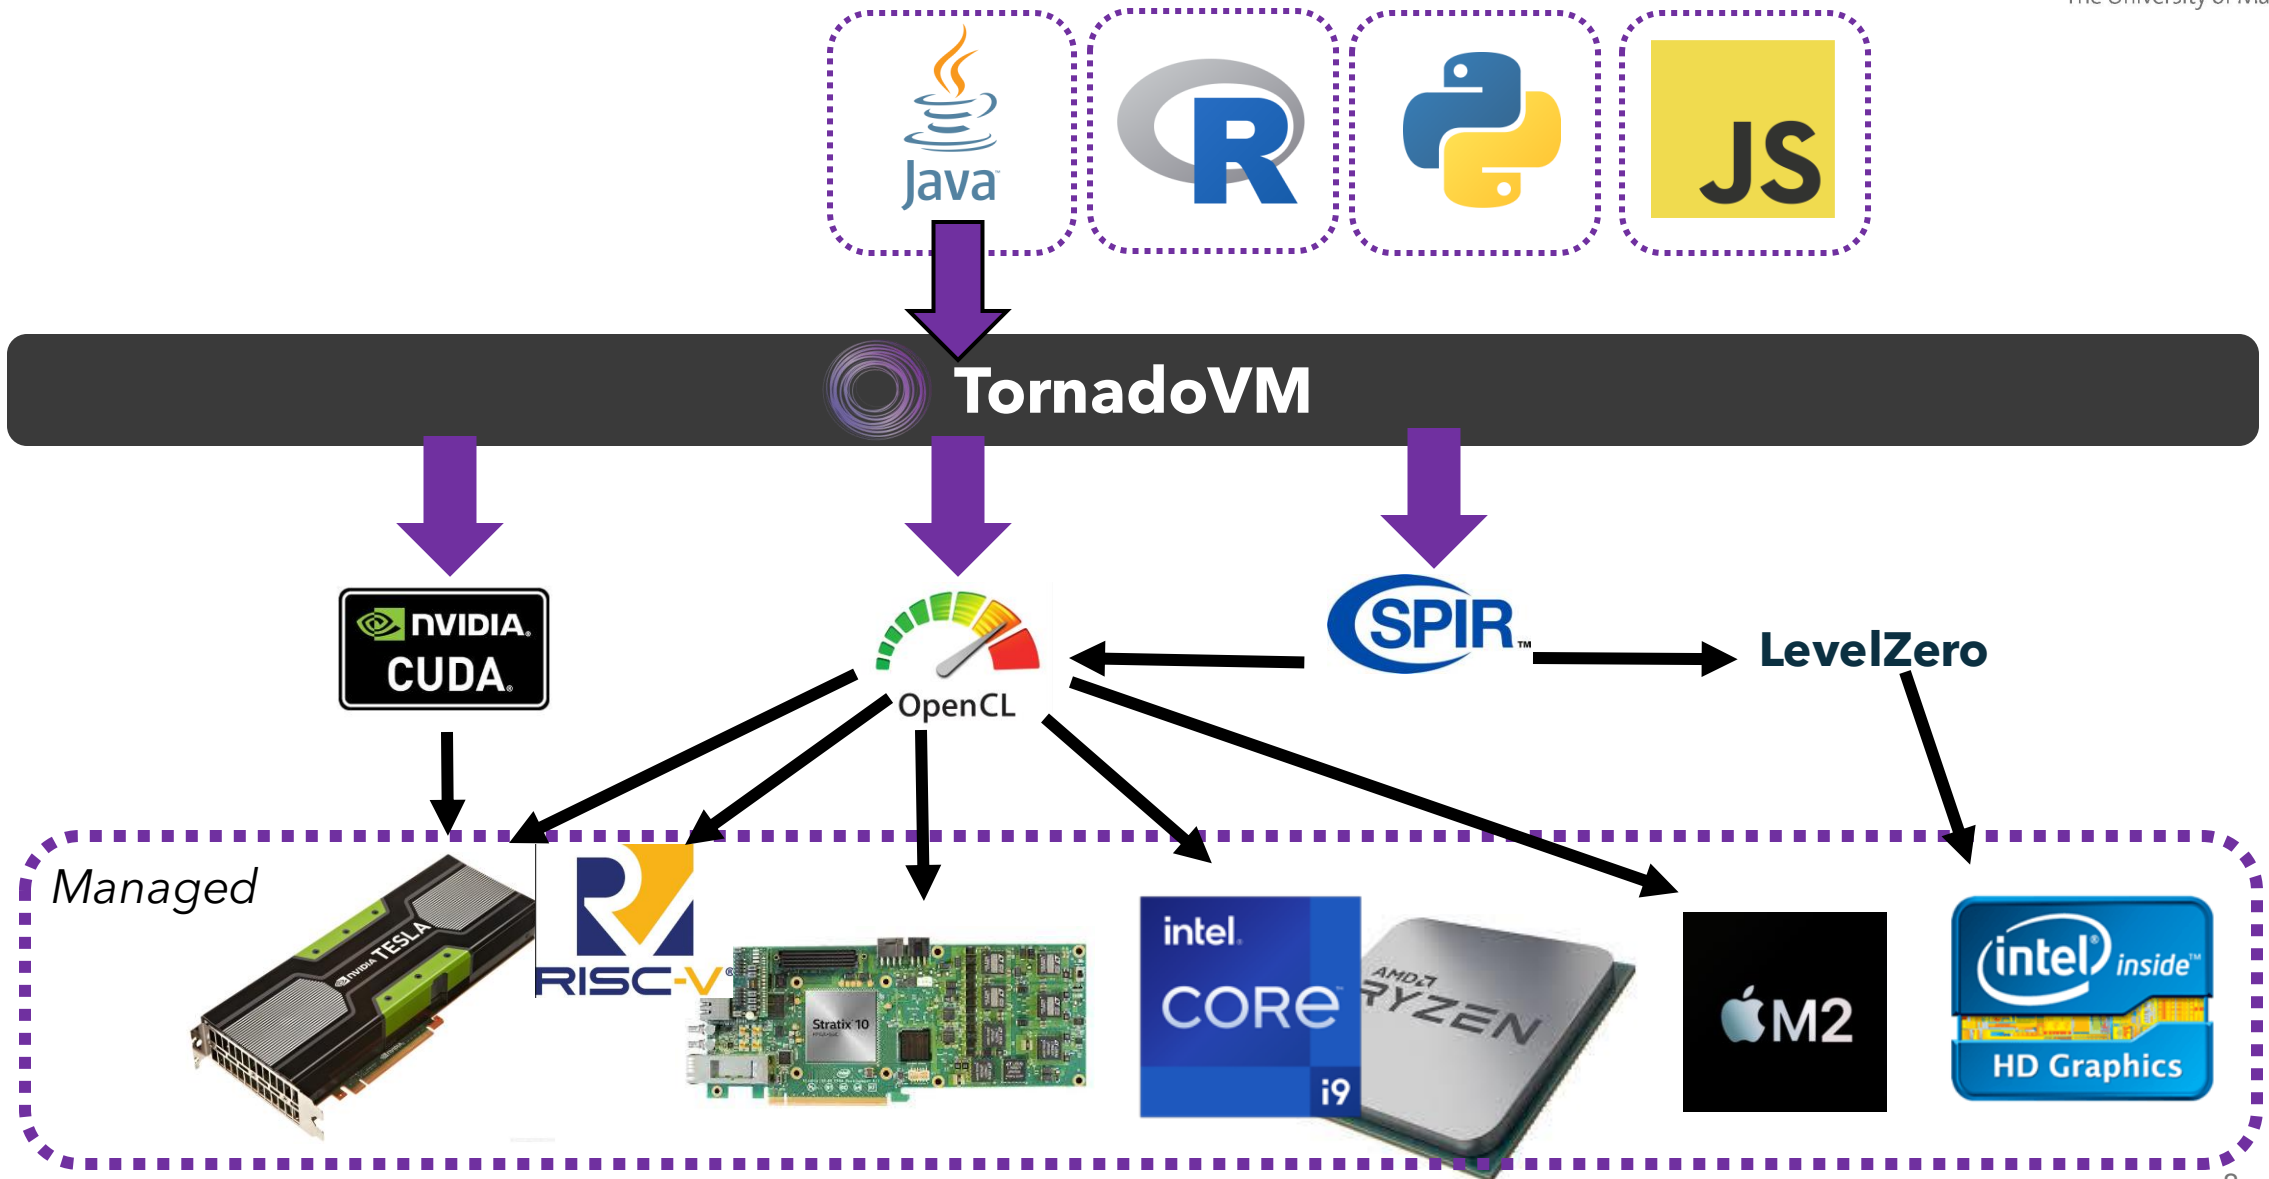


# Enabling Acceleration for Managed Runtime Languages





# Enabling Acceleration for Managed Runtime Languages

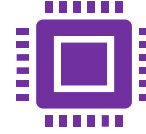




# TornadoVM



*Parallel Programming Framework  
for  
Accelerating  
Java Data Parallel Workloads  
on  
Heterogeneous Hardware*



*Parallel Programming API*

*Three JIT Compilers*

- OpenCL C
- CUDA PTX
- SPIR-V

*Optimizer*

- 100+ optimization phases

*Optimising Runtime System*

*Runs on JDK 21 & JDK 23*



*Dynamic Task Reconfiguration*

*Muti-device*

*Automatic Batch Processing*

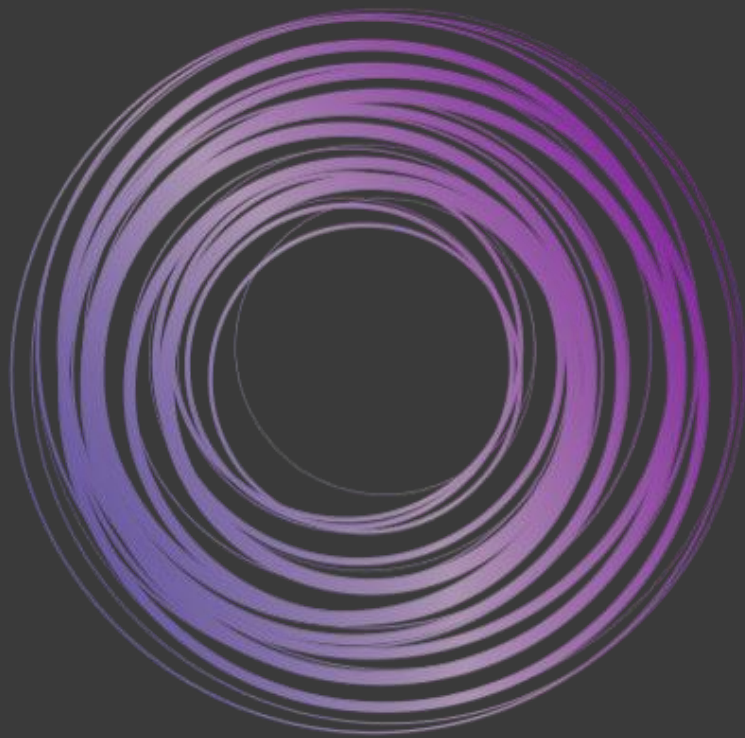
*Automatic Data Management*

*Multi-backend*

and it is **Open Source**:



<https://github.com/beehive-lab/TornadoVM>

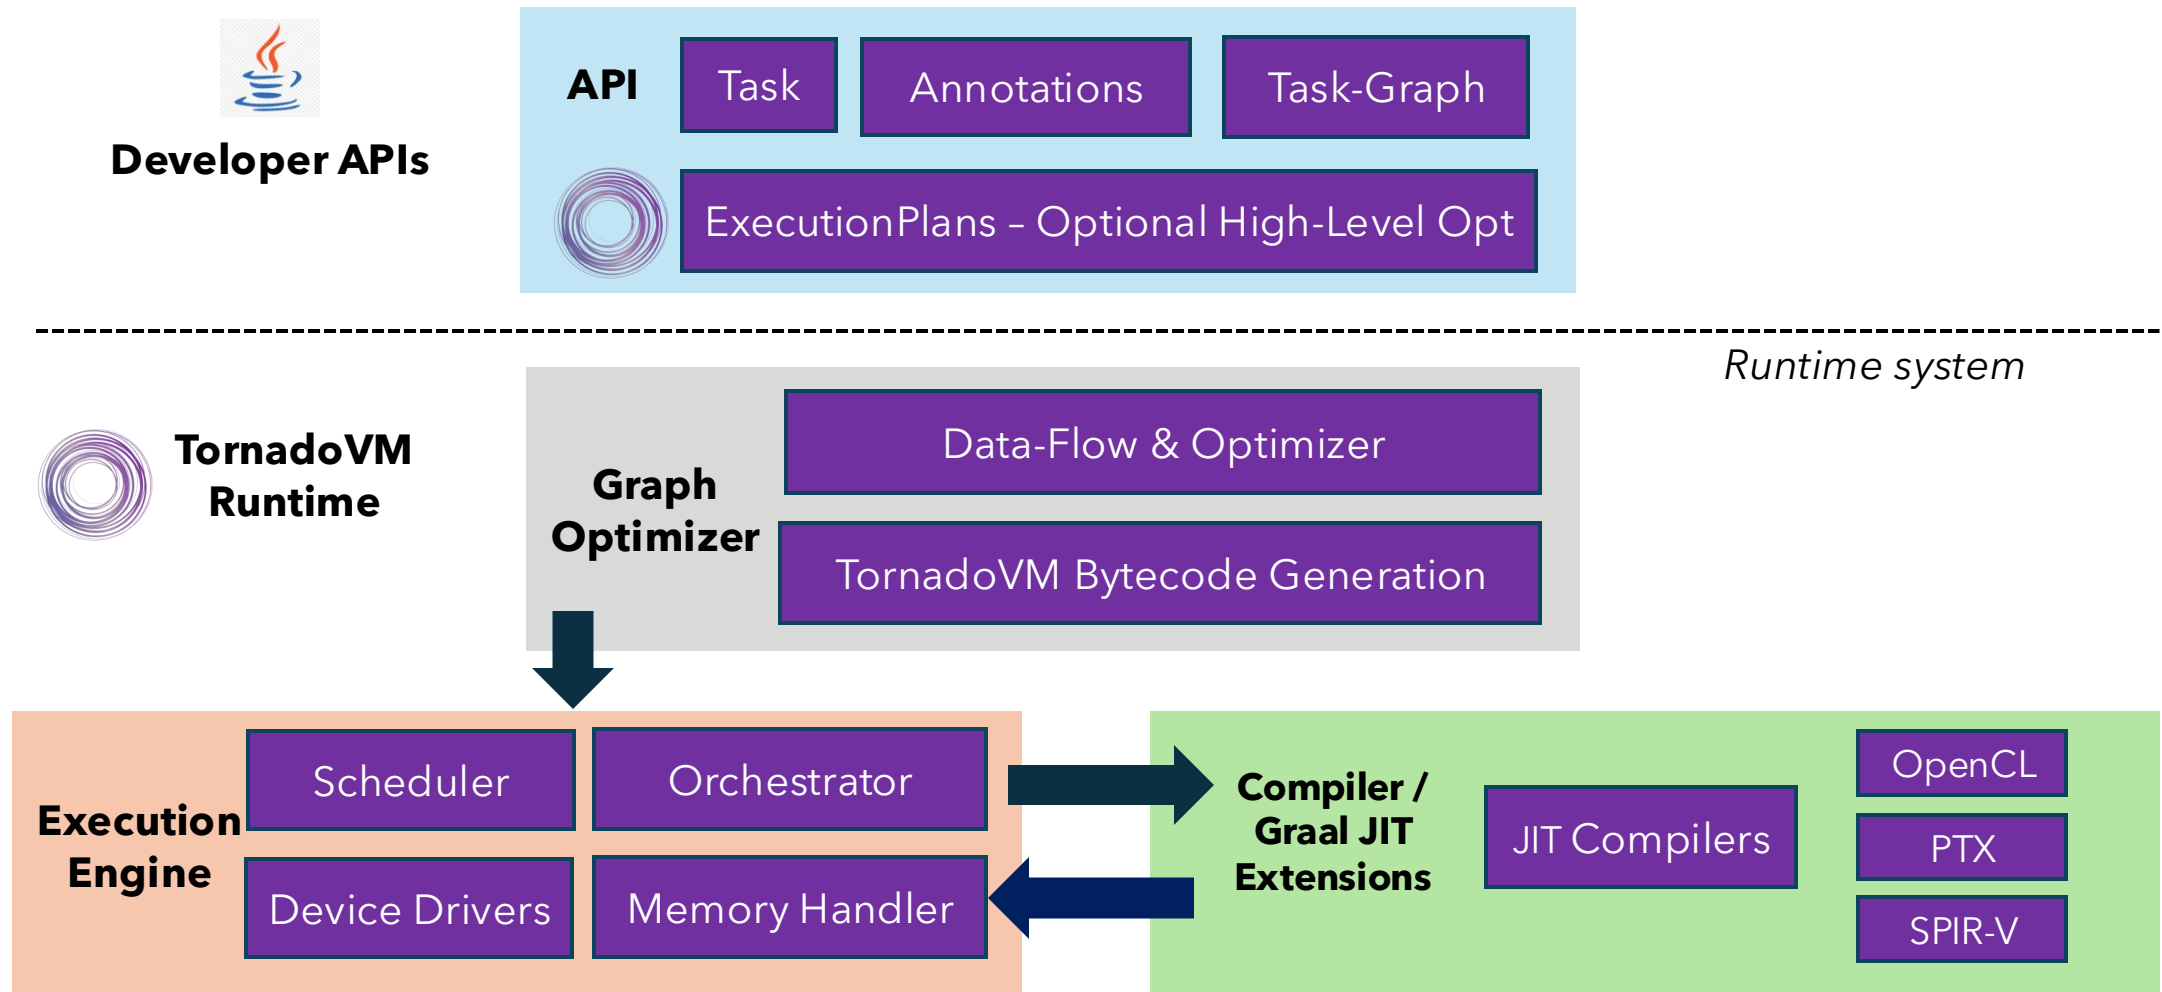


**TORNADOVM**

[www.tornadovm.org](http://www.tornadovm.org)

TornadoVM APIs

# TornadoVM's Software Stack



# Different components of the User APIs

## a) How to represent parallelism within functions/methods?

- A.1: Java annotations for expressing parallelism (**@Parallel, @Reduce**) for Non-Experts
- A.2: Kernel API for GPU experts (use of **kernel context** object)

## b) How to define which methods to accelerate?

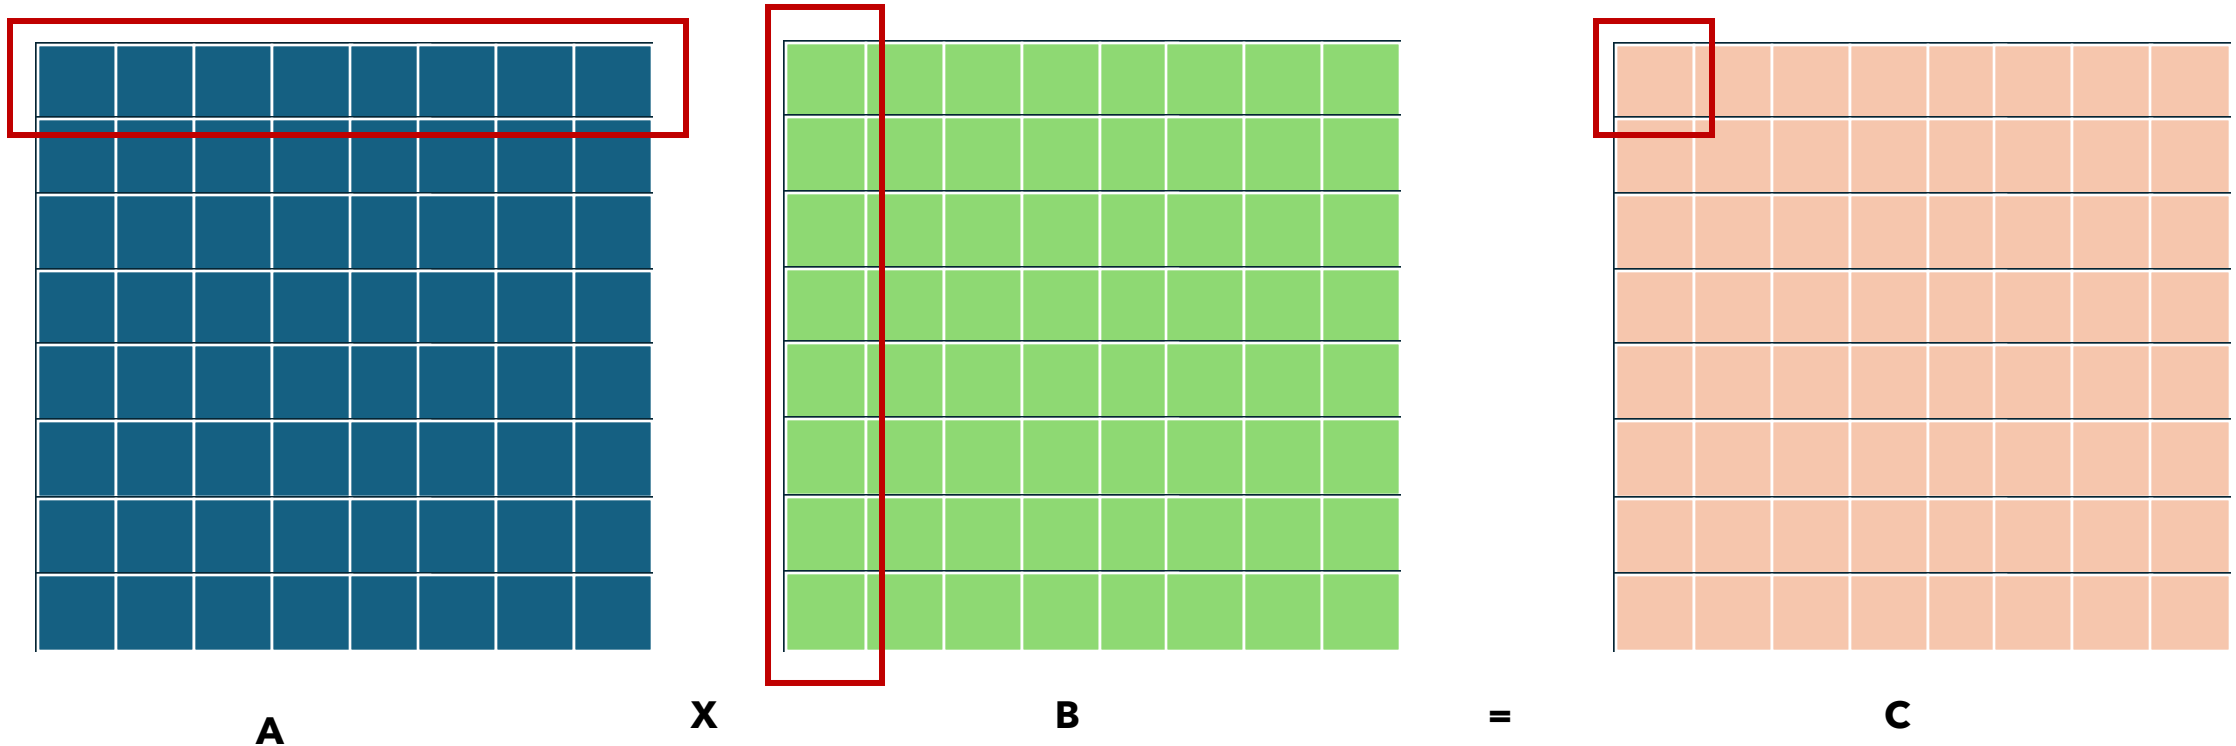
Build a **Task-Graph API** to define data In/Out and the code to be accelerated

## c) How to explore different optimizations?

**Execution Plan**

# Let's Learn the Different API Components with an Example

## Matrix Multiplication



- Widely used on ML and AI
- Linear Algebra Kernel



<https://github.com/beehive-lab/TornadoVM/blob/master/tornado-examples/src/main/java/uk/ac/manchester/tornado/examples/compute/MatrixMultiplication2D.java>

# Example using the TornadoVM Loop Parallel API

```
class Compute {  
    public static void mxm(Matrix2DFloat A, Matrix2DFloat B,  
                           Matrix2DFloat C, final int size) {  
        for (int i = 0; i < size; i++) {  
            for (int j = 0; j < size; j++) {  
                float sum = 0.0f;  
                for (int k = 0; k < size; k++) {  
                    sum += A.get(i, k) * B.get(k, j);  
                }  
                C.set(i, j, sum);  
            }  
        }  
    }  
}
```



<https://github.com/beehive-lab/TornadoVM/blob/master/tornado-examples/src/main/java/uk/ac/manchester/tornado/examples/compute/MatrixMultiplication2D.java>

# Tornado API - example using Annotations

Panama-Based Memory Segments for Off-Heap Data (\*)

```
class Compute {  
    public static void mxx(Matrix2DFloat A, Matrix2DFloat B,  
                           Matrix2DFloat C, final int size) {  
        for (@Parallel int i = 0; i < size; i++) {  
            for (@Parallel int j = 0; j < size; j++) {  
                float sum = 0.0f;  
                for (int k = 0; k < size; k++) {  
                    sum += A.get(i, k) * B.get(k, j);  
                }  
                C.set(i, j, sum);  
            }  
        }  
    }  
}
```

We add the parallel annotation as a hint for the compiler

We only have 2 annotations:

**@Parallel**  
**@Reduce**



We add the @Parallel Annotation



# Tornado API - example using Kernel Context

```
class Compute {  
    public static void mxm(Matrix2DFloat A, Matrix2DFloat B,  
                           Matrix2DFloat C, final int size,  
                           KernelContext context) {  
        int idx = context.globalIdx;  
        int jdx = context.globalIdx;  
        float sum = 0.0f;  
        for (int k = 0; k < size; k++)  
            sum += A.get(idx, k) * B.get(k, jdx);  
        C.set(idx, jdx, sum);  
    }  
}
```

Kernel-Context accesses  
thread ids, local memory  
and barriers

It needs a **Grid of Threads** to  
be passed during the kernel  
launch



Alternative API for expert programmers. It offers more control

## How to identify which methods to accelerate? --> TaskGraph

```
TaskGraph taskGraph = new TaskGraph("myComputeGraph")  
  
    .transferToDevice(DataTransferMode.EVERY_EXECUTION , matrixA, matrixB)  
  
    .task("parallelTask", Compute::mxm, matrixA, matrixB, matrixC, size)  
  
    .transferToHost(DataTransferMode.EVERY_EXECUTION, matrixC);
```



Host Code

TaskGraph is a new TornadoVM object exposed to developers to define :

- a) **The code to be accelerated** (which Java methods?)
- b) **The data (Input/Output)** and how data should be streamed



# Defining and Running Execution Plans

## How to run/explore different optimizations? --> ExecutionPlan

```
TornadoExecutionPlan executionPlan = new TornadoExecutionPlan(taskGraph.snapshot());  
  
// optional: change runtime parameters/optimize the runtime  
executionPlan.withWarmUp()  
              .withProfiler(...)  
              .withDynamicReconfiguration(PERFORMANCE, ...);  
  
// blocking call: the execute is where JIT compilation + execution happen  
TornadoExecutionResult result = executionPlan.execute();  
  
long elapsedKernelTime = result.getProfiler().getDeviceKernelTime(); // in ns
```



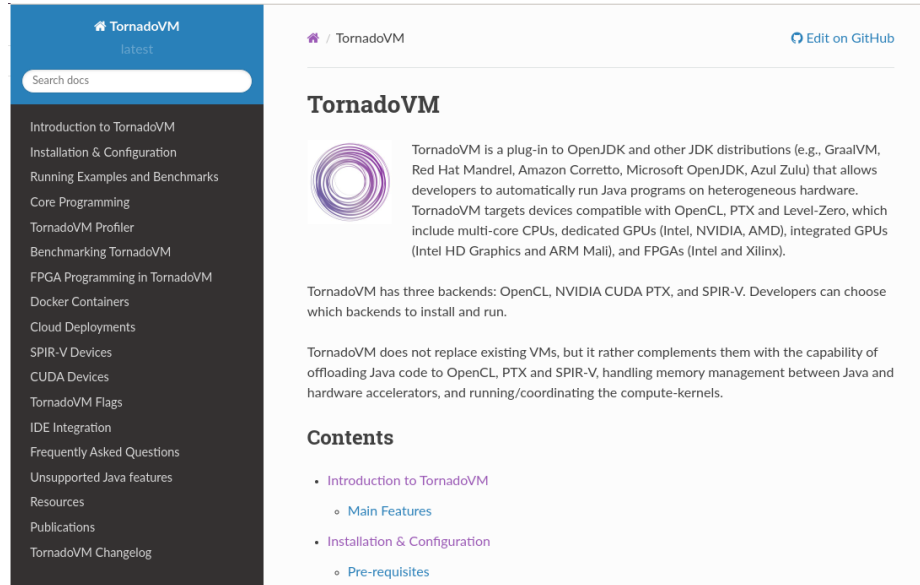
Host Code



### Optional High-Level Optimization Pipelines:

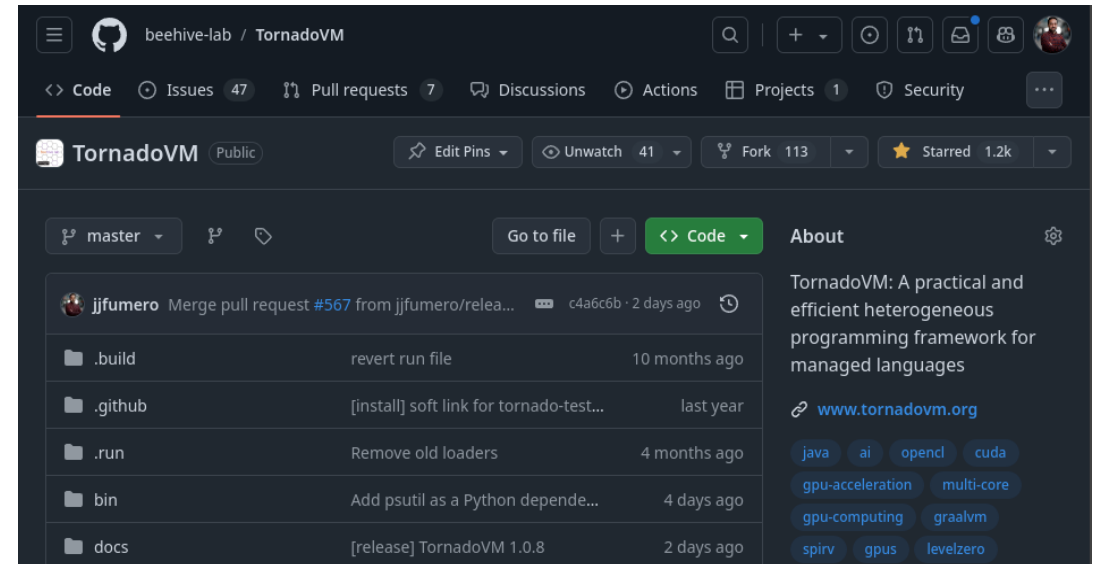
- Enable/Disable Profiler
- Enable Warmup
- Enable Dynamic Reconfiguration
- Enable Batch Processing
- Enable Thread Scheduler (no need for recompilation for different grids schedulers)

# To learn more about the APIs



The screenshot shows the TornadoVM documentation website. The header includes the TornadoVM logo and a search bar. The left sidebar lists various topics: Introduction to TornadoVM, Installation & Configuration, Running Examples and Benchmarks, Core Programming, TornadoVM Profiler, Benchmarking TornadoVM, FPGA Programming in TornadoVM, Docker Containers, Cloud Deployments, SPIR-V Devices, CUDA Devices, TornadoVM Flags, IDE Integration, Frequently Asked Questions, Unsupported Java features, Resources, Publications, and TornadoVM Changelog. The main content area features the TornadoVM logo, a description of the framework as a plug-in to OpenJDK and other JDK distributions, and a list of contents including Introduction to TornadoVM, Main Features, Installation & Configuration, and Pre-requisites.

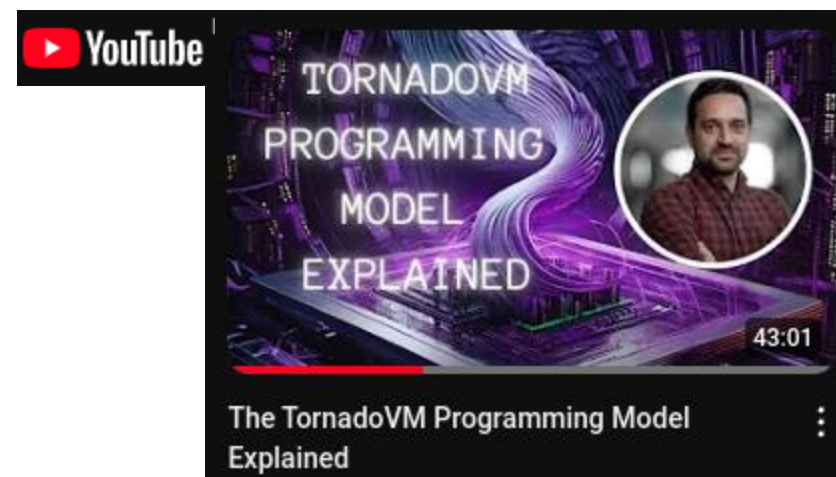
<https://tornadovm.readthedocs.io/en/latest/>

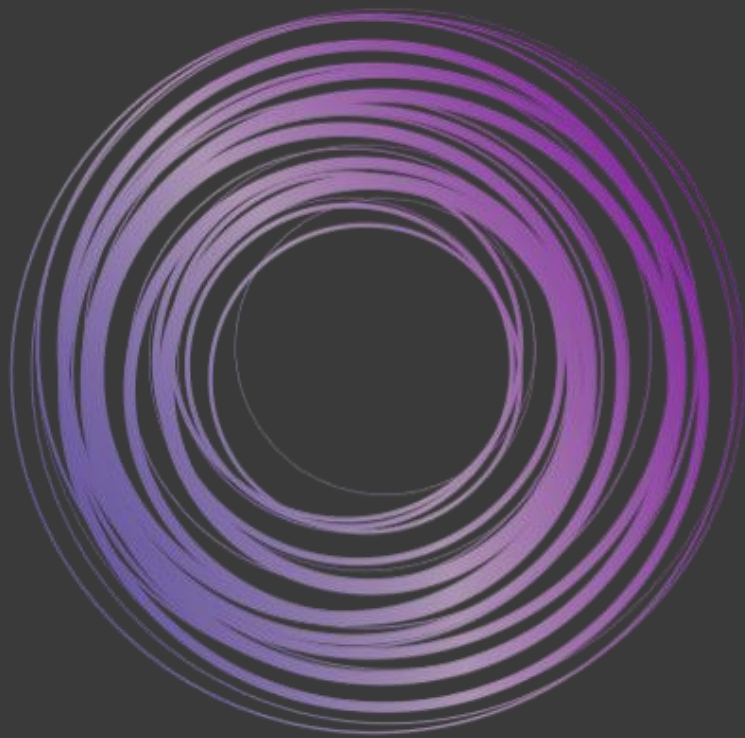


The screenshot shows the TornadoVM GitHub repository page. The header includes the repository name, a search bar, and navigation links for Code, Issues (47), Pull requests (7), Discussions, Actions, Projects (1), and Security. The repository is public and has 41 unwatched items, 113 forks, and 1.2k stars. The main content area shows a list of files and folders, including .build, .github, .run, bin, and docs, with their respective commit messages and dates. The right sidebar provides an overview of the repository, including a description, a link to the website, and tags for various technologies like java, ai, opengl, cuda, gpu-acceleration, multi-core, gpu-computing, graalvm, spirv, gpus, and levelzero.



<https://github.com/beehive-lab/TornadoVM>





**TORNADOVM**

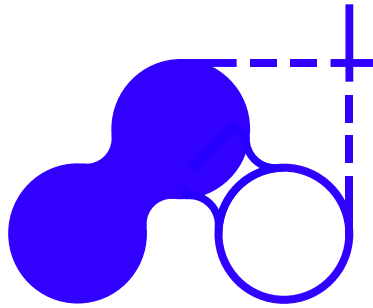
[www.tornadovm.org](http://www.tornadovm.org)

How TornadoVM can be  
used on RISC-V for RVV  
1.0?

# oneAPI Construction Kit (OCK)



<https://github.com/uxlfoundation/oneapi-construction-kit>



*"Programming Framework to enable hardware platforms to access open standards (e.g., SYCL, OpenCL, etc) and it can be used to extend the oneAPI software ecosystem to custom compute architectures [1]"*



Part of the UXL Foundation: <https://uxlfoundation.org/>

Working with Codeplay in the AERO EU Project [2]



[1] <https://developer.codeplay.com/products/oneapi/construction-kit/home/>

[2] <https://aero-project.eu/>

# Running on Multi-Vendor CPUs: ARM, RISC-V, Intel Using the **oneAPI Construction Kit (OCK)**



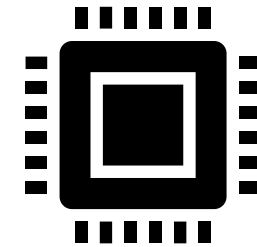
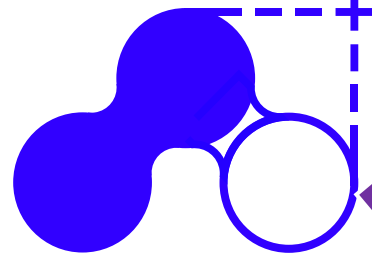
<https://jjfumero.github.io/posts/2024/09/10/tornadovm-ock>



Instructions how to build



arm



*Custom accelerators*



# Running on Multi-Vendor CPUs: ARM, RISC-V, Intel Using the **oneAPI Construction Kit (OCK)**



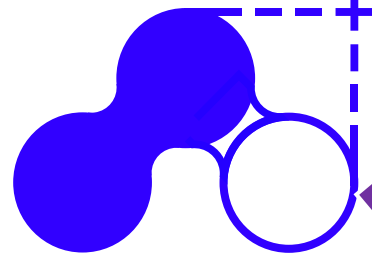
<https://jjfumero.github.io/posts/2024/09/10/tornadovm-ock>



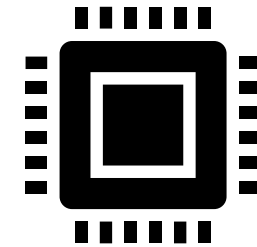
Instructions how to build



arm

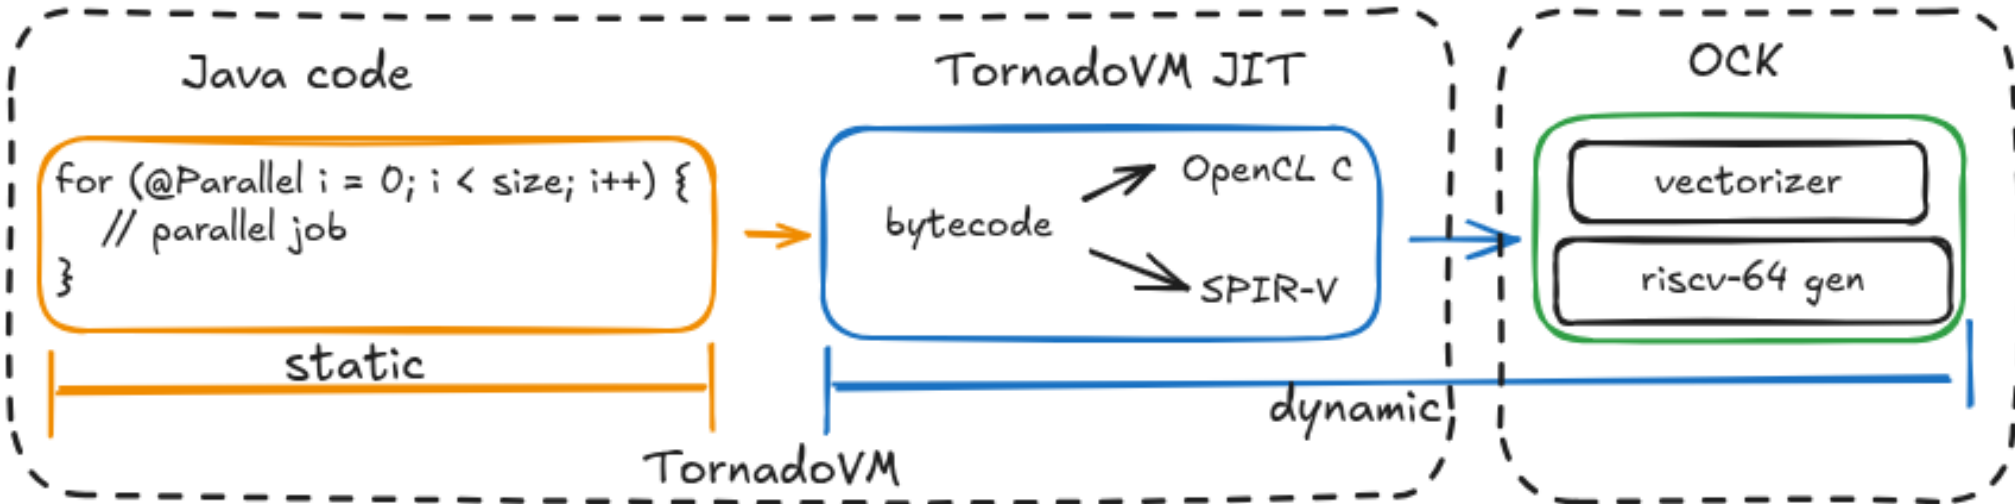


*This talk*



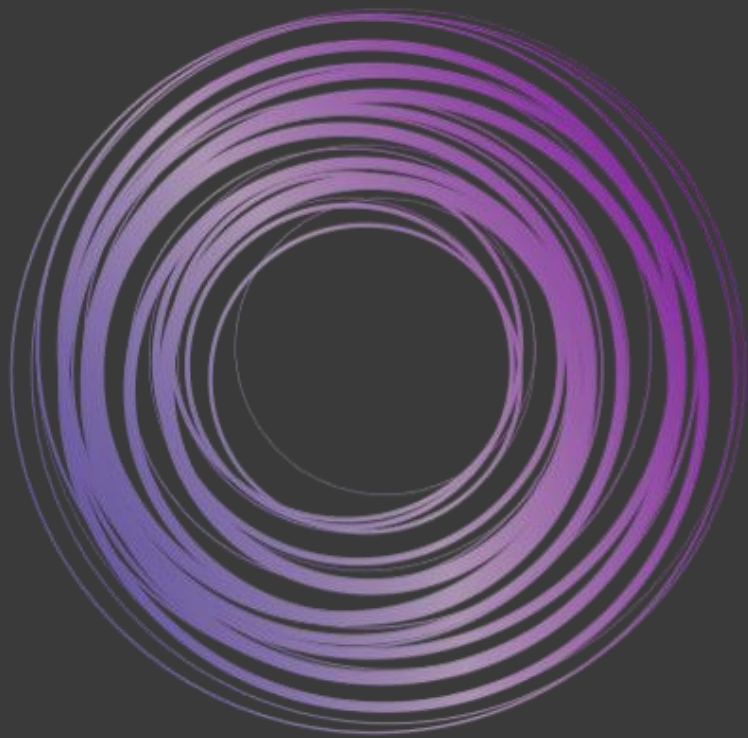
*Custom accelerators*

# JIT Compilation Process



TornadoVM contains more than 50 compilation phases

Compilation is per-device, per architecture. This is crucial to get performance portability



**TORNADOVM**

[www.tornadovm.org](http://www.tornadovm.org)



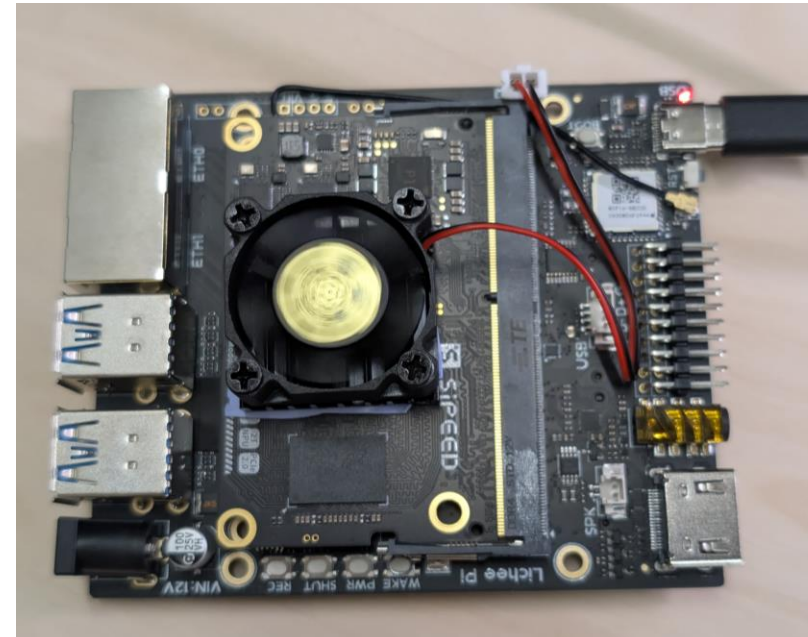
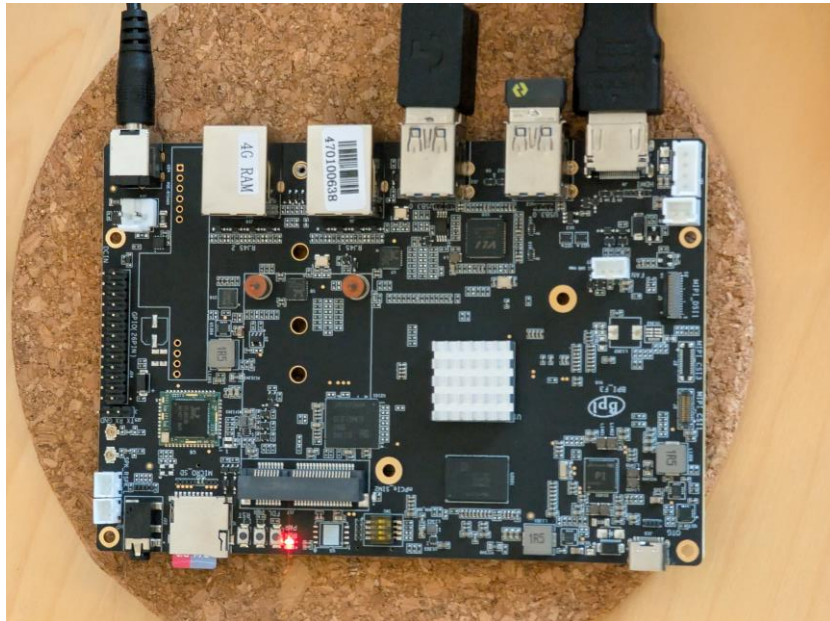
Performance

## SBC used

RISC-V Spacemit K1 Processor

- Banana PI BPI F3
- Sipeed Lichee PI 3A

Octacore RISC-V @ 1.6GHz  
RVA22 and 256-bit RVV 1.0  
1MB shared L2 Cache  
Max 16GB of RAM

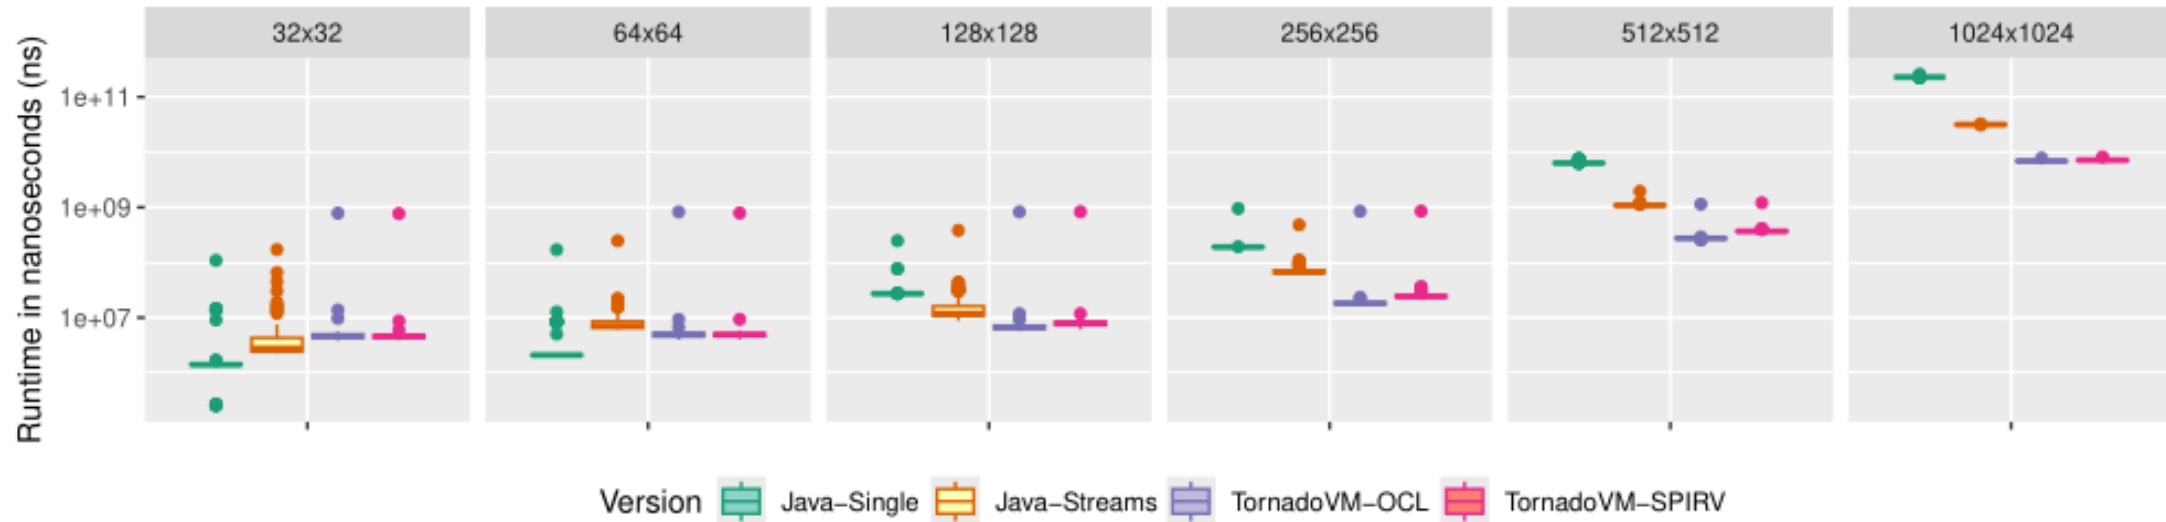


# Things to be aware of

- Need to be extremely patient. This development board is extremely slow.
- ~4 days to compile LLVM on the RISC-V
  - o Single thread compilation to avoid swapping
- ~2 days to compile OCK
  - o Multiple configurations:
    - **Refsi** (RISC-V Emulator on RISC-V)
    - No vectorization
    - **With auto-vectorization**
- Use active cooling to avoid throttling



# Performance of MxM on RISC-V (Banana PI 4GB)



- For small matrices, stay on single core.
  - No auto-vectorization for Java on RISC-V as in FEB 2025 (we checked the generated Assembly code)
- Java streams is 2.3x-7.2x faster than single core for larger matrices
- TornadoVM + OCK achieves 4.1x-33x faster than Java seq, and up to 4.6x vs Streams
- SPIR-V backend is bit slower compared to OpenCL C

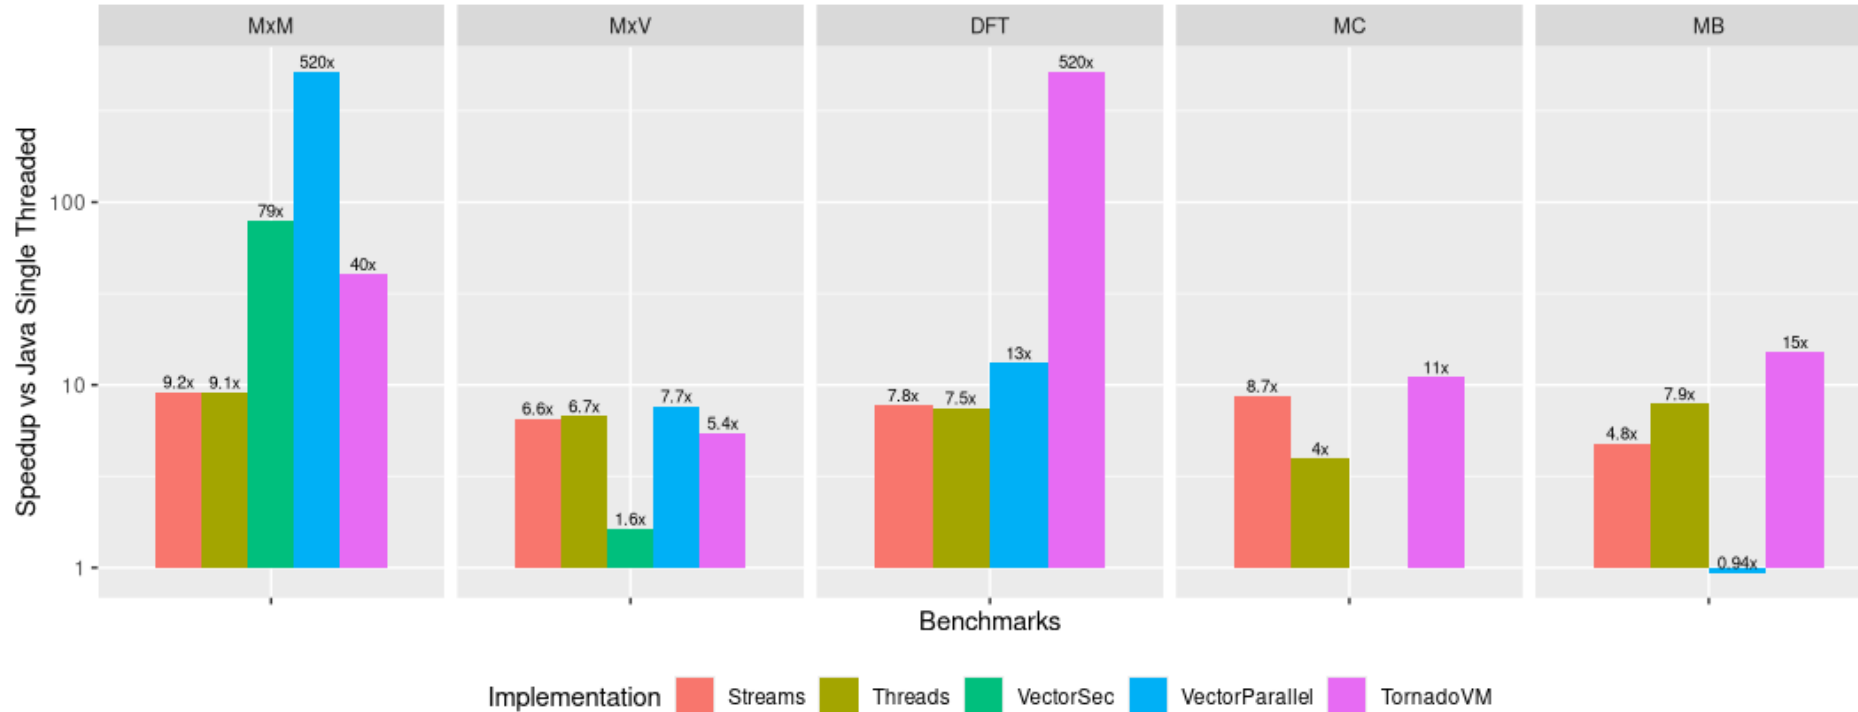
TornadoVM 1.0.10-dev  
(ec667bd65)  
LLVM 19.1.5  
GCC 13.2  
OpenJDK 21.0.5  
Banana PI F3



# Performance on RISC-V Spacemit K1 (Lichee PI3A)

Performance of TornadoVM on RISC-V Spacemit K1

The higher, the better.



TornadoVM Version: 05539e7

Explicit Vector API is quite competitive  
Hard to express Vector API in some applications

TornadoVM: 0553ae7 (1.0.11-dev)  
LLVM 19.1.5  
GCC 13.2  
OpenJDK 21.0.5  
Lichee PI 3A



# Vector API uses Intrinsic

Java JIT Compiled Code (C2), using `-XX:+UnlockDiagnosticVMOptions -XX:+PrintAssembly`

```
0x00007f469cdb7a0f:    cmp    %r11,%r10
0x00007f469cdb7a12:    jne     0x00007f469cdb8ad4
0x00007f469cdb7a18:    mov     %rcx,%rax    ;*invokestatic reductionCoerced {reexecute=0 rethrow=0 return_oop=0}
                                ; - jdk.incubator.vector.FloatVector::reduceLanesTemplate@78 (line 2644)
                                ; - jdk.incubator.vector.Float256Vector::reduceLanes@2 (line 324)
                                ; - tornadovm.benchmarks.DFT::lambda$computeWithParallelVectorAPI$1@263 (line
136)
```

Java Source Code:

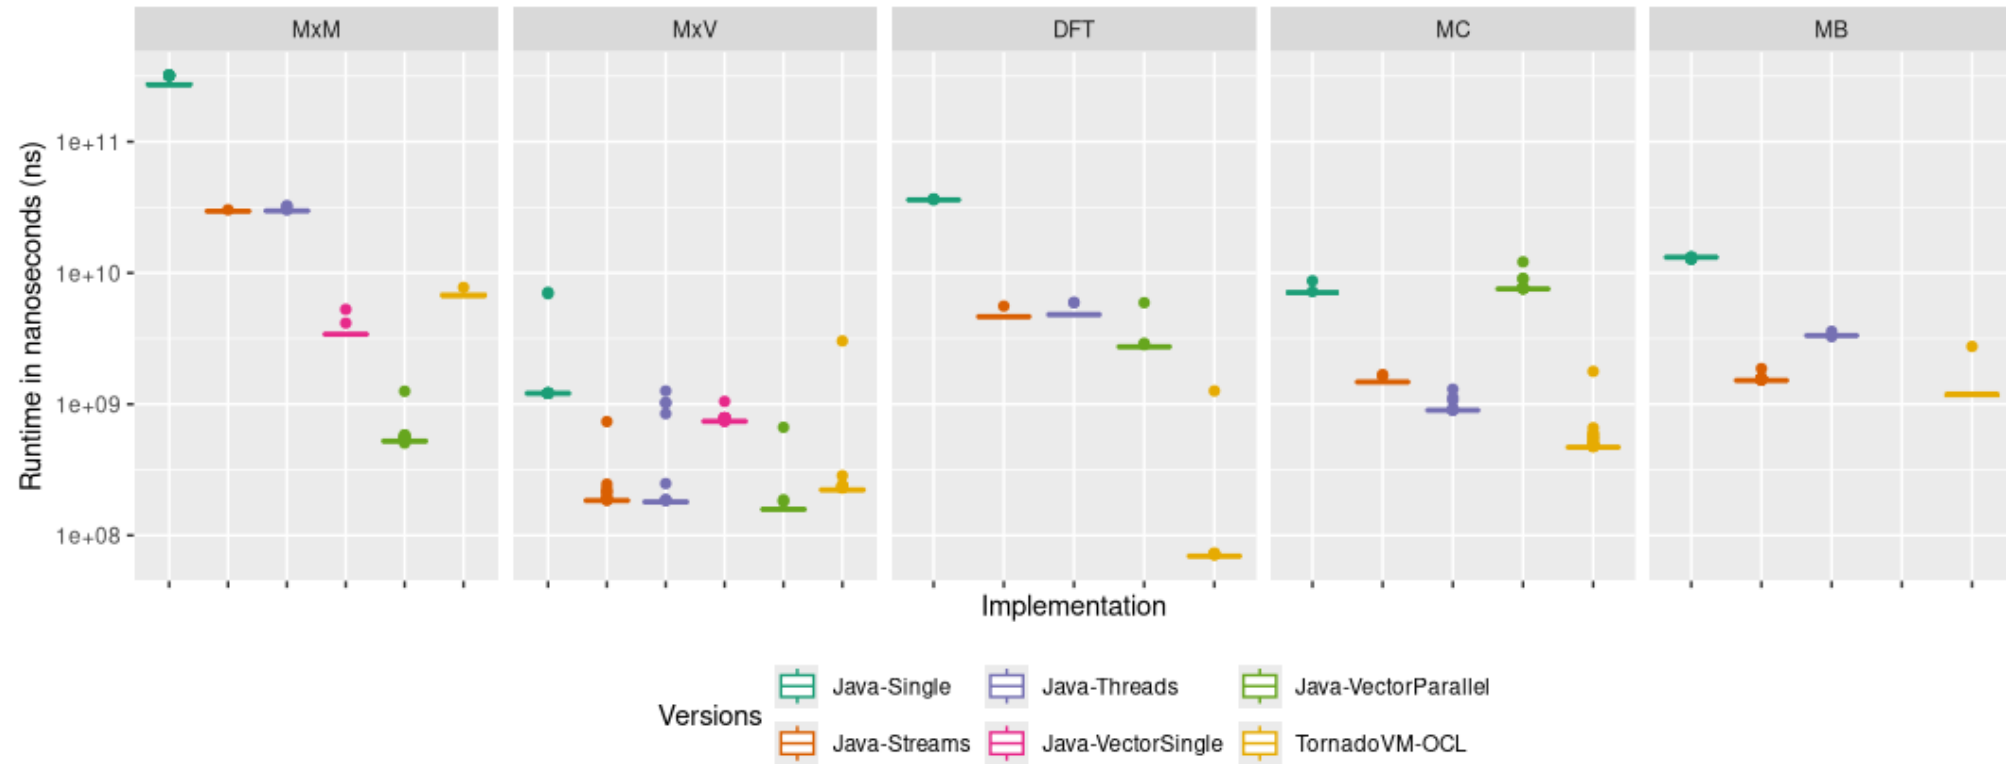
```
136 sumimag += -1 * vInReal.mul(FloatVector.fromArray(species, sinAngles, 0))
        .add(vInImag.mul(FloatVector.fromArray(species, cosAngles, 0)))
        .reduceLanes(VectorOperators.ADD);
```



# Runtime Distribution on RISC-V Spacemit K1 (Lichee PI3A)

Performance of MxM using TornadoVM OCK vs Java on RISC-V Spacemit K1.

The lower, the better

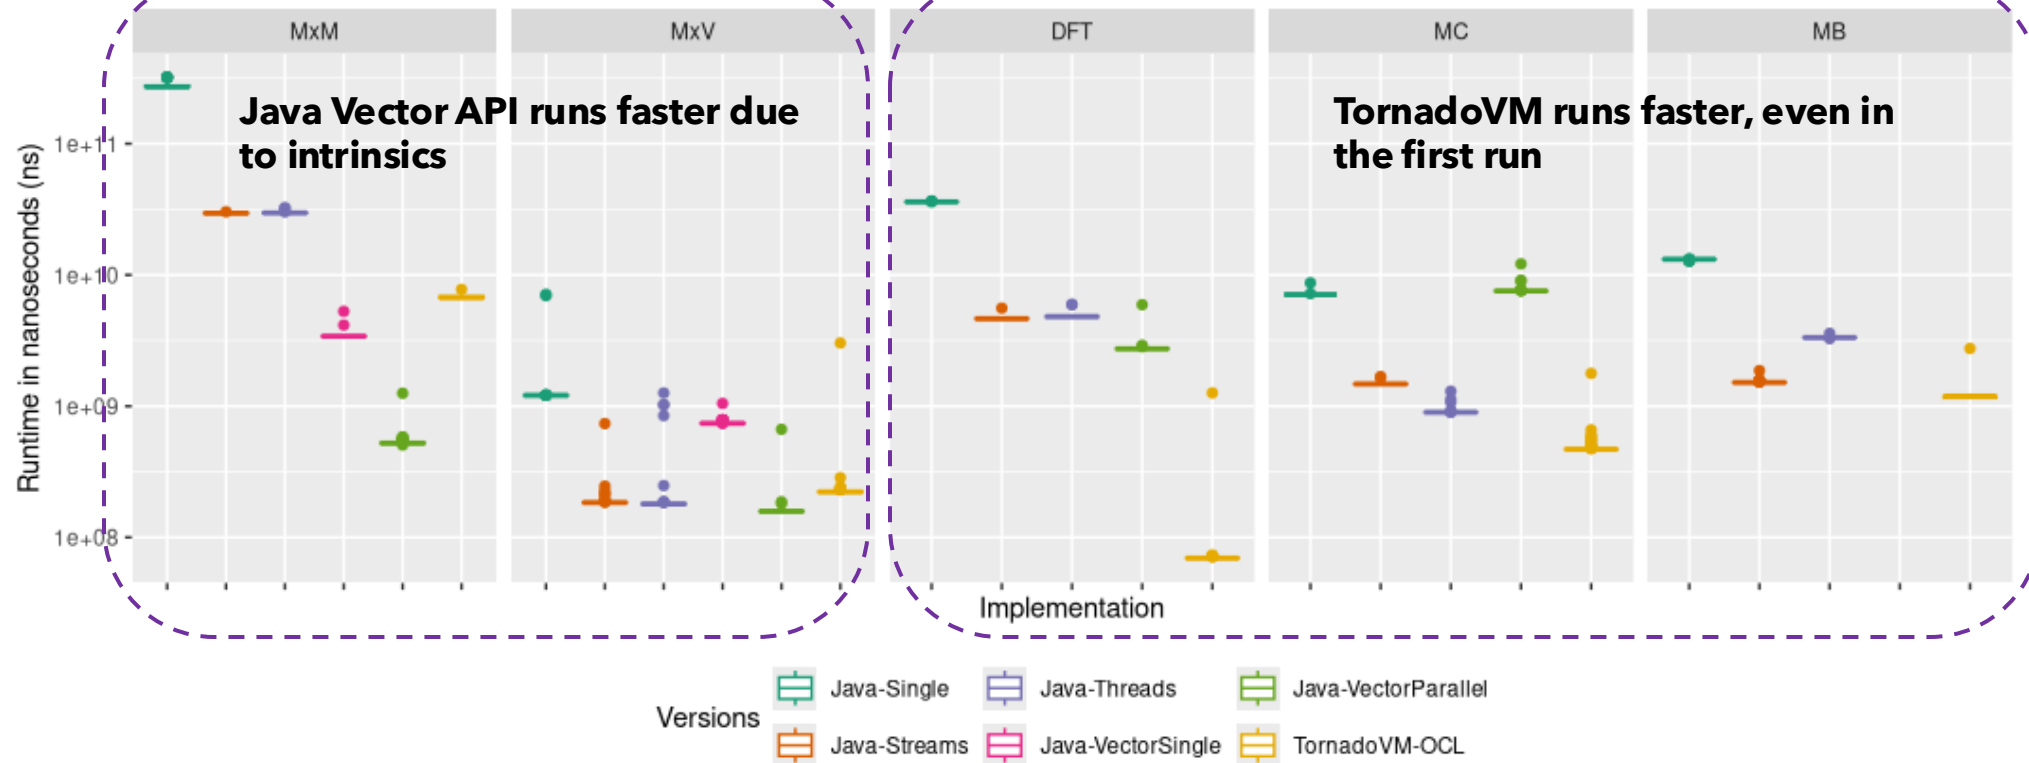


TornadoVM Version: 05539e7

# Runtime Distribution on RISC-V Spacemit K1 (Lichee PI3A)

Performance of MxM using TornadoVM OCK vs Java on RISC-V Spacemit K1.

The lower, the better

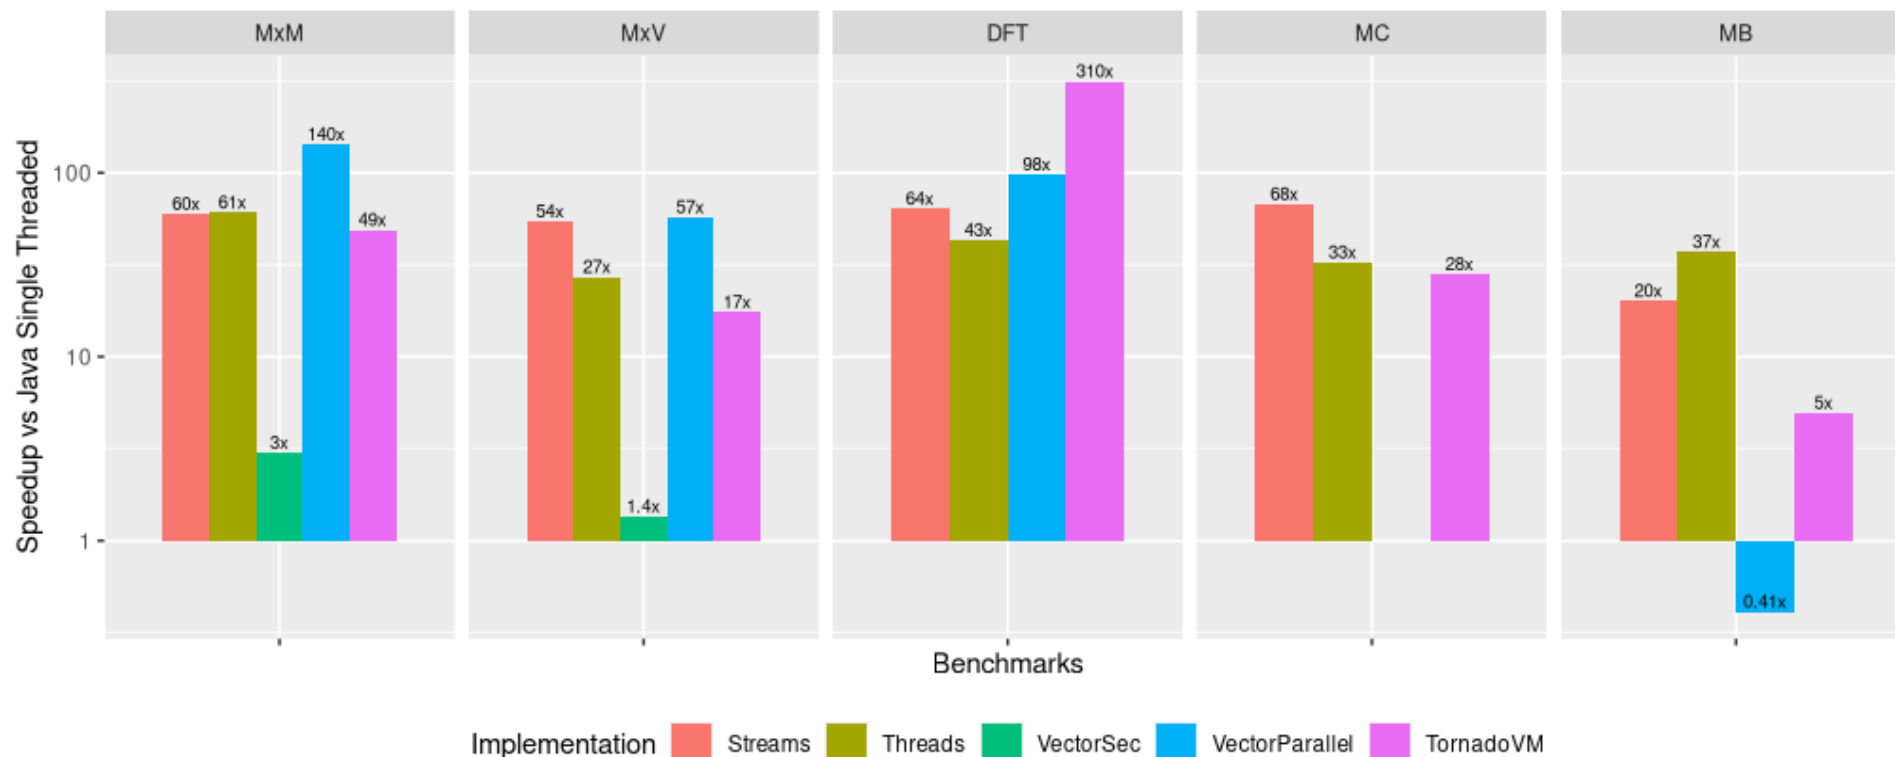


TornadoVM Version: 05539e7

# [ARM CPUs] ARM Neoverse V2 (Grace Hopper)

Performance of TornadoVM on ARM Neoverse-V2 (72 cores)

The higher, the better.



TornadoVM Version: 649ba9c

ARM Neoverse V2 CPU (NVIDIA Grace Hopper)

Ubuntu 22.04.4 LTS

Kernel 6.2.0-1015-nvidia-64k

OCK 4.0: [5be5a8d](#)

TornadoVM 1.0.11-dev:

[649ba9c](#)

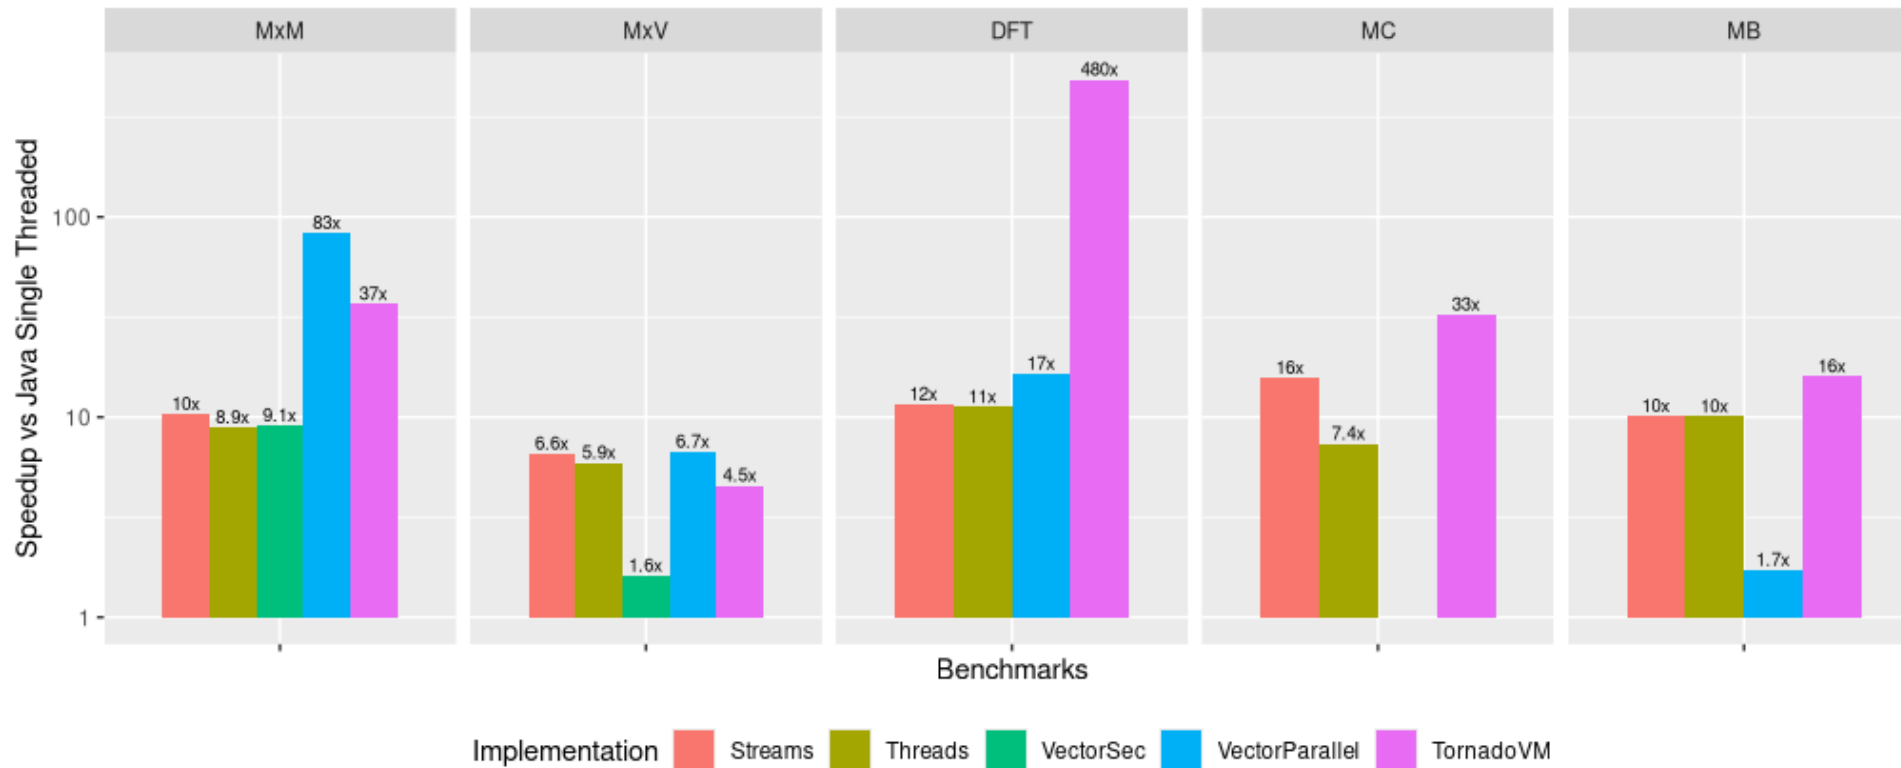
OpenJDK 21.0.3+7-LTS

Java multi-threaded is very competitive in this platform, since OCK does not use SVE instructions. Explicit vectorization is a mixed: if auto-vectorization works, there is no major benefit, at least for these benchmarks.

# [INTEL CPUs] TornadoVM + OCK/oneAPI

Performance of TornadoVM on Intel i7-12700K

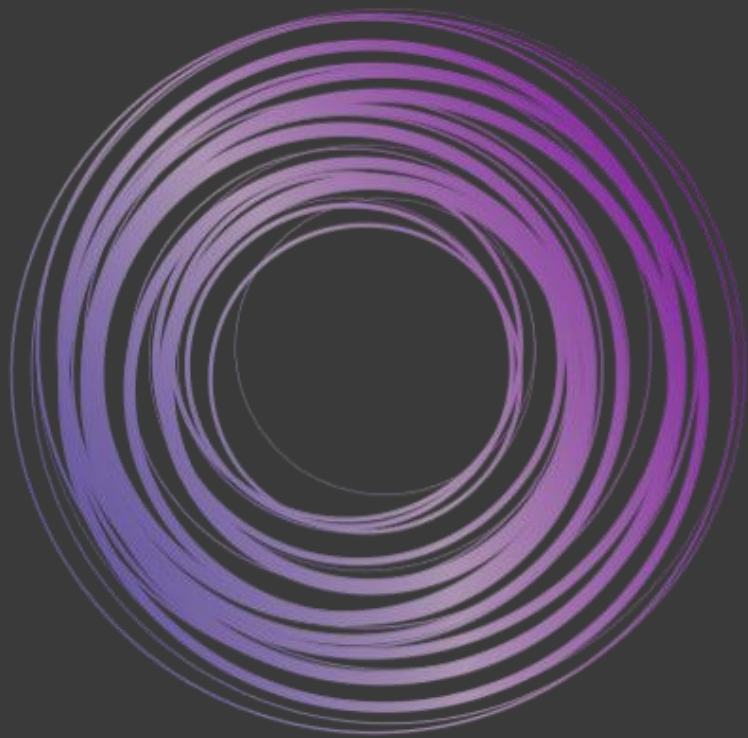
The higher, the better.



TornadoVM Version: 649ba9c

Intel i7-12700K  
Fedora 41  
Kernel 6.2.15-200  
OCK 4.0: [a537ec99](#)  
TornadoVM 1.0.11-dev:  
[649ba9c](#)  
OpenJDK 21.0.6

For Intel CPUs, the Java JIT compiler generates vector instructions  
TornadoVM + OCK generates efficient code just for the kernel of interest

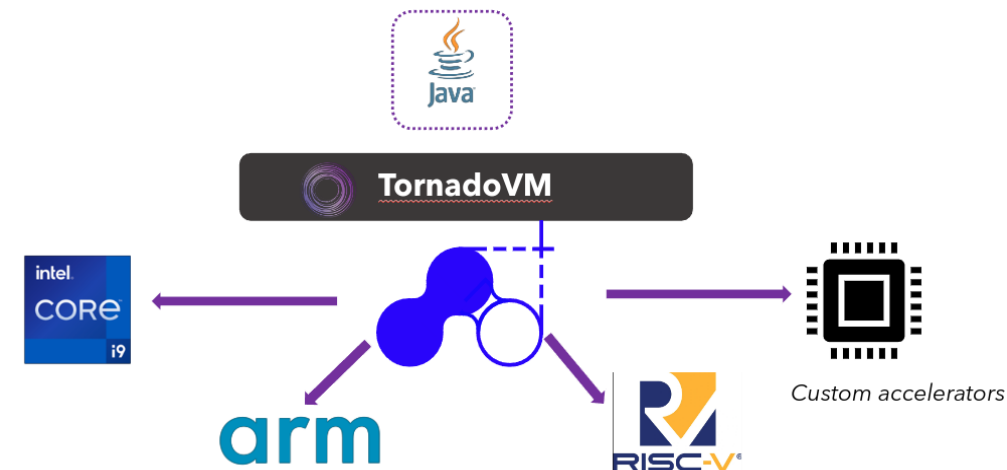
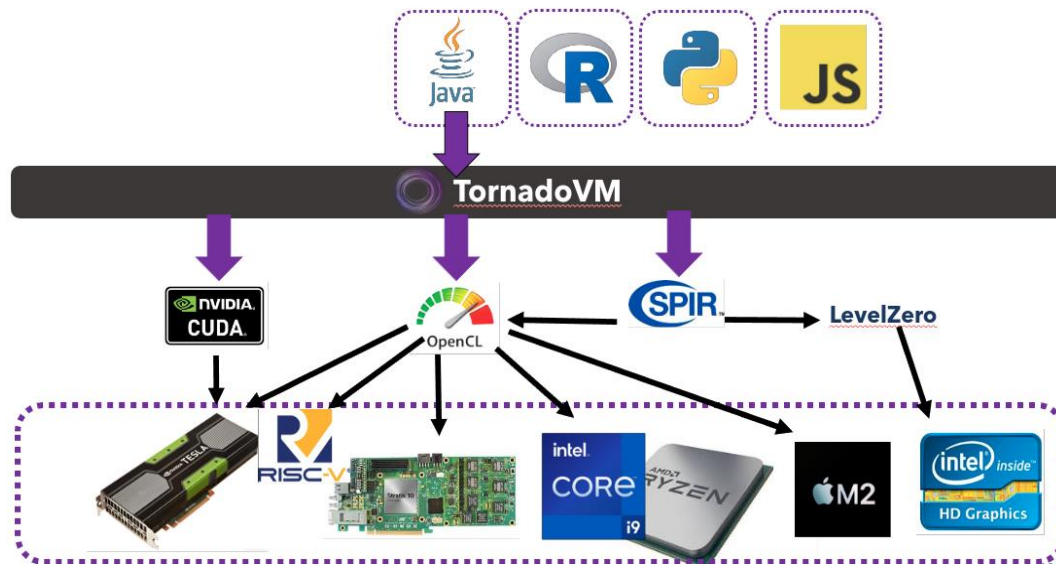


**TORNADOVM**

[www.tornadovm.org](http://www.tornadovm.org)

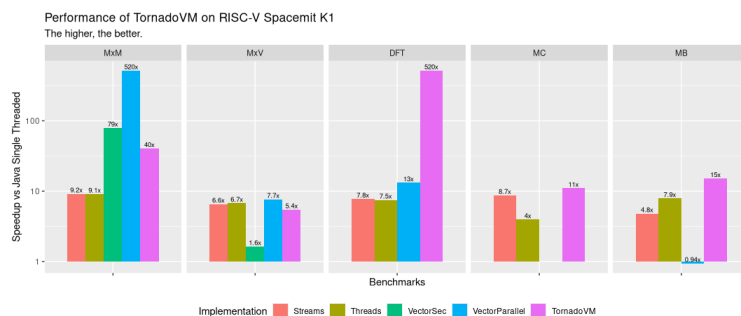
Conclusions

# Key Takeaways



Faster than Java  
multi-core

Usually slower than  
explicit Java  
vectorization



  
tornadovm.org



# Collaborations and Projects





Thank you!



[juan.fumero@manchester.ac.uk](mailto:juan.fumero@manchester.ac.uk)



[@jfumero.bsky.social](https://bsky.social/@jfumero)



[@snatverk](https://twitter.com/snatverk)



Discussions